



PES
UNIVERSITY
ONLINE

RISC V Architecture

Mahesh Awati & Vinay Reddy

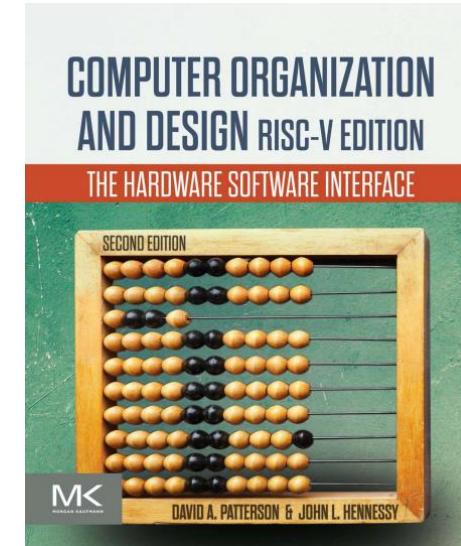
Department of Electronics and Communication Engg.

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering



Computer Abstractions and Technology

Syllabus/Topics



1. Introduction
2. Seven Great Ideas in Computer Architecture
3. Below Your Program
4. Technologies for Building Processors and Memory
5. Performance
6. The Power Wall
7. The Sea Change: The Switch from Uniprocessors to Multiprocessors
8. Real Stuff: Benchmarking the Intel Core i7
9. Going Faster: Matrix Multiply in Python
10. Fallacies and Pitfalls

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 1 – Introduction

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Can you think of a place where computers are not used today?



Applications that were economically infeasible suddenly become practical.

In the recent past, few applications were “computer science fiction.”



What where those “computer science fiction” applications?

1. Computers in automobiles

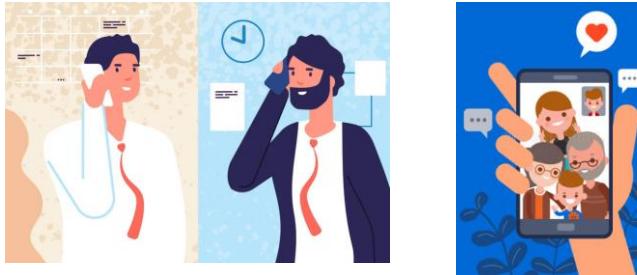
- Reduce pollution
- Improve fuel efficiency via engine controls
- Increase safety through
 - blind spot warnings,
 - lane departure warnings,
 - moving object detection, and
 - air bag inflation to protect occupants in a crash.



What where those “computer science fiction” applications?

2. Cell phones

Allowing person-to-person communication to almost anyone anywhere in the world.



3. Human genome project

You will soon be able to acquire your own genome, allowing medical care to be tailored to you.

Computer Abstractions and Technology

Introduction

What where those “computer science fiction” applications?

4. World Wide Web

The web has replaced libraries and newspapers.



5. Search engines

As the content of the web grew in size and in value, finding relevant information became increasingly important - rely on search engines



Is that all?

More of it....

Today's science fiction suggests tomorrow's killer applications: already on their way are –

- Glasses that augment reality



- The cashless society



- Cars that can drive themselves

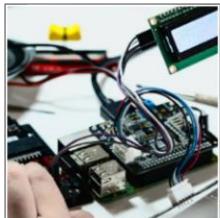
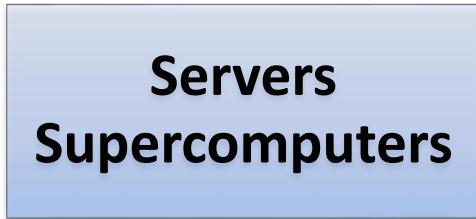


Traditional Classes of Computing Applications and Their Characteristics

Broadly computers are used in three dissimilar classes of applications



**Personal
Computers**



**Embedded
Computers**

Traditional Classes of Computing Applications and Their Characteristics

1. Personal computers (PCs):

- A computer designed for use by an individual.
- Personal computers emphasize delivery of good performance to single users.
- Low costs and usually execute third-party software.



Personal Laptops

Traditional Classes of Computing Applications and Their Characteristics

2. Servers:

- A computer used for running larger programs for multiple users.
- Typically accessed only via a network.



Supercomputer:

- A class of computers with the highest performance and cost.
- They are configured as servers.
- These are used for high-end scientific and engineering calculations, such as weather forecasting, oil exploration, protein structure determination, and other large-scale problems

The Cray-1 supercomputer was designed by American electrical engineer and supercomputer architect Seymour Cray in the 1970s.

Traditional Classes of Computing Applications and Their Characteristics

3. Embedded computer:

- A computer inside another device used for running one predetermined application or collection of software.
- A popular term today is Internet of Things (IoT) which suggests many small devices that all communicate wirelessly over the Internet.



Purpose-built computing platforms designed for a specific task or fixed functionality

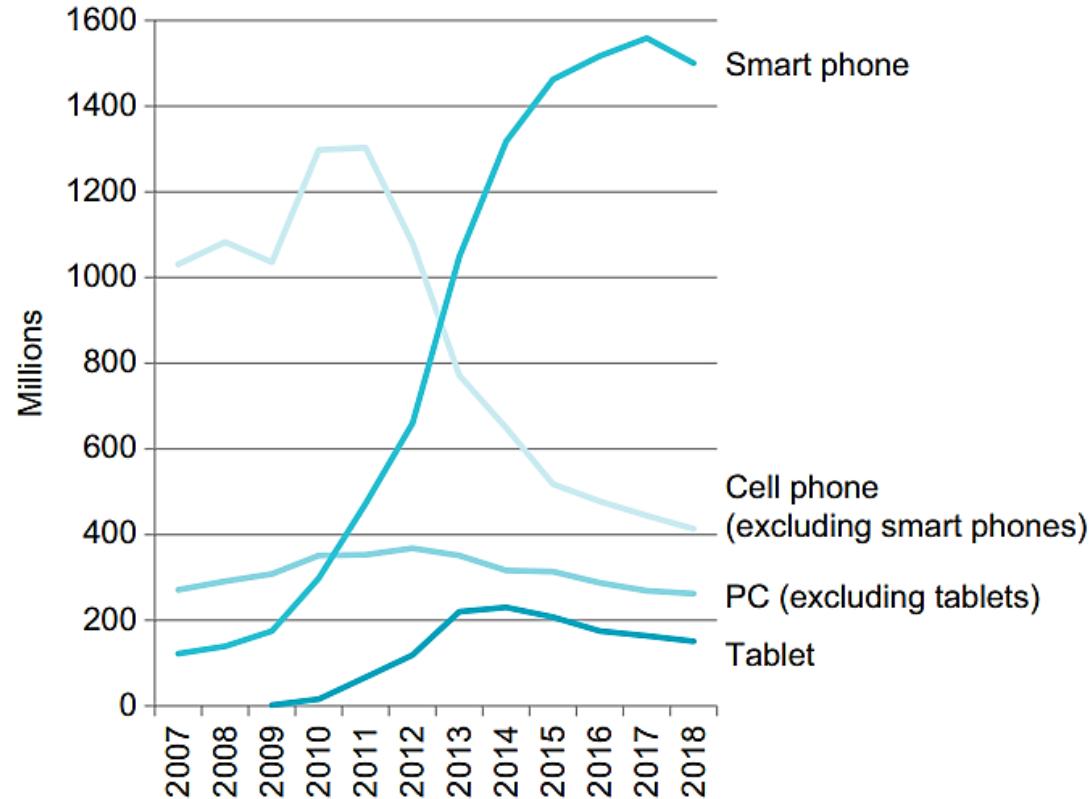
Can you call FPGA as an embedded computer?



The 2^x vs. 10^y bytes

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	10^3	kibibyte	KiB	2^{10}	2%
megabyte	MB	10^6	mebibyte	MiB	2^{20}	5%
gigabyte	GB	10^9	gibibyte	GiB	2^{30}	7%
terabyte	TB	10^{12}	tebibyte	TiB	2^{40}	10%
petabyte	PB	10^{15}	pebibyte	PiB	2^{50}	13%
exabyte	EB	10^{18}	exbibyte	EiB	2^{60}	15%
zettabyte	ZB	10^{21}	zebibyte	ZiB	2^{70}	18%
yottabyte	YB	10^{24}	yobibyte	YiB	2^{80}	21%
ronnabyte	RB	10^{27}	robibyte	RiB	2^{90}	24%
queccabyte	QB	10^{30}	quebibyte	QiB	2^{100}	27%

Post-PC Era



Graph shows the rapid growth over time of tablets and smart phones versus that of PCs and traditional cell phones.

FIGURE 1.2 The number manufactured per year of tablets and smart phones, which reflect the post-PC era, versus personal computers and traditional cell phones.

Post-PC Era

Replacing the PC - The Personal Mobile Device (PMD)

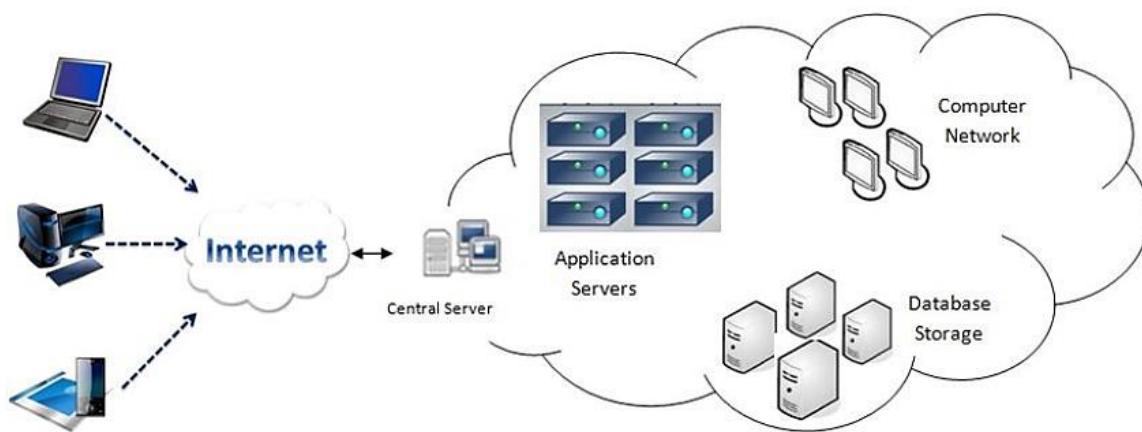
- Small wireless devices to connect to the Internet.
- They rely on batteries for power, and software is installed by downloading apps.
- Conventional examples are smart phones and tablets.
- Unlike PCs, they no longer have a keyboard and mouse



Post-PC Era

Taking over from the conventional server - Cloud Computing

- Refers to large collections of servers that provide services over the Internet.
- Some providers rent dynamically varying numbers of servers as a utility.
- Relies upon giant datacenters that are now known as Warehouse Scale Computers (WSCs).



Companies like Amazon and Google build these WSCs containing 50,000 servers and then let companies rent portions of them so that they can provide software services to PMDs without having to build WSCs of their own.

Post-PC Era

Software as a Service (SaaS)

- Delivers software and data as a service over the Internet
- Done via a thin program such as a browser that runs on local client devices, instead of binary code that must be installed, and runs wholly on that device.
- Examples include web search and social networking.



- *A SaaS application may be accessed through a browser or through an app.*
- *Online email applications that users access through a browser, such as Gmail and Office 365, are common examples of SaaS applications.*

What You Can Learn in this Course

- How are programs written in a high-level language, such as C or Java, translated into the language of the hardware, and how does the hardware execute the resulting program?



Understanding the aspects of both the hardware and software

- What is the interface between the software and the hardware, and how does software instruct the hardware to perform needed functions?



Understanding how to write many kinds of software

What You Can Learn in this Course

- What determines the performance of a program, and how can a programmer improve the performance?



Software translation of that program into the computer's language, and the effectiveness of the hardware in executing the program.

- What techniques can be used by hardware designers to improve performance?



This course will introduce the basic concepts of modern computer design.

What You Can Learn in this Course

- What techniques can be used by hardware designers to improve energy efficiency?



What can the programmer do to help or hinder energy efficiency?

- What are the reasons for and the consequences of the switch from sequential processing to parallel processing?



This book gives the motivation, describes the current hardware mechanisms to support parallelism, and surveys the new generation of “multicore” microprocessors

Understanding Program Performance

The performance of a program depends on a combination of -

- the **effectiveness of the algorithms** used in the program.
- the **software systems used** to create and translate the program into machine instructions, and
- the **effectiveness of the computer** in executing those instructions, which may include input/output (I/O) operations.

Hardware or software component	How this component affects performance	Where is this topic covered?
Algorithm	Determines both the number of source-level statements and the number of I/O operations executed	Other books!
Programming language, compiler, and architecture	Determines the number of computer instructions for each source-level statement	Chapters 2 and 3
Processor and memory system	Determines how fast instructions can be executed	Chapters 4, 5, and 6
I/O system (hardware and operating system)	Determines how fast I/O operations may be executed	Chapters 4, 5, and 6

Refers to the chapters from the prescribed textbook



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 2 – Seven Great Ideas in Computer Architecture

Topic 3 – Below your Program

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering



Abstraction



Common Case Fast



Parallelism



Pipelining



Prediction

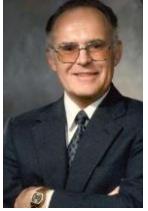


Hierarchy



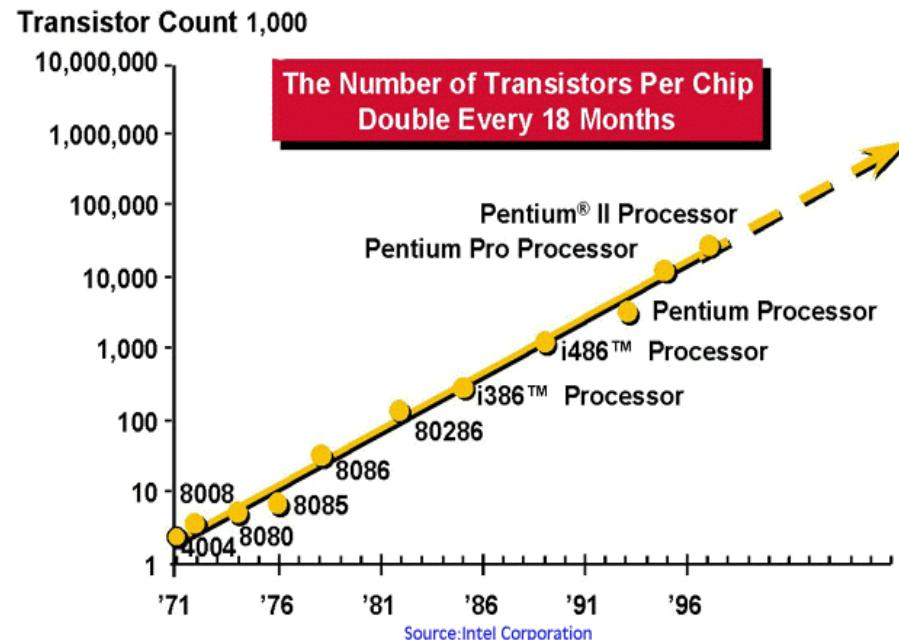
Dependability

Moore's Law



Gordon Moore

- Integrated circuit resources (Number of Transistors) would double every 18-24 Months.
- Computer architects must anticipate where the technology will be when the design finishes rather than design for where it starts.

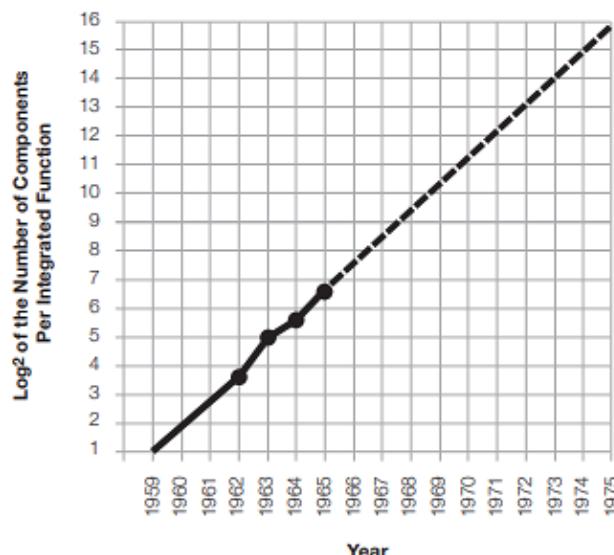
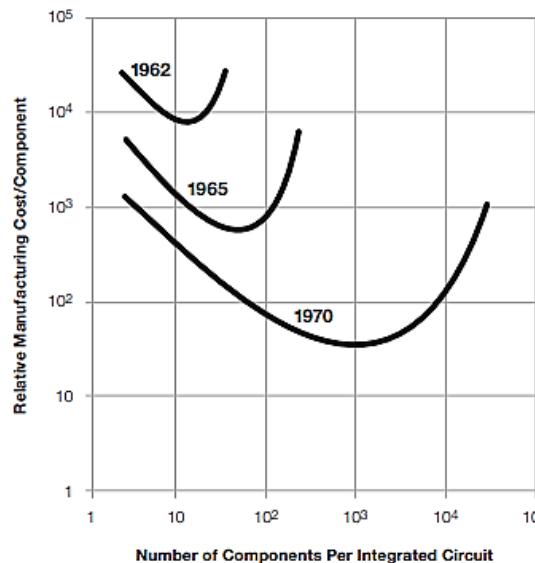


Moore's Law – The original paper - 1965

Paper: https://hasler.ece.gatech.edu/Published_papers/Technology_overview/gordon_moore_1965_article.pdf

Topic: Cramming more components onto integrated circuits

Prediction: With unit cost falling as the number of components per circuit rises, by 1975 economics may dictate squeezing as many as 65,000 components on a single silicon chip



*Where are we with respect to
Moore's law in 2022?*



Use Abstraction to Simplify Design/Below your Program

A major **productivity technique** for hardware and software is to **use abstractions** to characterize the design at different levels of representation, lower-level details are hidden to offer a simpler model at higher levels.

An example of the great idea of abstraction

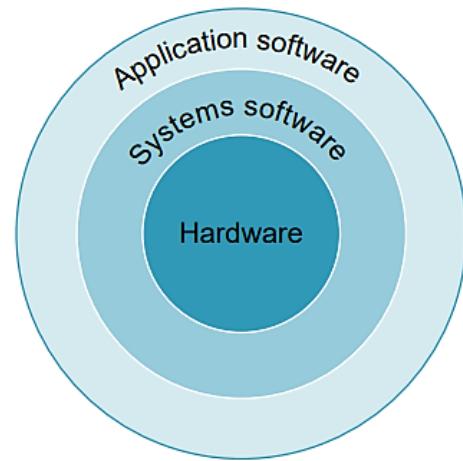
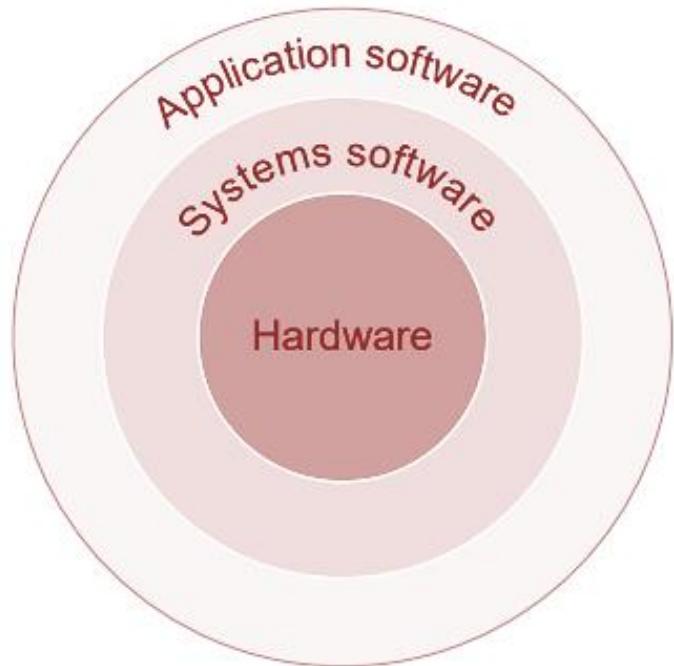


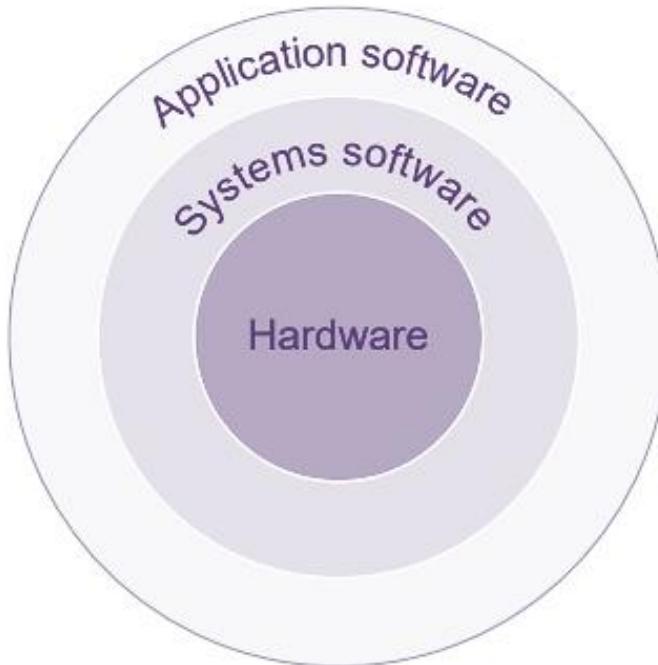
FIGURE 1.3 A simplified view of hardware and software as hierarchical layers, shown as concentric circles with hardware in the center and application software outermost.

Use Abstraction to Simplify Design/Below your Program



- **Hardware** is the center with **application software** being the outermost ring
- **Systems software** sitting between the hardware and the application software.
- System software
 - Operating system and a compiler

Use Abstraction to Simplify Design/Below your Program

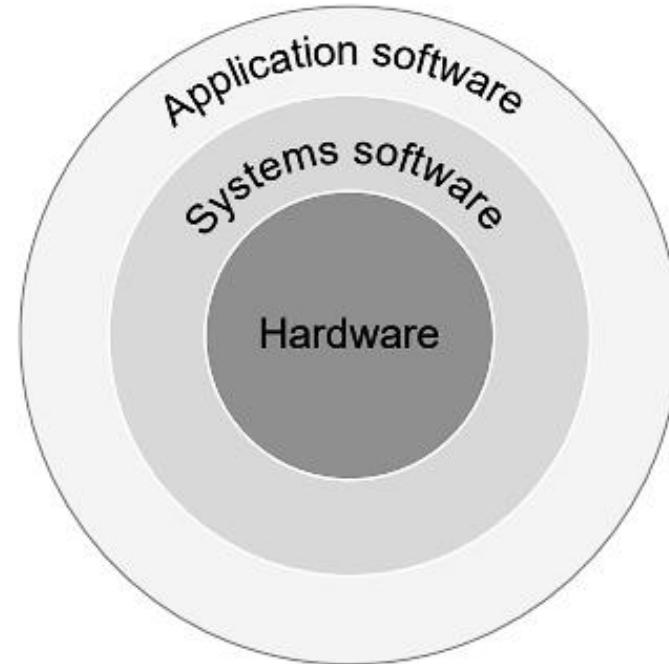


An **operating system** interfaces between a user's program and the hardware and provides a variety of services and supervisory functions.

Important functions are:

- Handling basic input and output operations
- Allocating storage and memory
- Providing for protected sharing of the computer among multiple applications using it simultaneously

Use Abstraction to Simplify Design/Below your Program



Compilers perform vital function

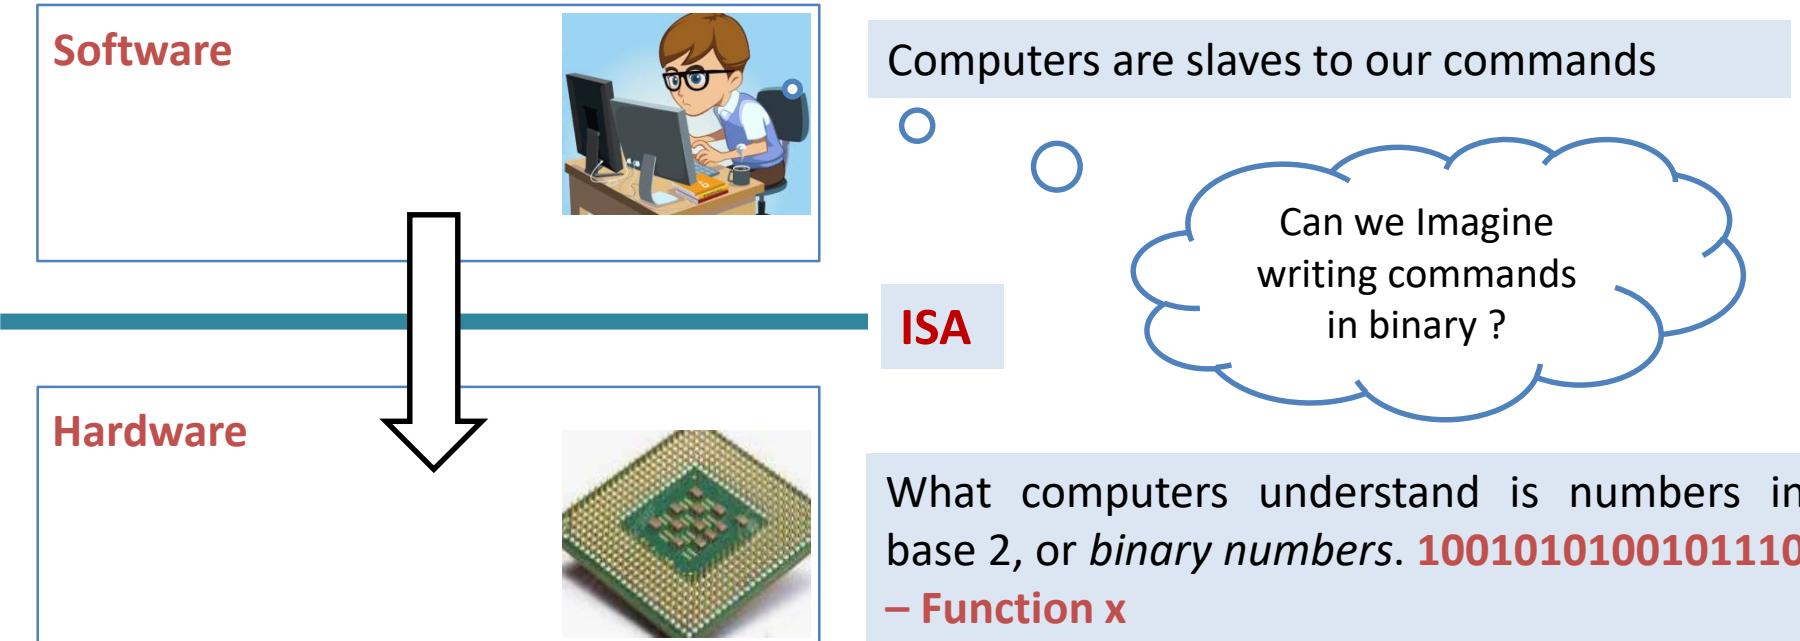
- Translation of a program written in a high-level language, such as C, C++, Java, or Visual Basic into instructions that the hardware can execute.

- Compiler - A program that translates high-level language statements into assembly language statements.

Seven Great Ideas in Computer Architecture

Use Abstraction to Simplify Design/Below your Program

From a High-Level Language to the Language of Hardware



Use Abstraction to Simplify Design/Below your Program

Things to know

Binary digit

Also called a bit

One of the two numbers in base 2 (0 or 1) that are the components of information.

Instruction

A command that computer hardware understands and obeys.

Assembler

A program that translates a symbolic version of instructions into the binary version.

Assembly language

A symbolic representation of machine instructions.

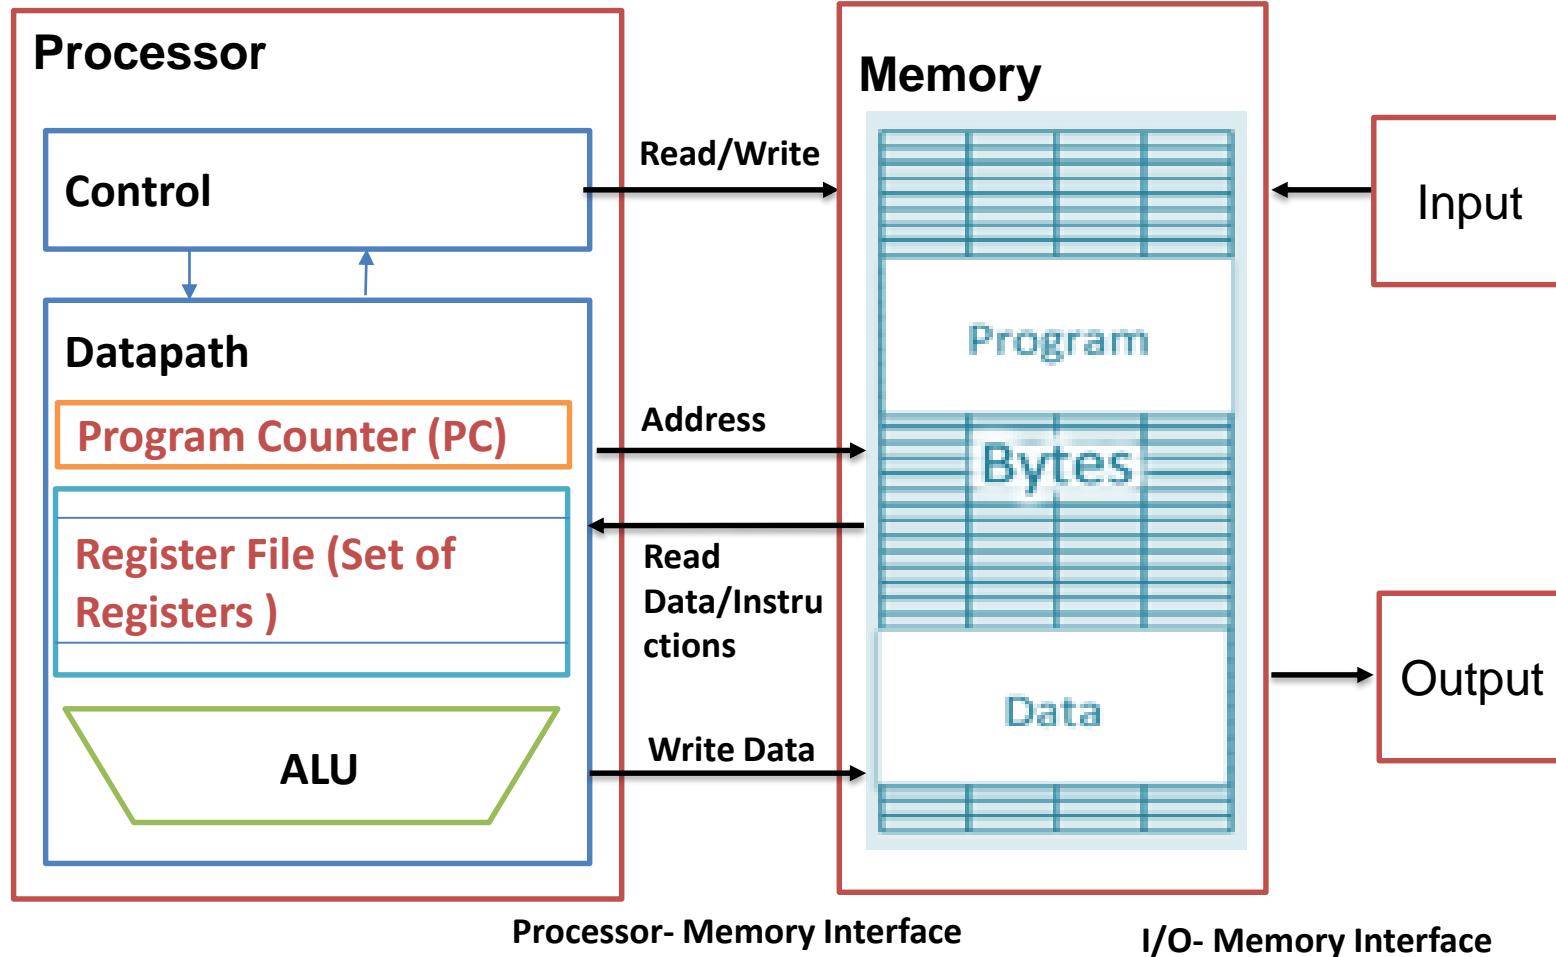
Machine language

A binary representation of machine instructions.

Computer Abstractions and Technology

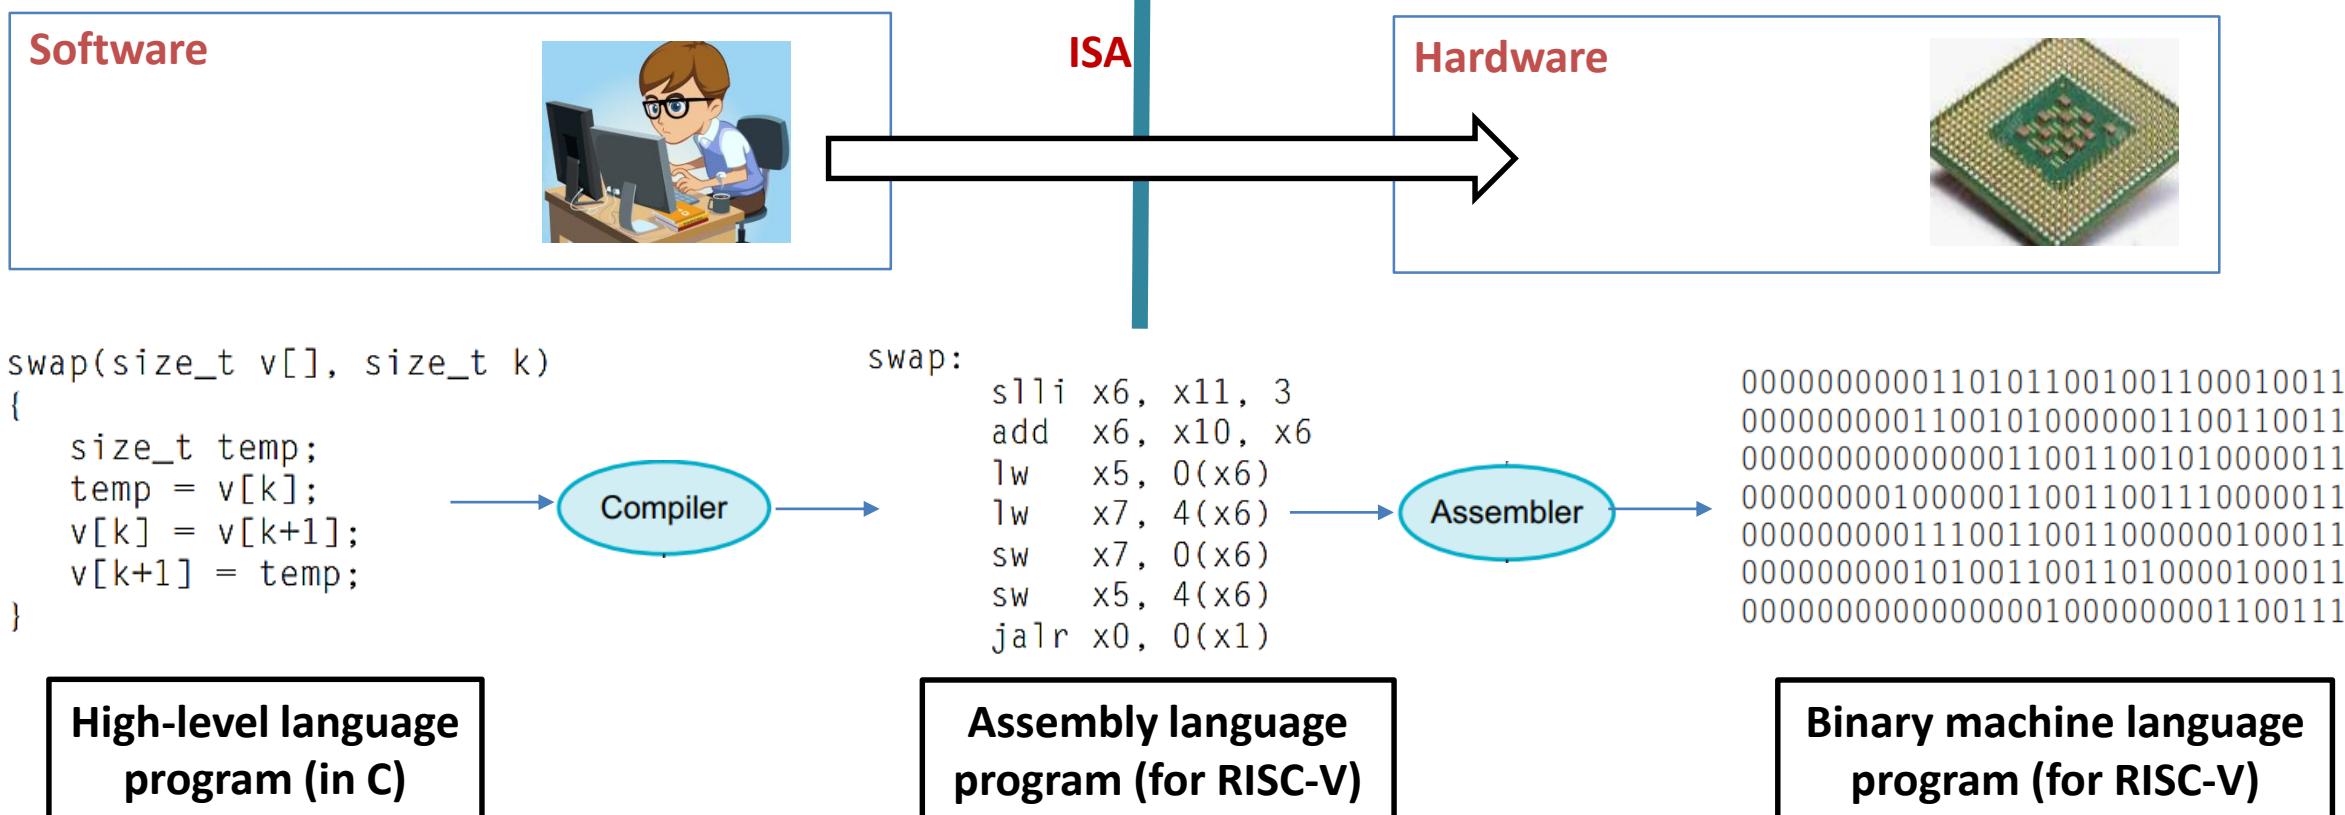
Components of Computer

Things to know



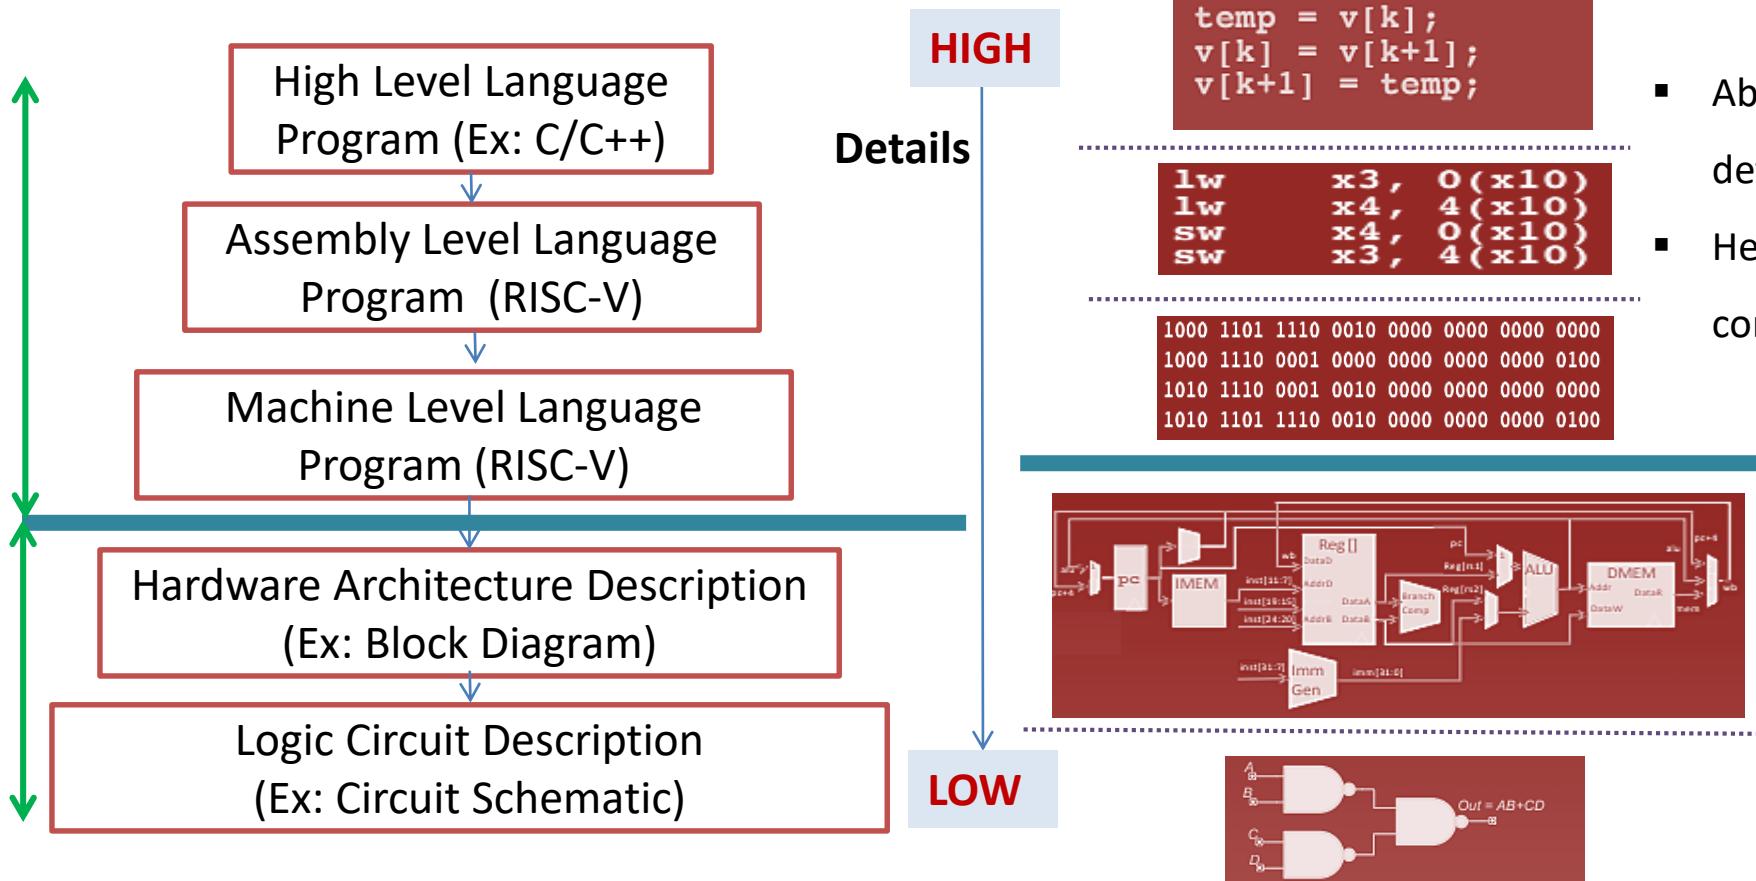
Use Abstraction to Simplify Design/Below your Program

From a High-Level Language to the Language of Hardware



Seven Great Ideas in Computer Architecture

Use Abstraction to Simplify Design/Below your Program



Use Abstraction to Simplify Design/Below your Program

Summary

- A compiler enables a programmer to write this high-level language expression: $A + B$
- The compiler would compile it into this assembly language statement: add A, B
- The assembler would translate this statement into the binary instructions that tell the computer to add the two numbers A and B.

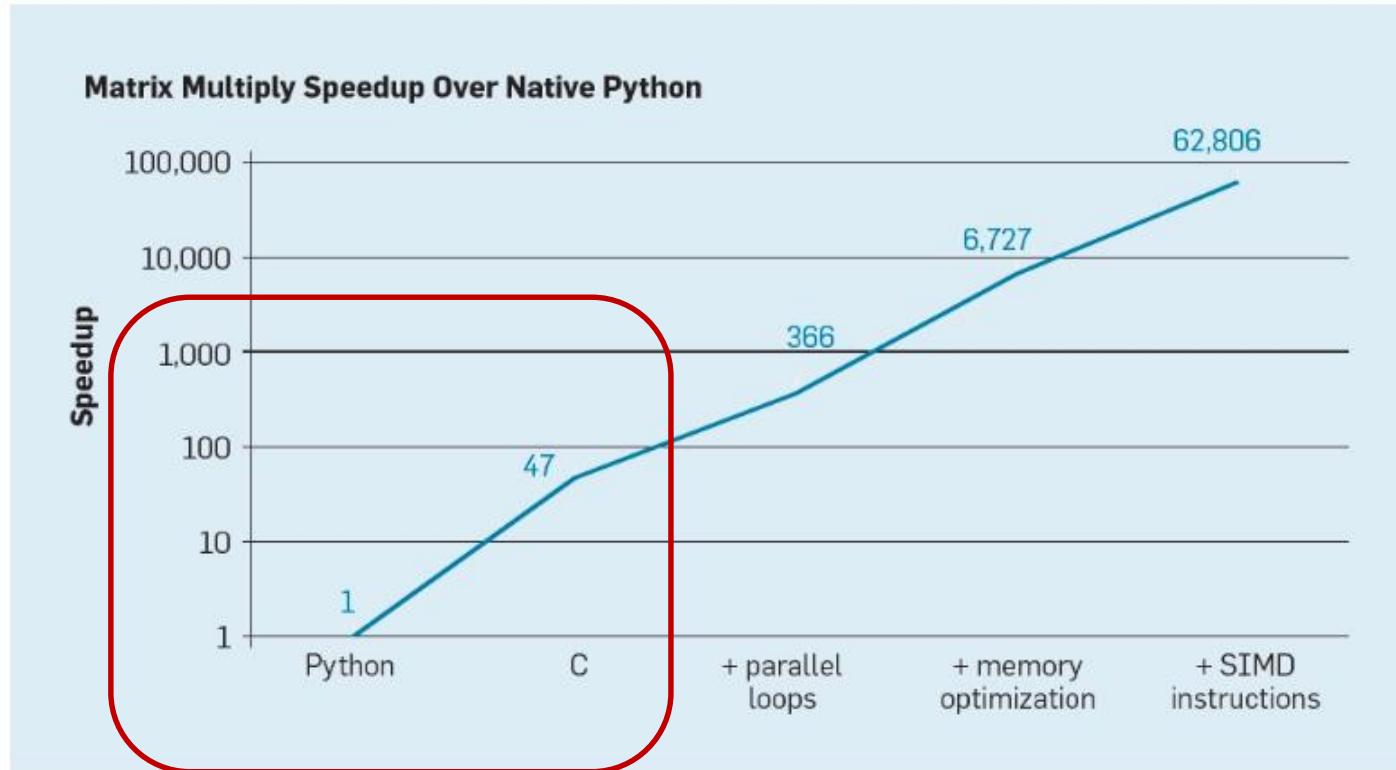
Benefits of High-level programming:

- Allow the programmer to think in a more natural language, using English words and algebraic notation
- Improved programmer productivity.
- Programming languages allow programs to be independent of the computer on which they were developed.

Computer Abstractions and Technology

Seven Great Ideas in Computer Architecture

Abstraction & inefficiency



Potential speedup of matrix multiply in Python for four optimizations



The C version of the code is 47 times faster than Python version.

Why??



Takeaway: Abstraction is Good but it should not result in inefficiency

Hierarchy of Memories

- Bigger memories have more capacity and reduce miss-rate.
- Their sizes often causes more delay per read or write operation.
- Bigger memories consume more energy to be refreshed to preserve the data.

These conflicting parameters can be solved by having a hierarchy of memories in the design.

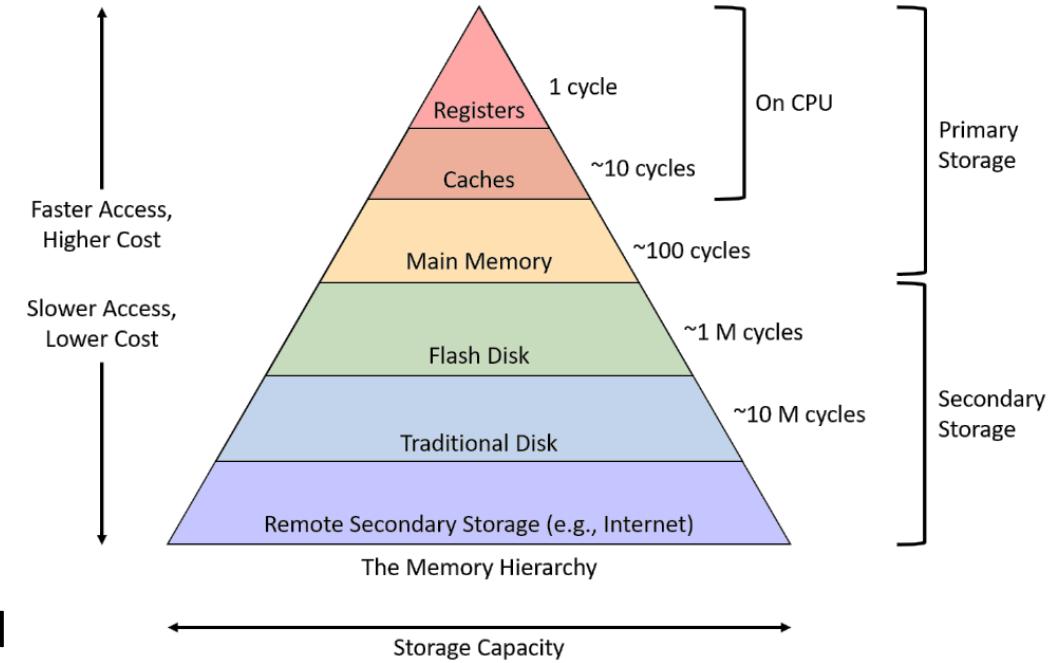
Hierarchy of Memories

Programmers want the memory to be

- Fast
- Large
- Cheap

Why memory is more important?

- Memory speed often shapes performance
- Capacity limits the size of problems that can be solved
- The cost of memory today is often the majority of computer cost





THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 2 – Seven Great Ideas in Computer Architecture

Part 2

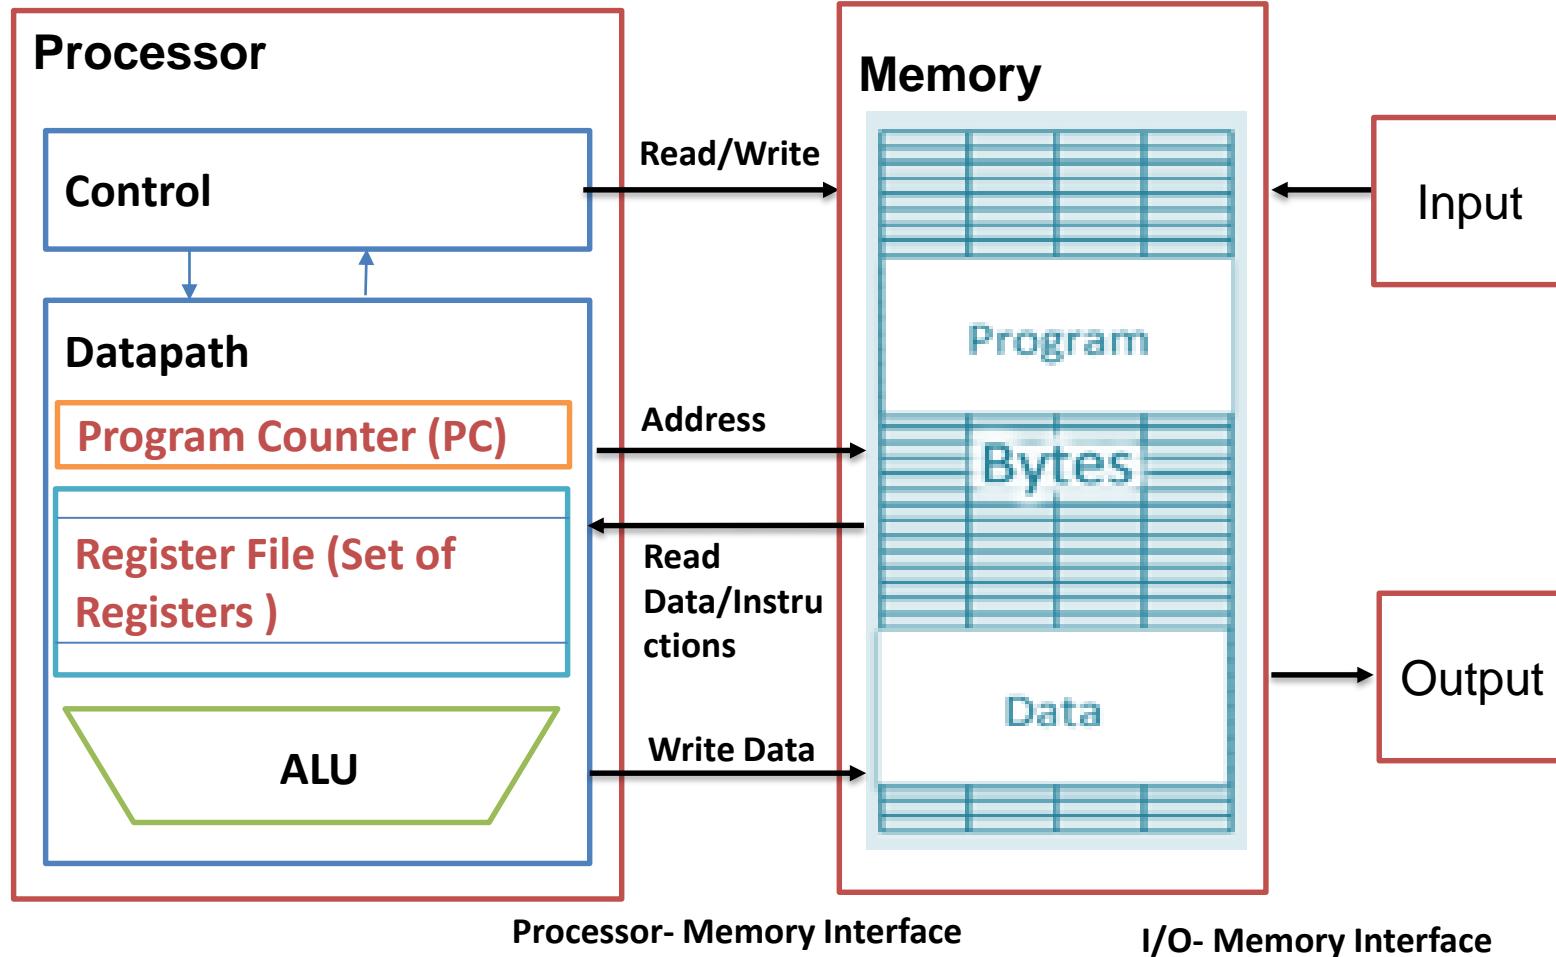
Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Computer Abstractions and Technology

Components of Computer

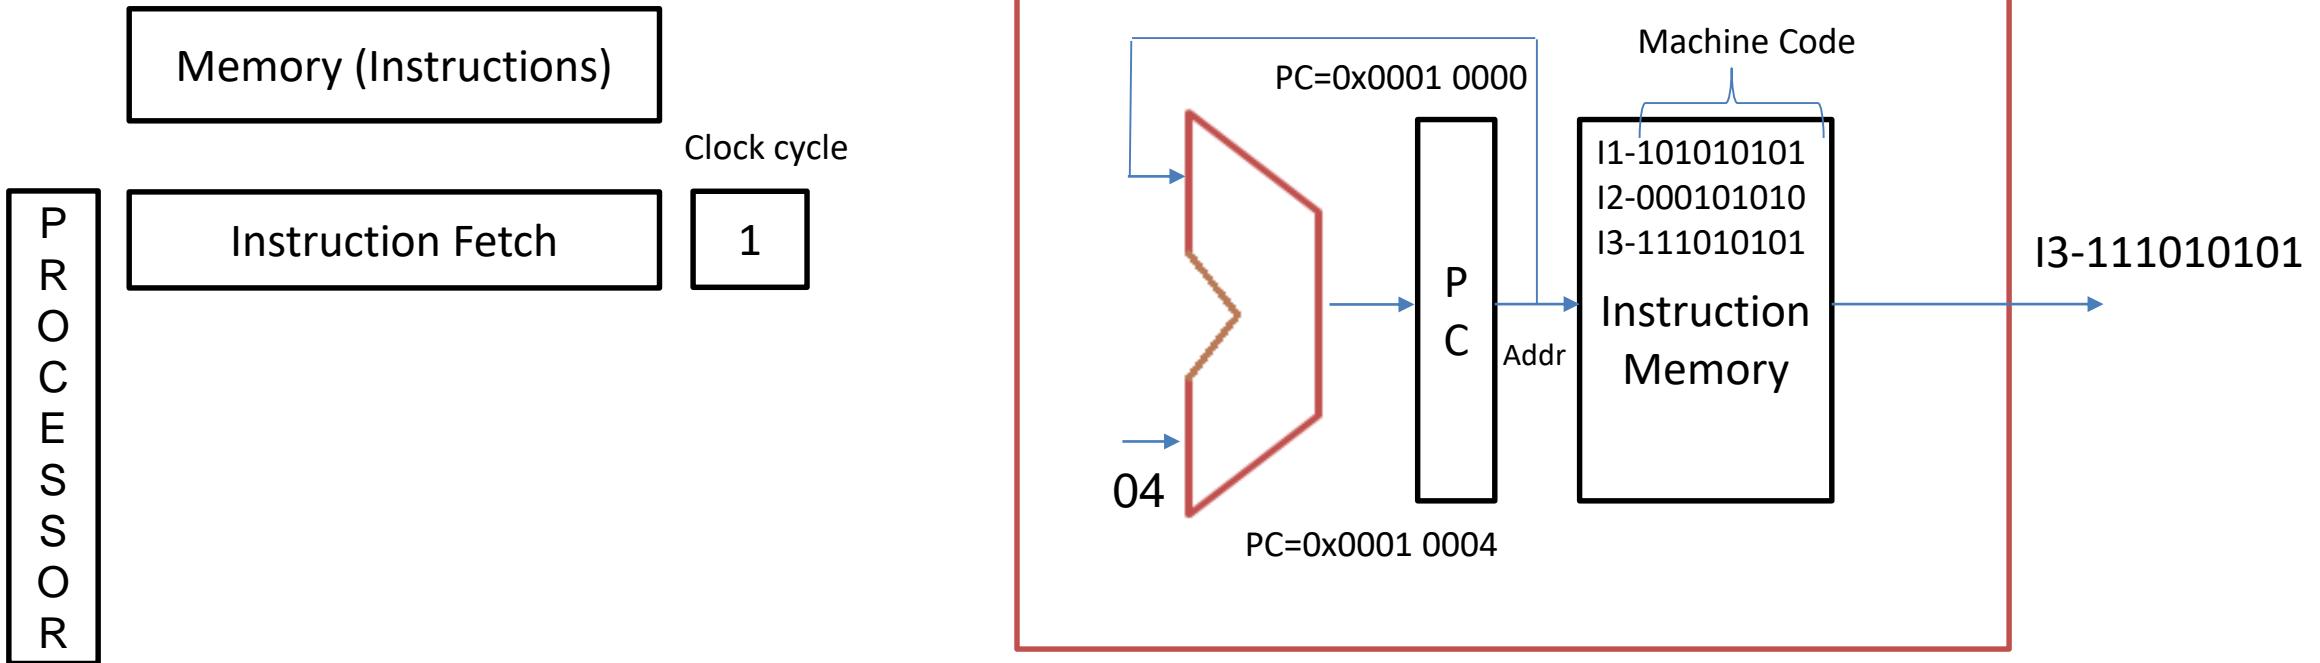
Things to know



Datapath: Anything that **stores data or operates on data within a processor** is called data path.

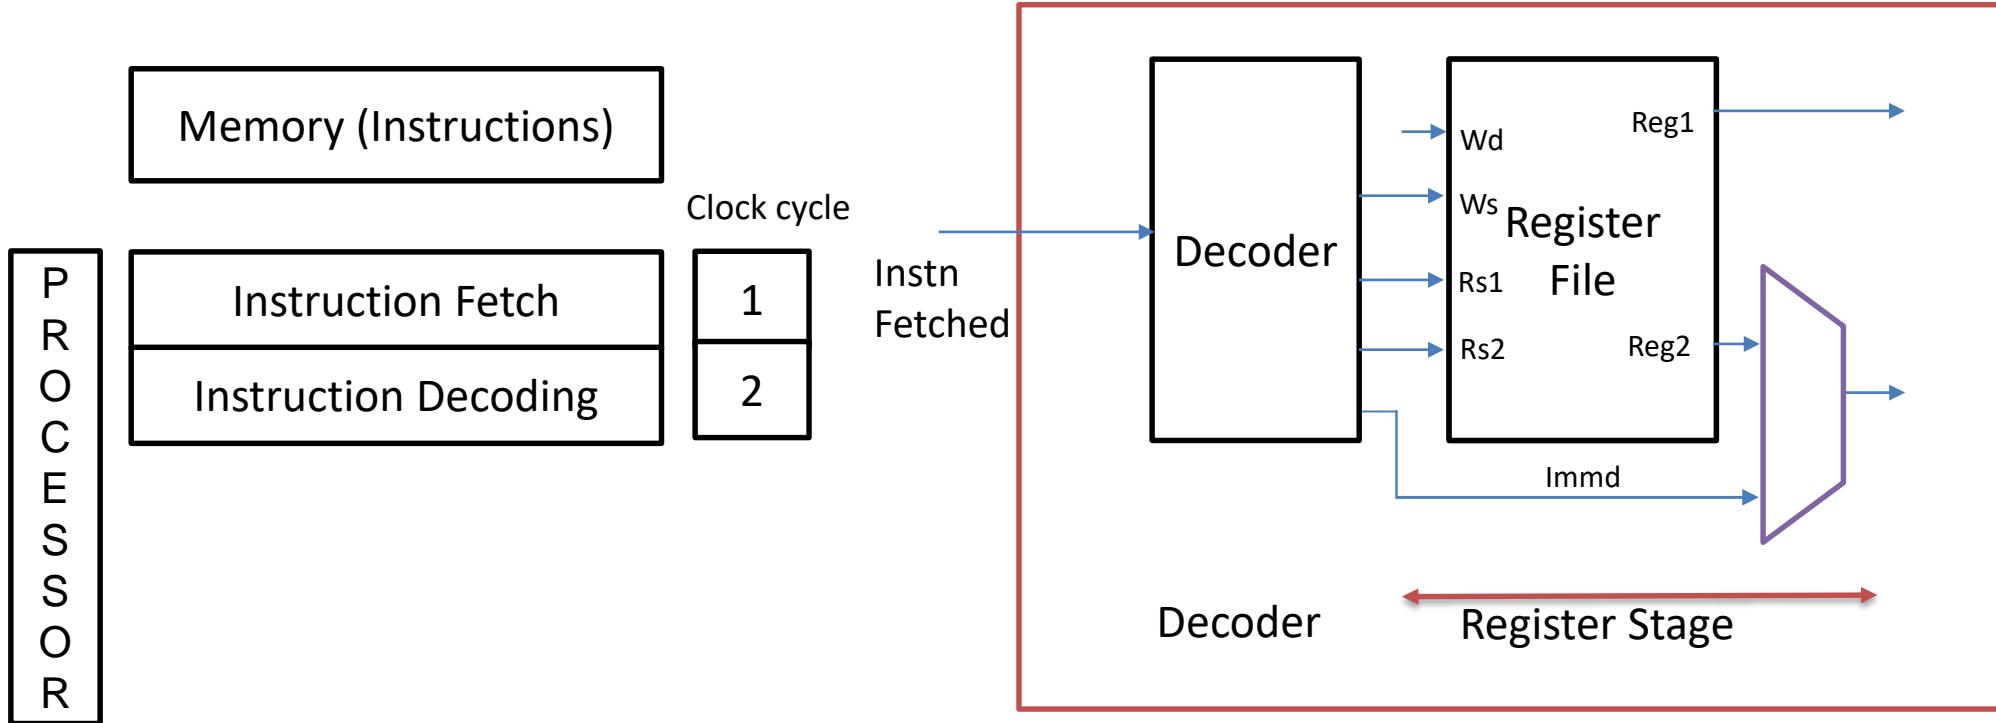
Performance via Pipelining

- First let us understand the steps involved in Non-Pipelined mode of execution.
- Instruction Execution involves different stages like Fetch , Decode, Execute, Memory access and Write back.



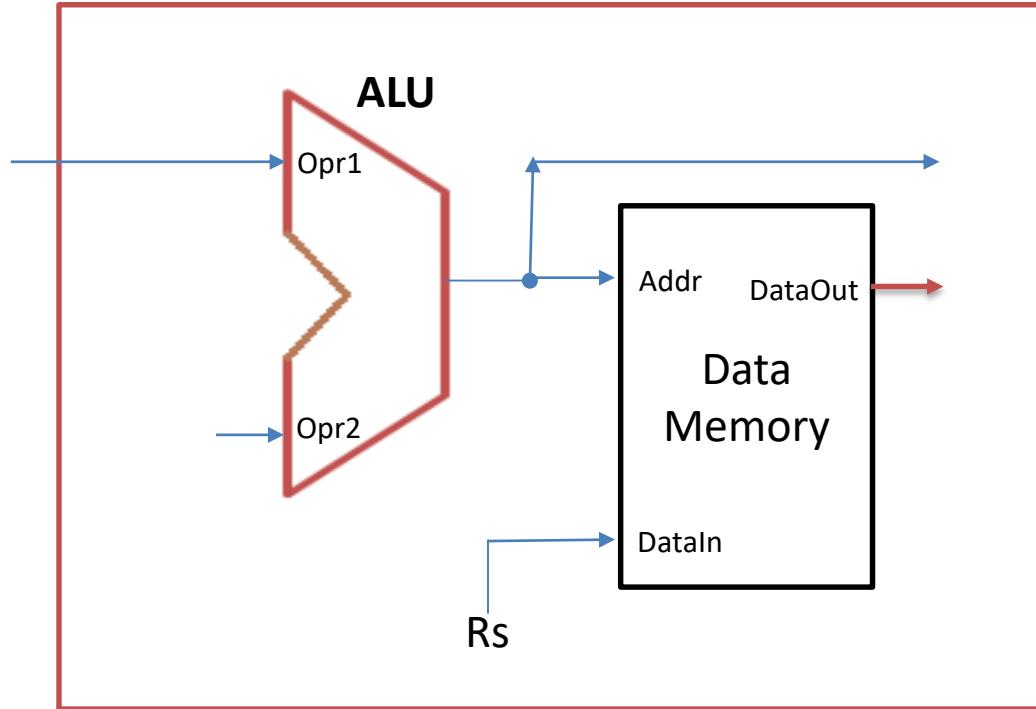
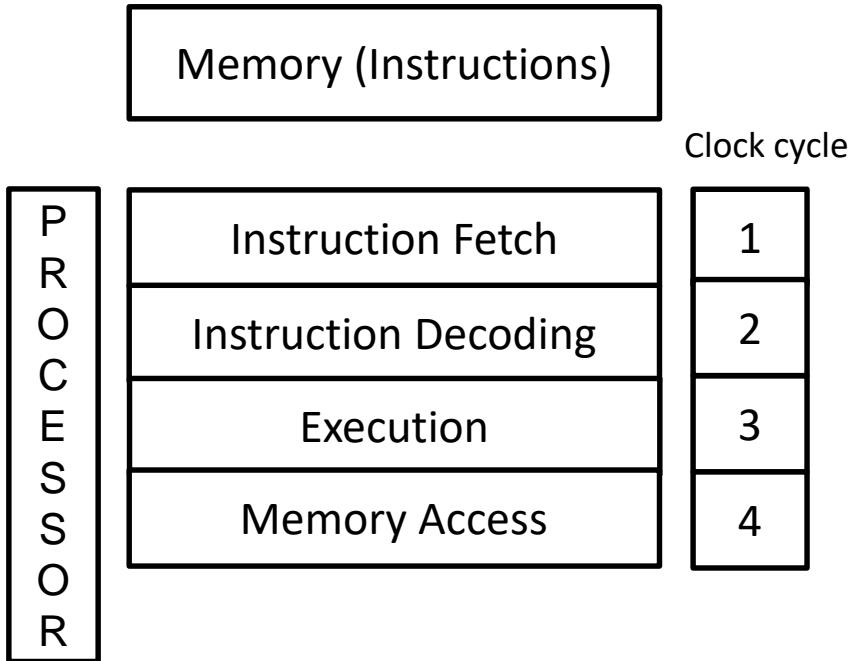
Performance via Pipelining

- First let us understand the steps involved in Non-Pipelined mode of execution.
- Instruction Execution involves different stages like Fetch , Decode, Execute, Memory access and Write back.



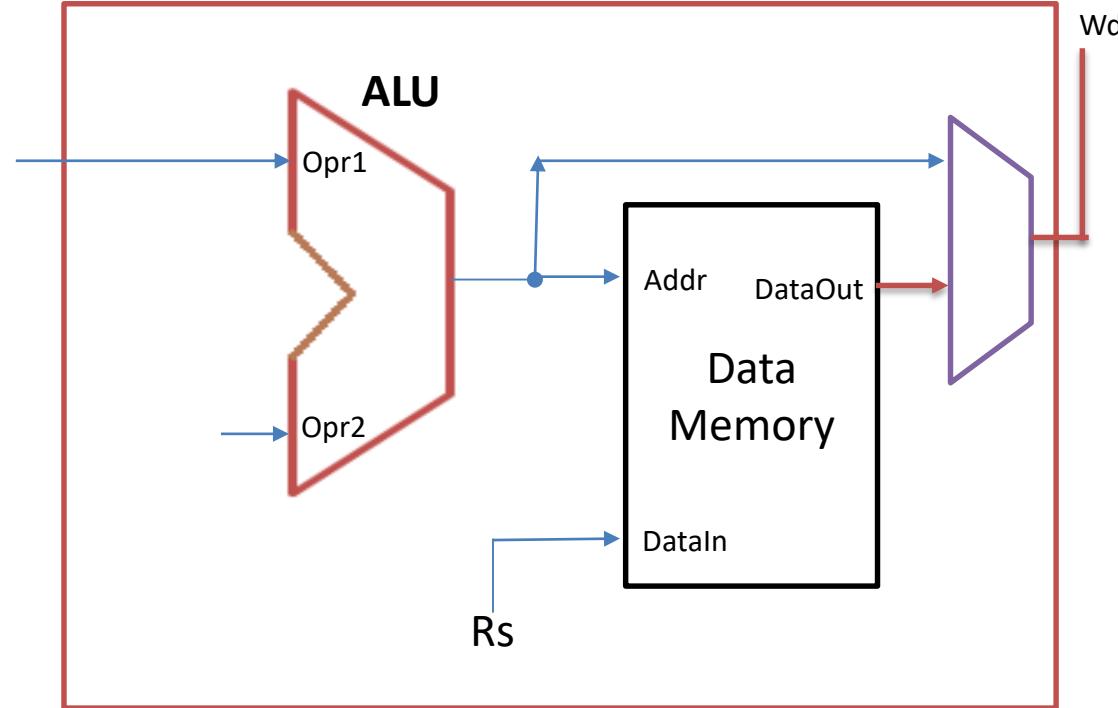
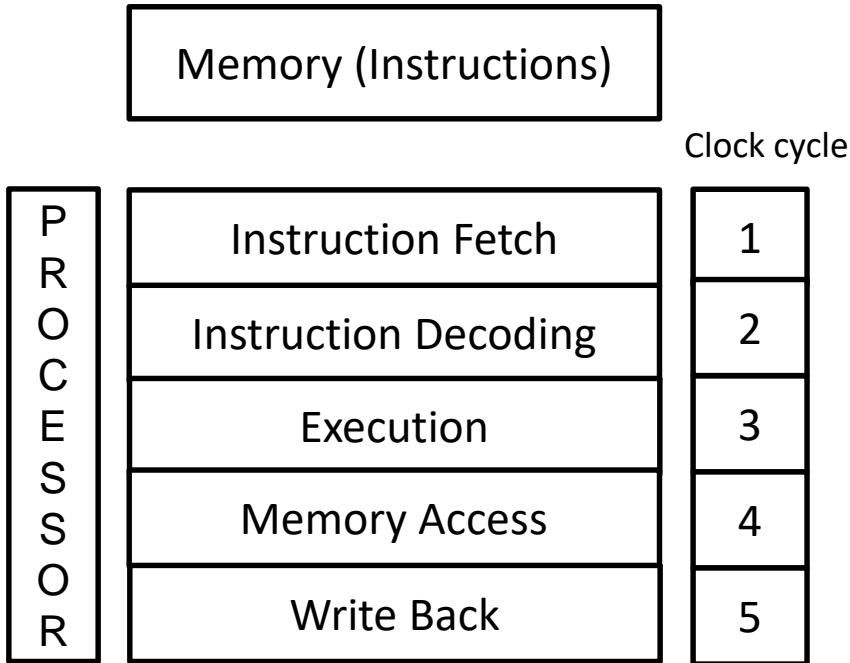
Performance via Pipelining

- First let us understand the steps involved in Non-Pipelined mode of execution.
- Instruction Execution involves different stages like Fetch , Decode, Execute, Memory access and Write back.



Performance via Pipelining

- First let us understand the steps involved in Non-Pipelined mode of execution.
- Instruction Execution involves different stages like Fetch , Decode, Execute, Memory access and Write back.

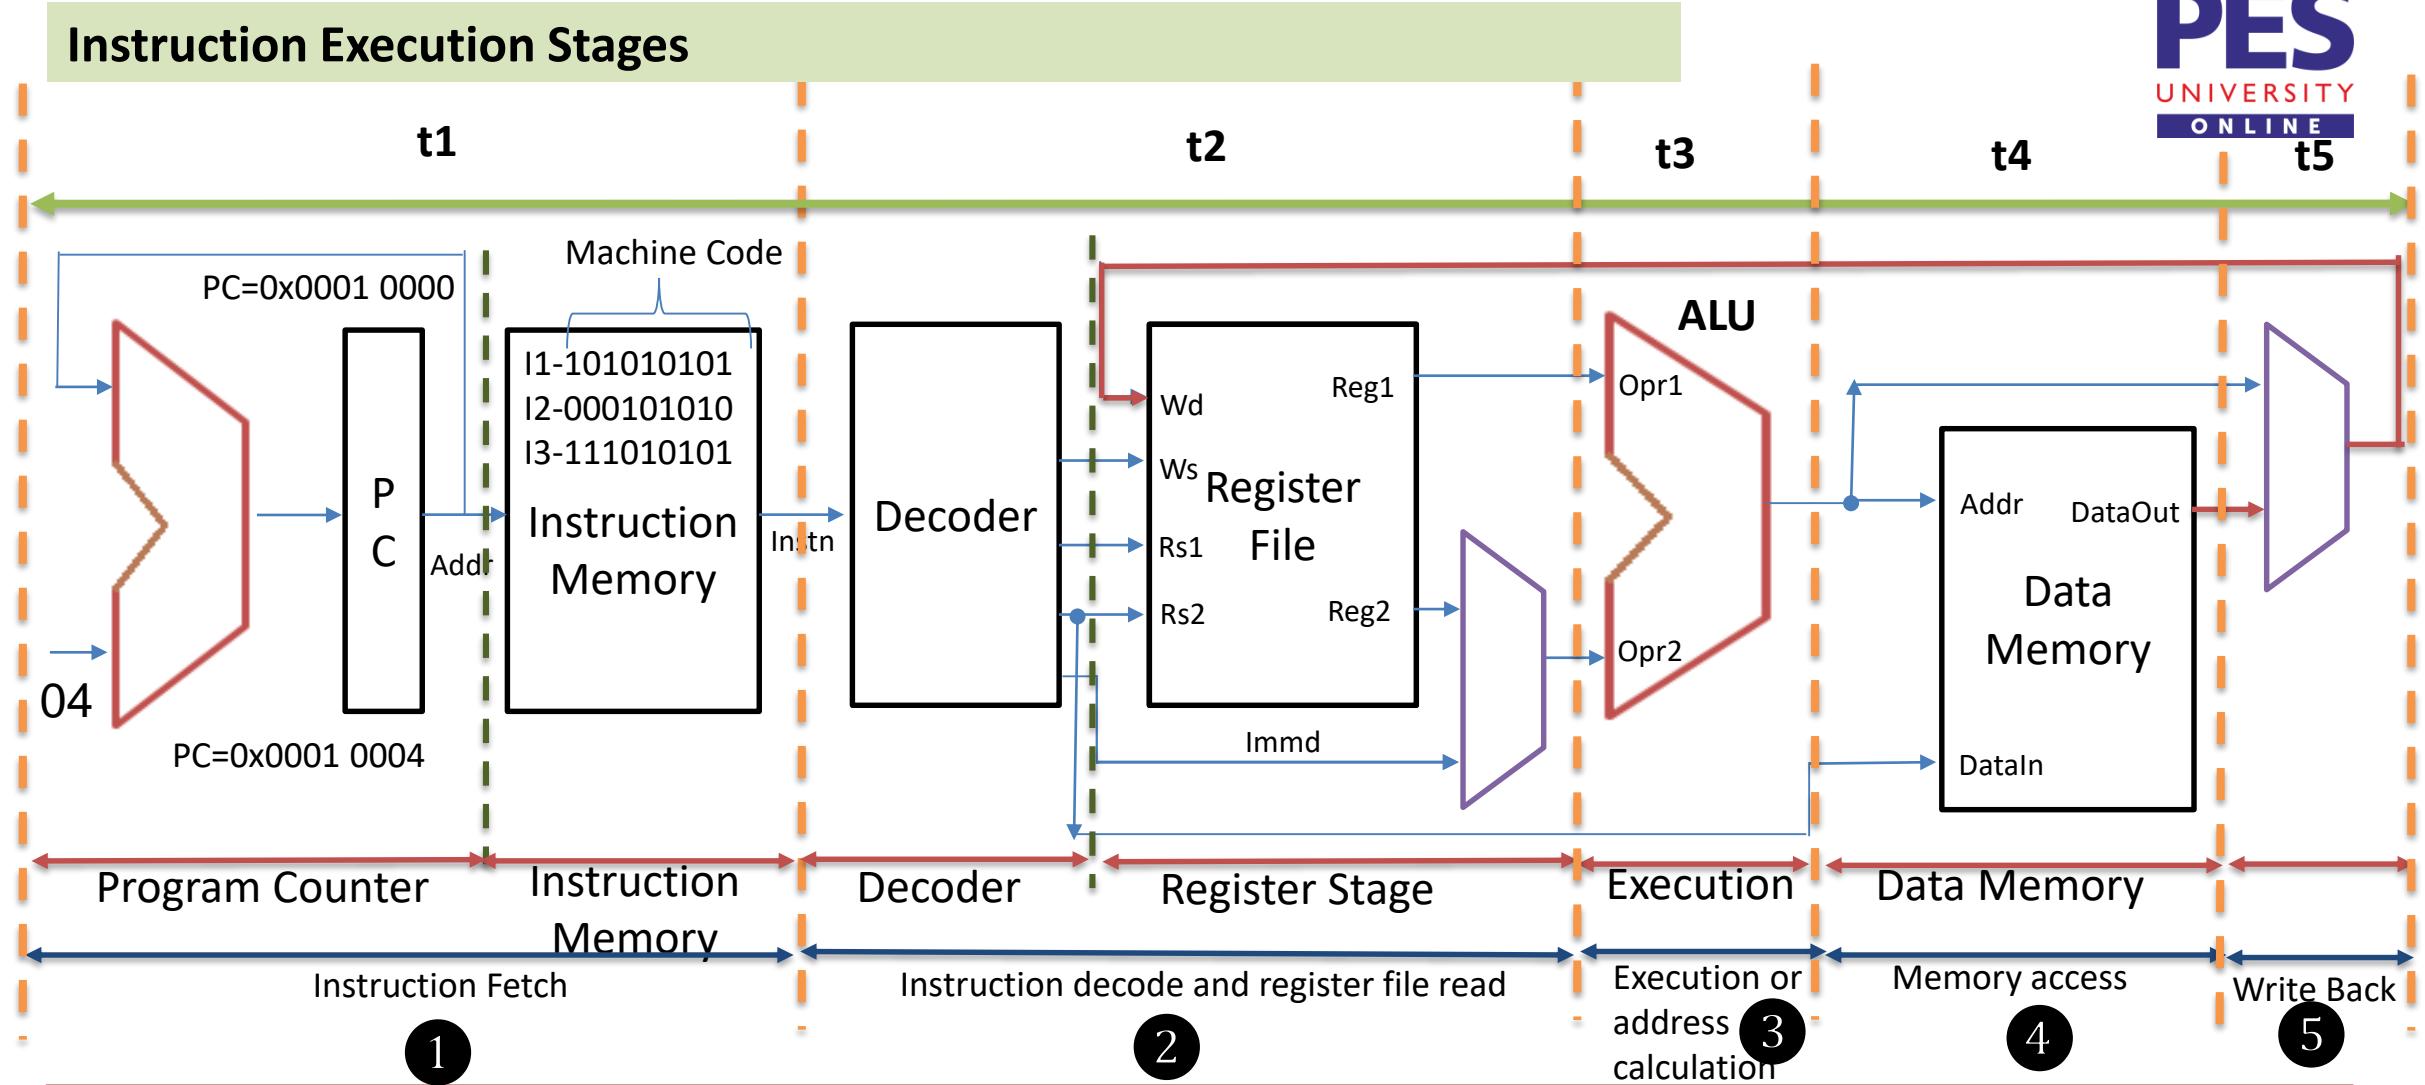


RISC V Instruction Set Architecture (ISA)

Single Cycle Datapath

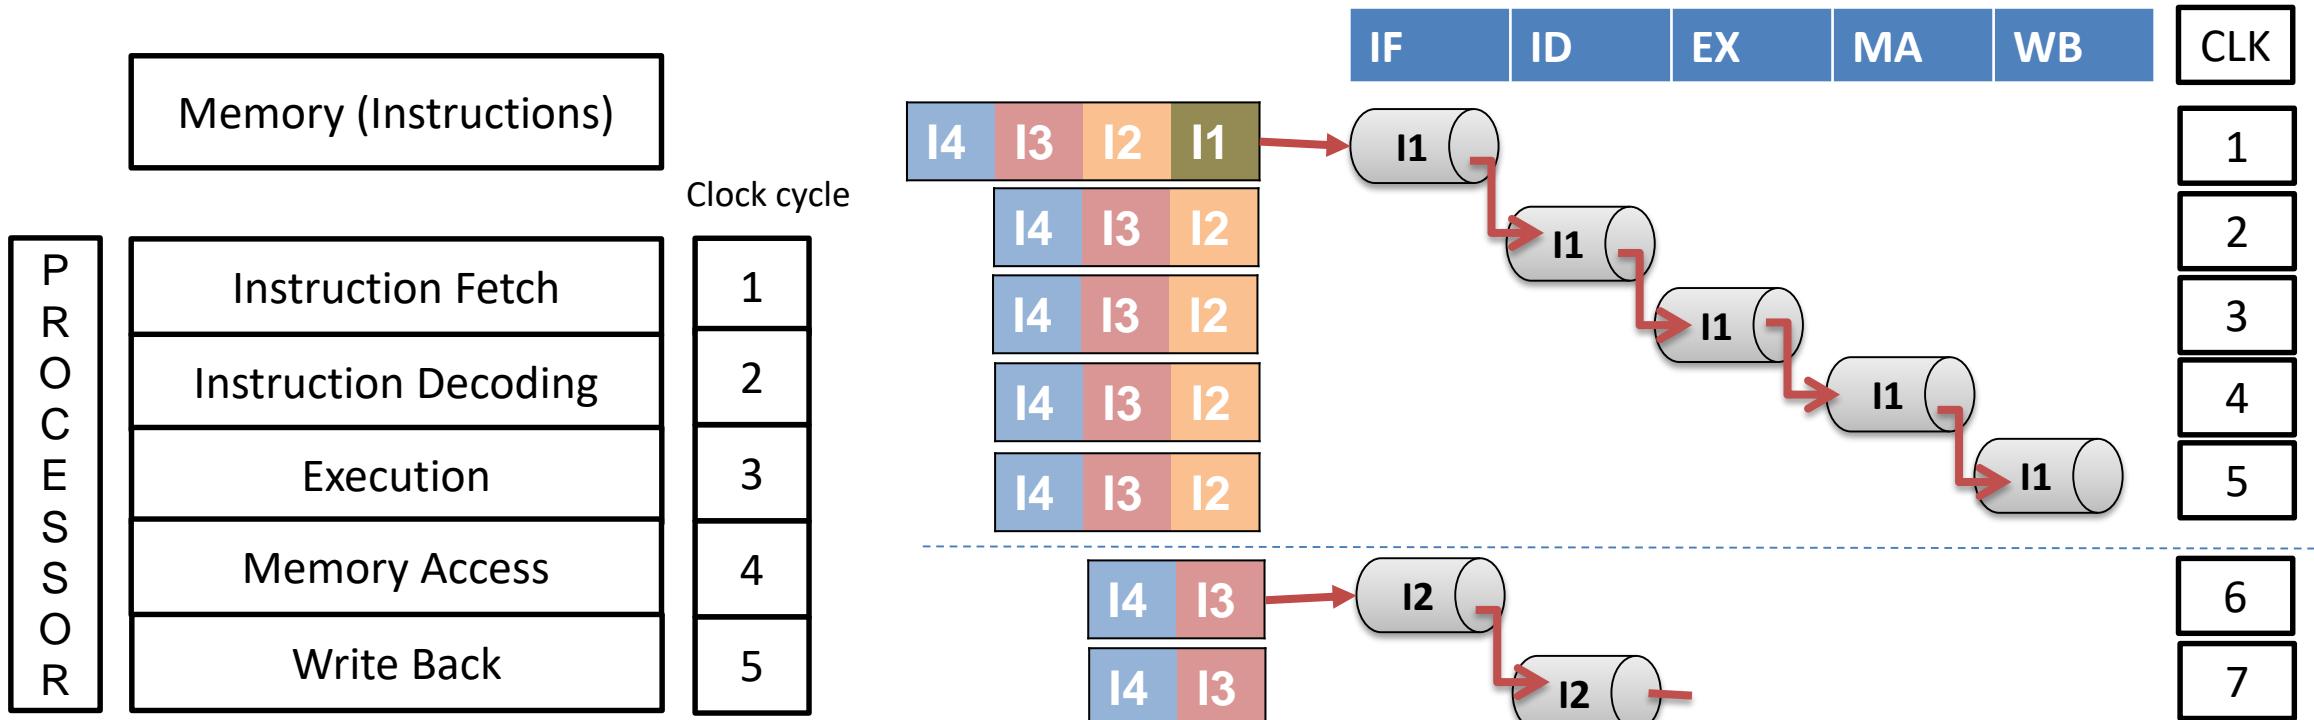


PES
UNIVERSITY
ONLINE



Performance via Pipelining

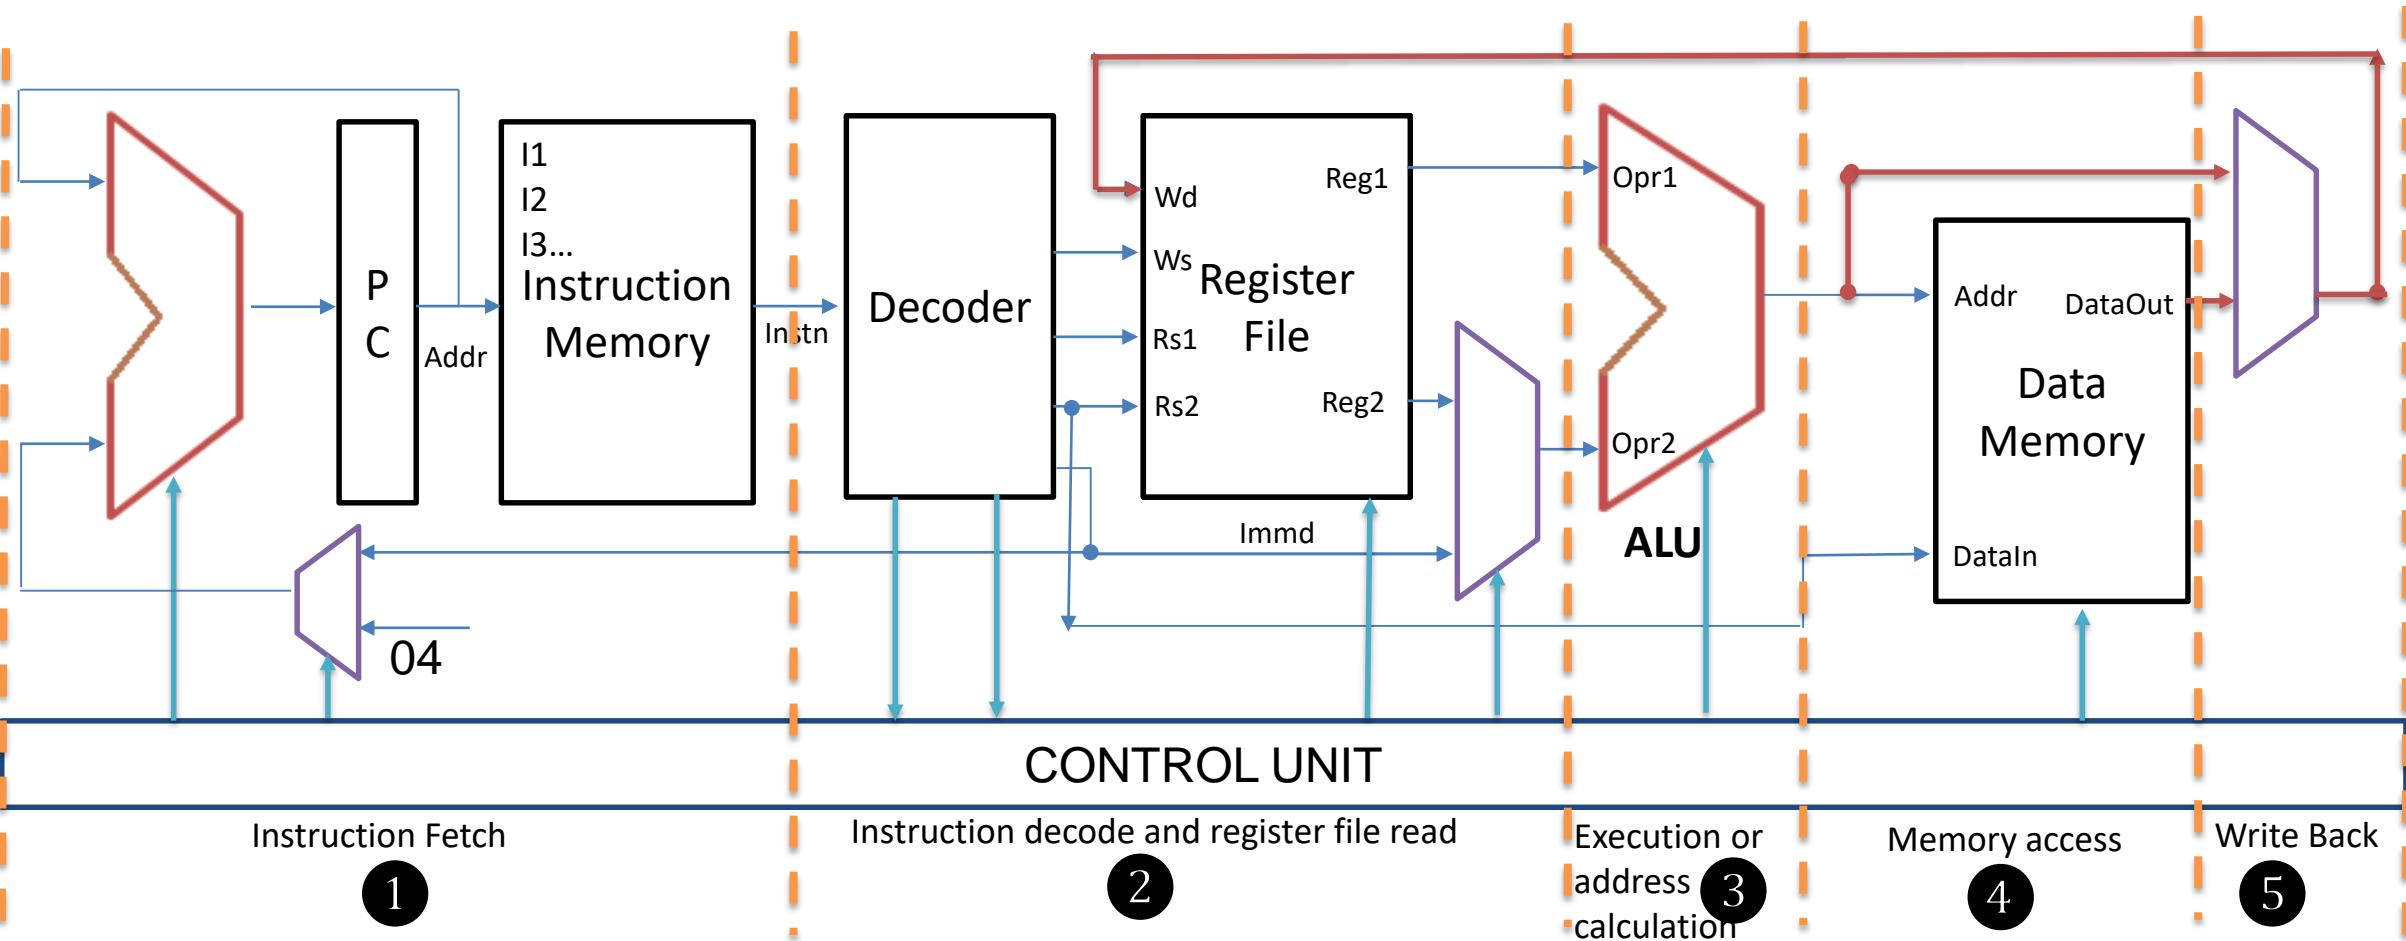
- First let us understand the steps involved in **Non-Pipelined mode** of execution.
- Instruction Execution involves different stages like Fetch , Decode, Execute, Memory access and Write back.



Computer Abstractions and Technology

Data Path

- Datapath - Designed to support data transfers required by instructions.
- Anything that stores data or operates on data within a processor is called data path.



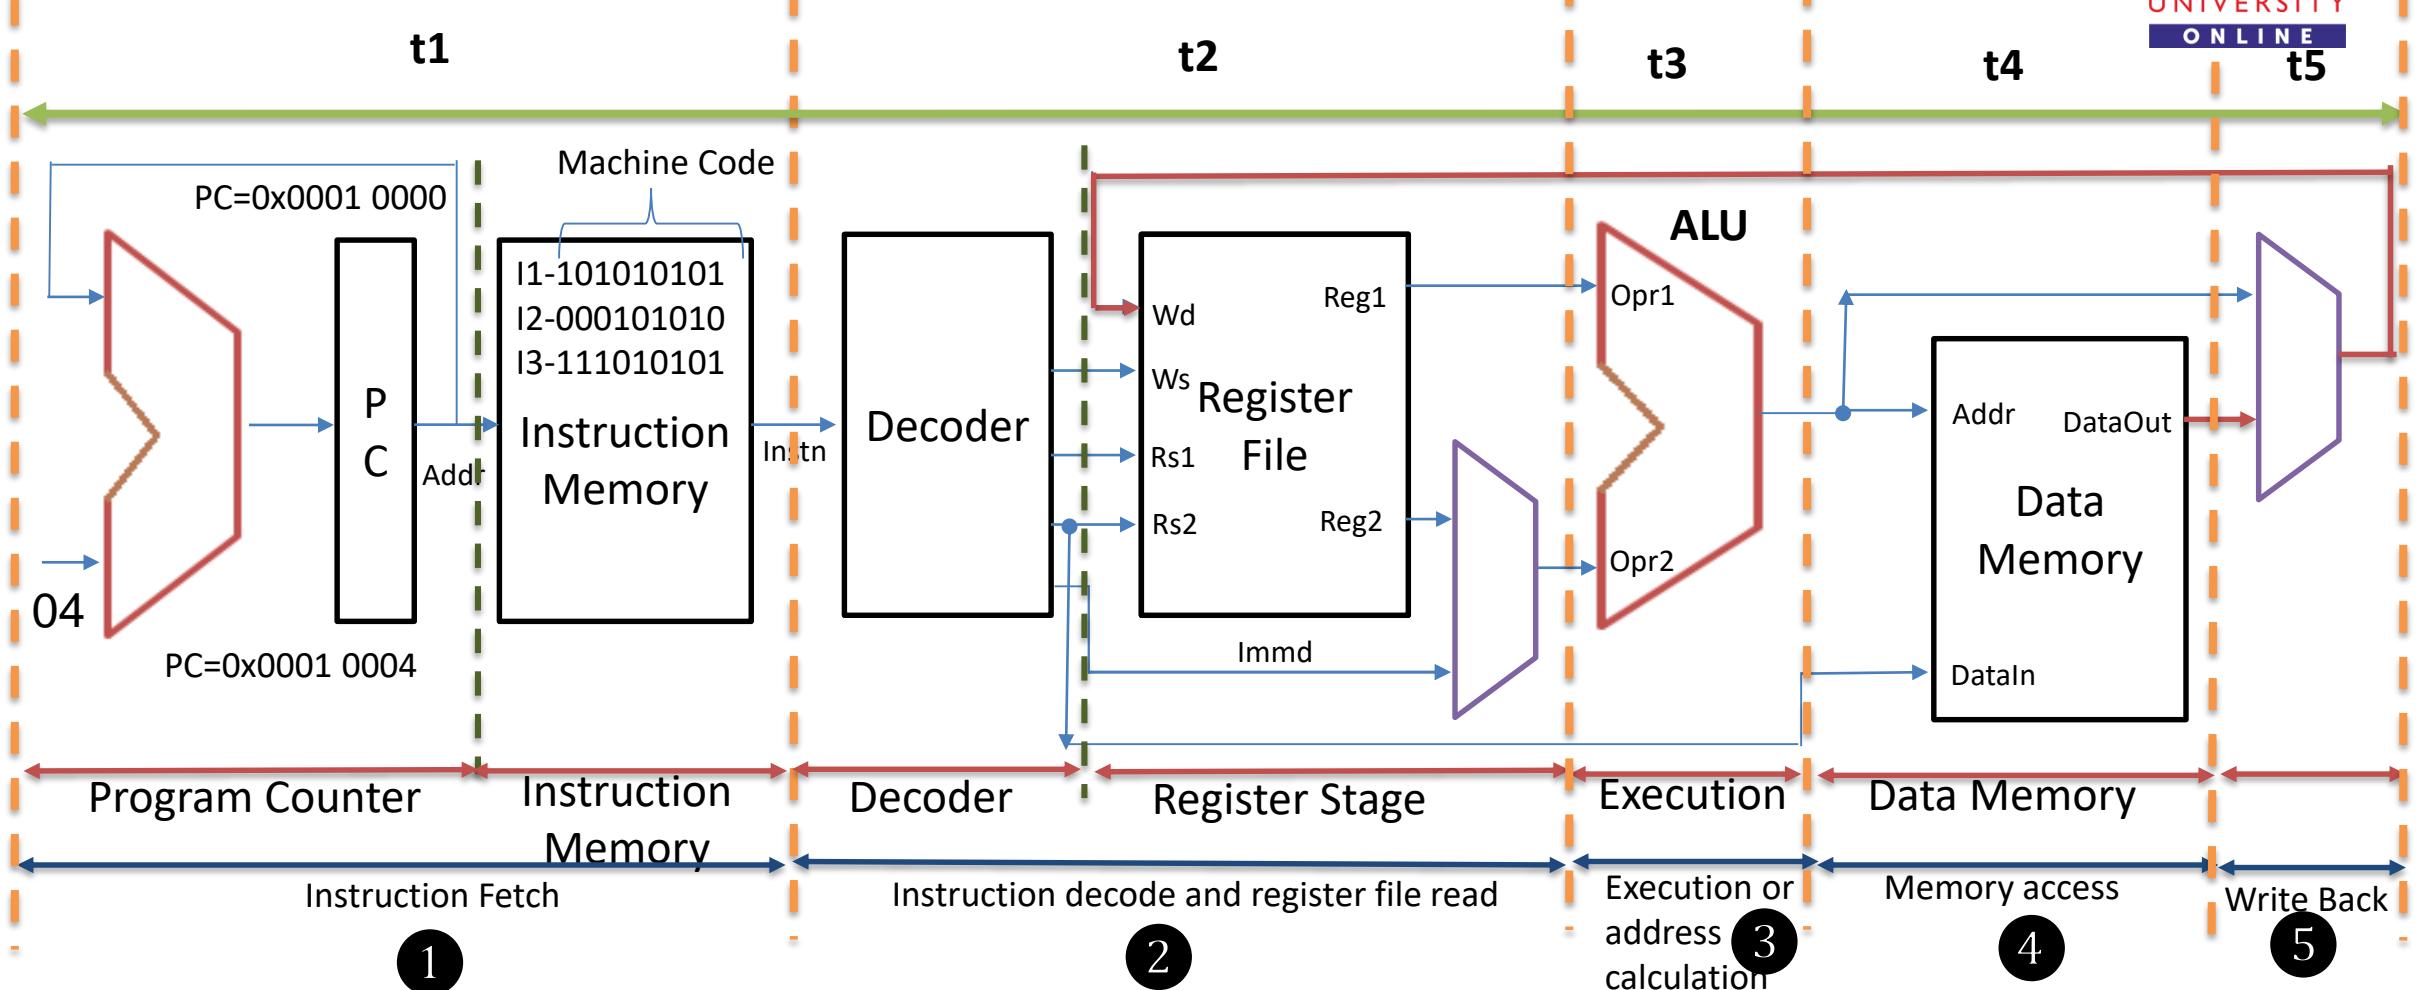
RISC V Instruction Set Architecture (ISA)

Single Cycle Processor



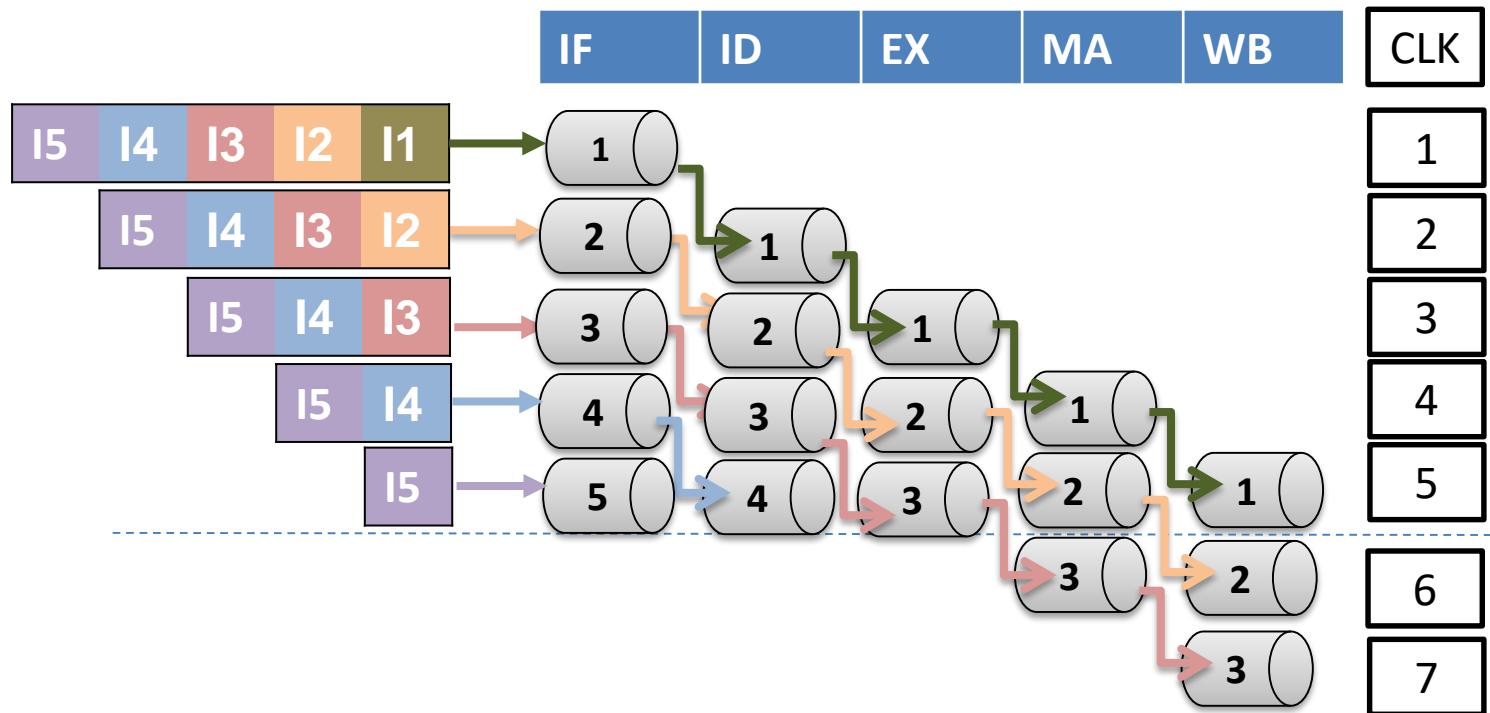
PES
UNIVERSITY
ONLINE

Instruction Execution Stages



Performance via Pipelining

- If these stages are independently performing the task in a sequence, then the **pipelined approach of execution** can be used.
- Processor may have Single, Two, Three, Five or Six stages of pipeline.



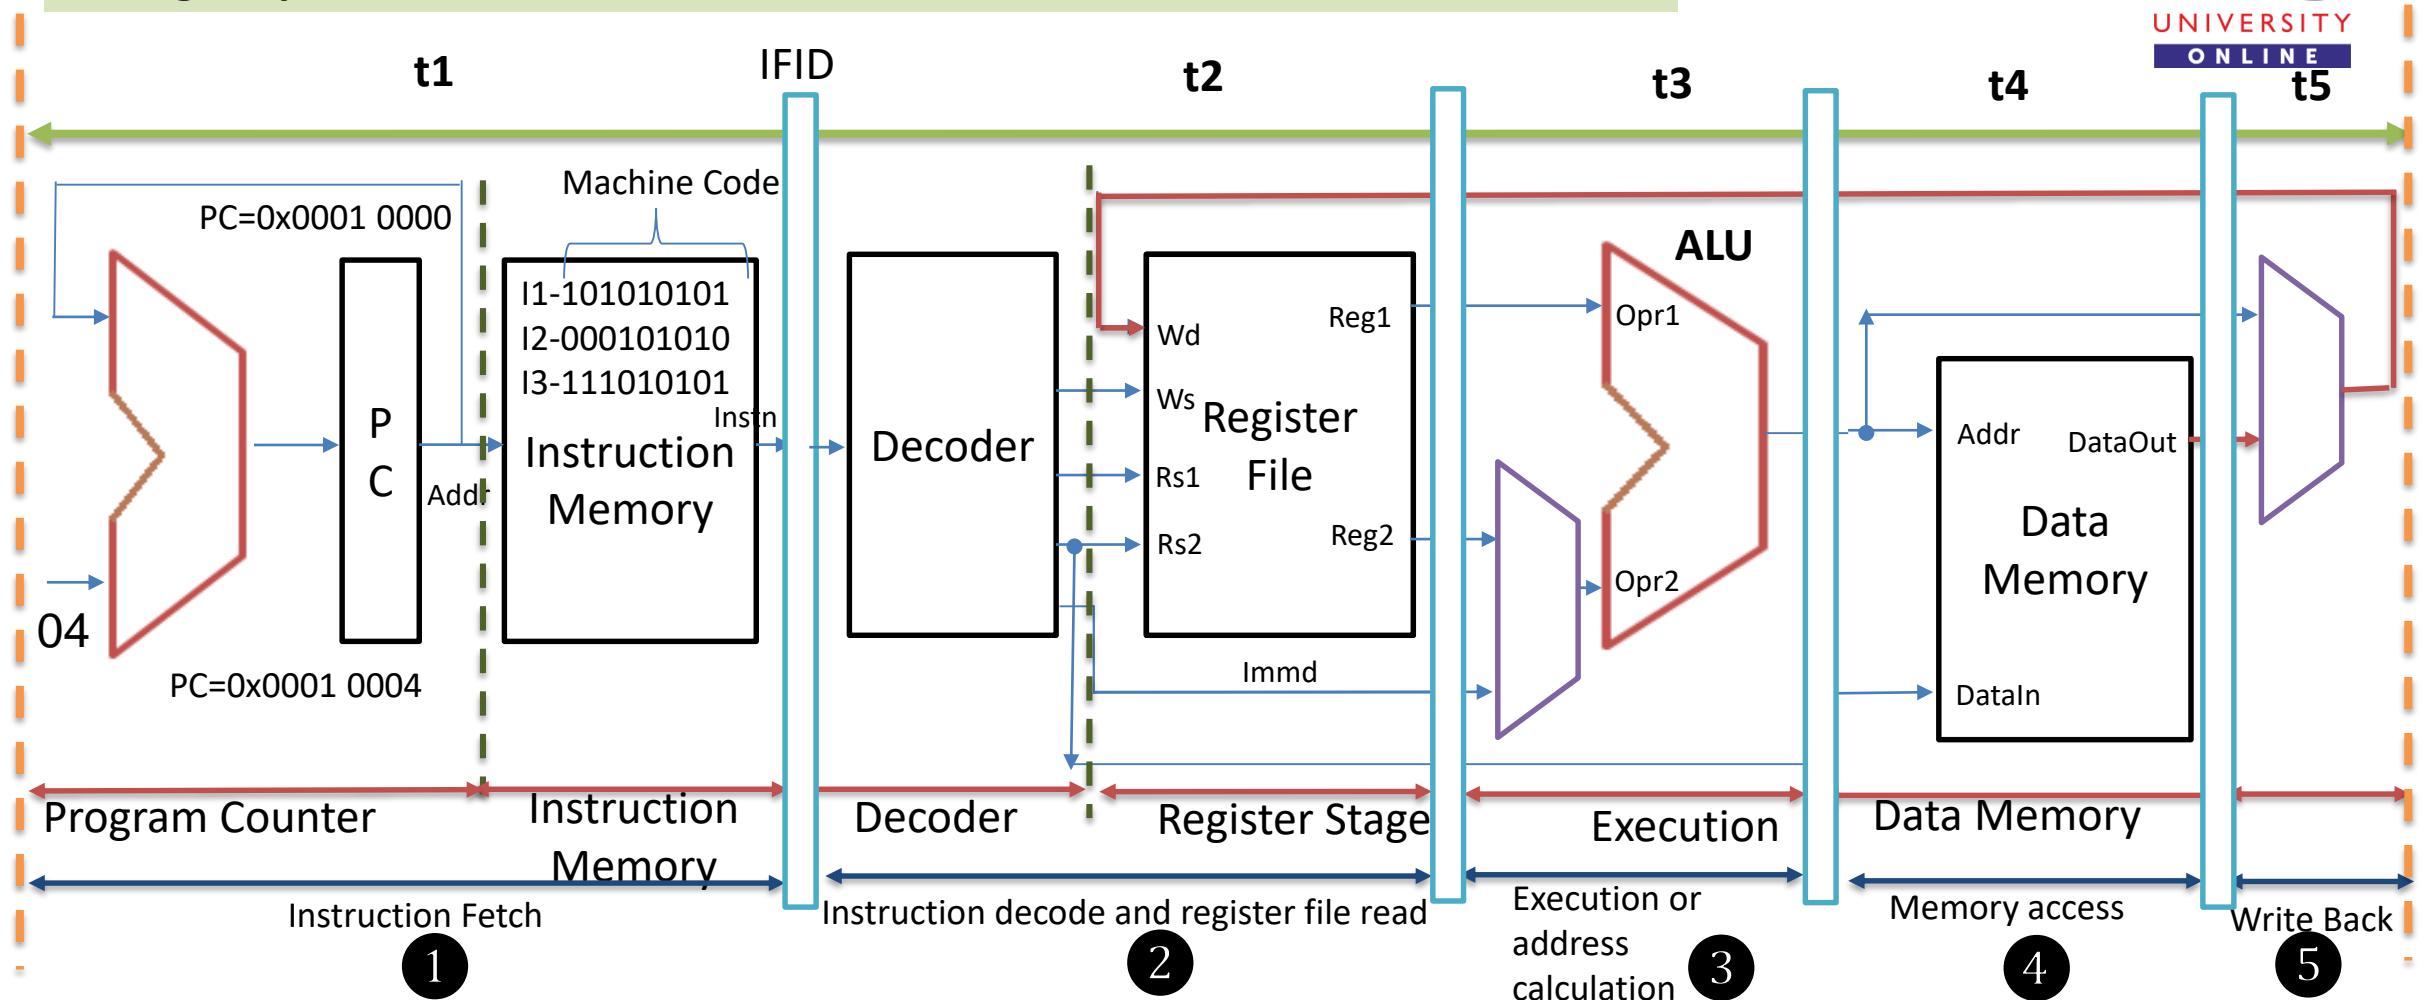
RISC V Instruction Set Architecture (ISA)

Pipelined Processor and Instruction Execution Stages



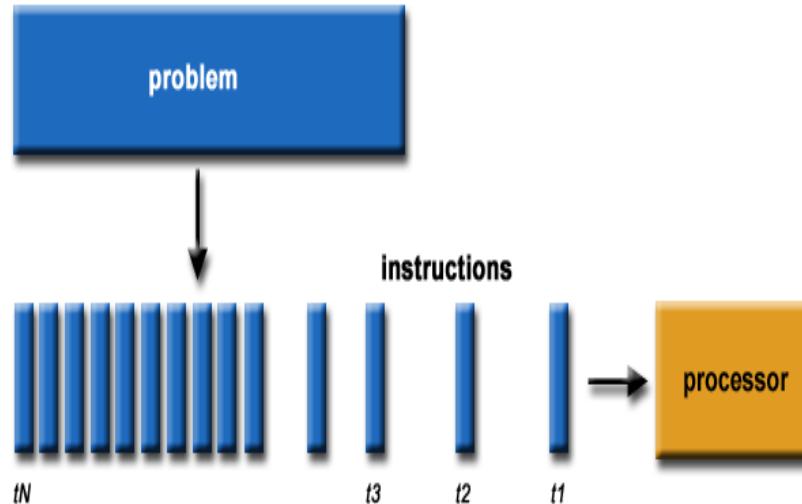
PES
UNIVERSITY
ONLINE

5 Stage Pipelined Processor

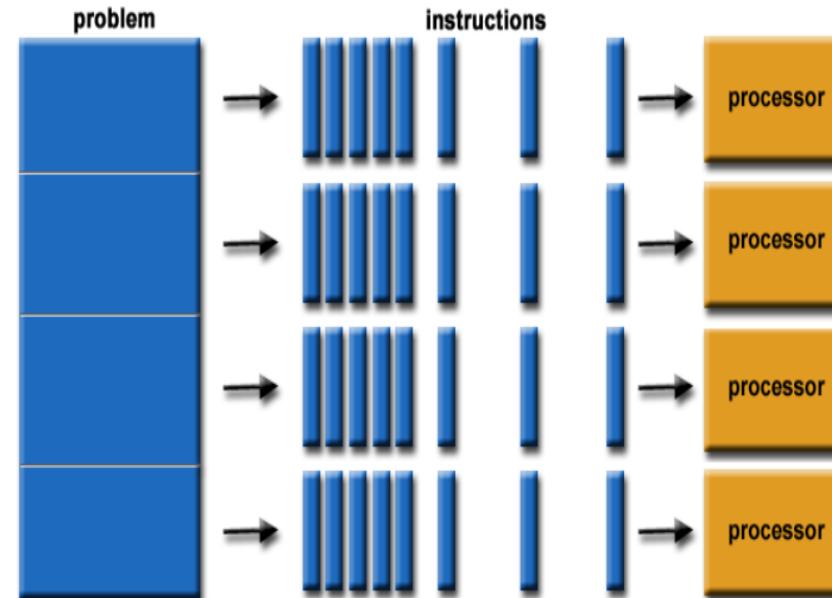


Performance via Parallelism

Sequential computing



Parallel computing



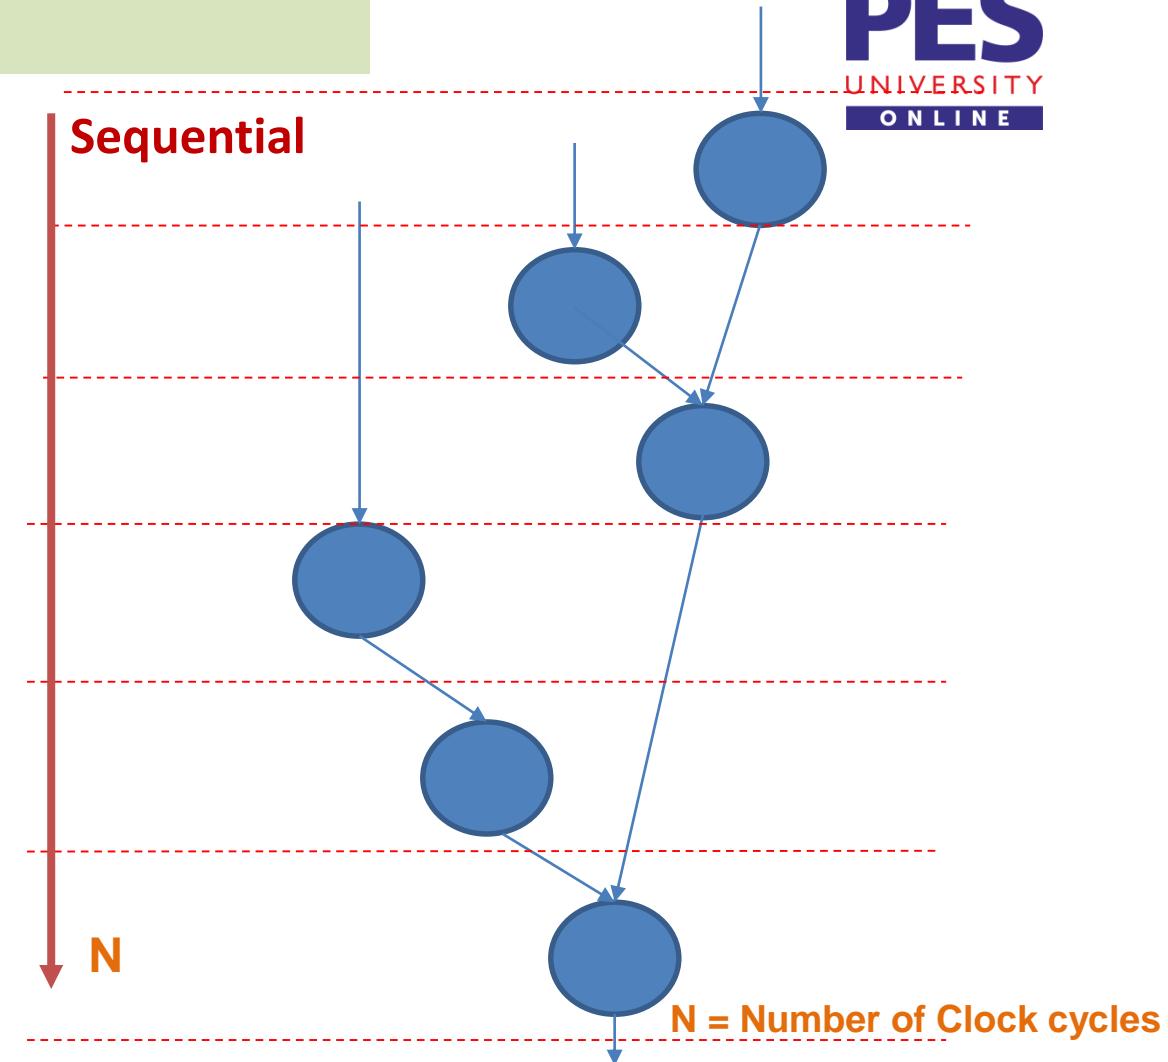
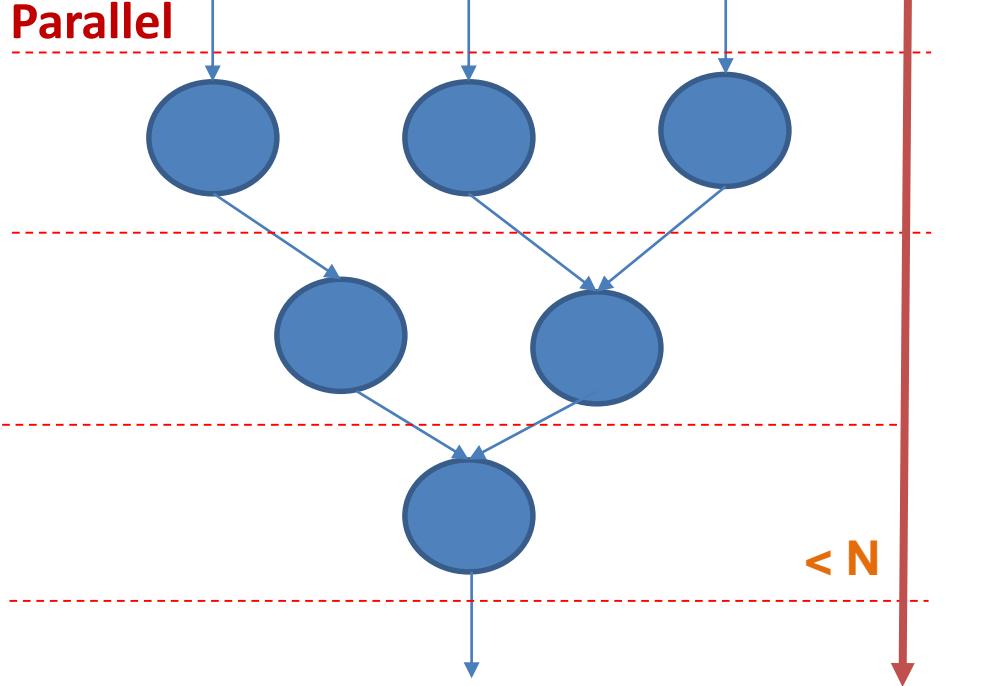
Doing different parts of a task in parallel accomplishes the task in less time than doing them sequentially

Sequential computing is a computational model in which operations are performed in order, one at a time on one processor or computer.

Parallel computing is a computational model where a problem or program is broken into multiple smaller sequential computing operations some of which are performed simultaneously in parallel.

Performance via Parallelism

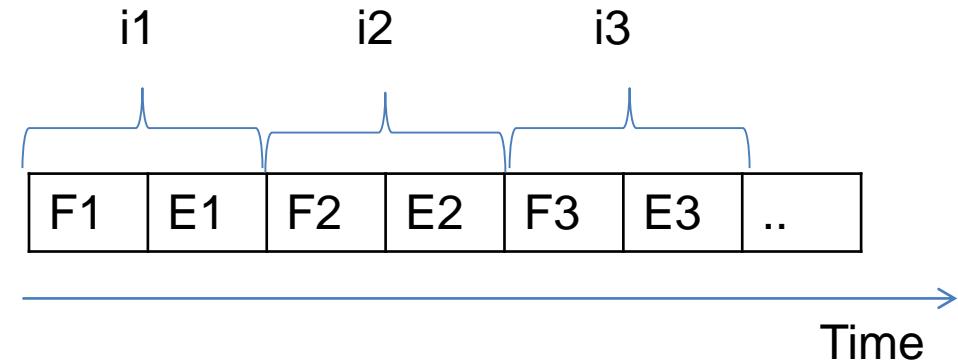
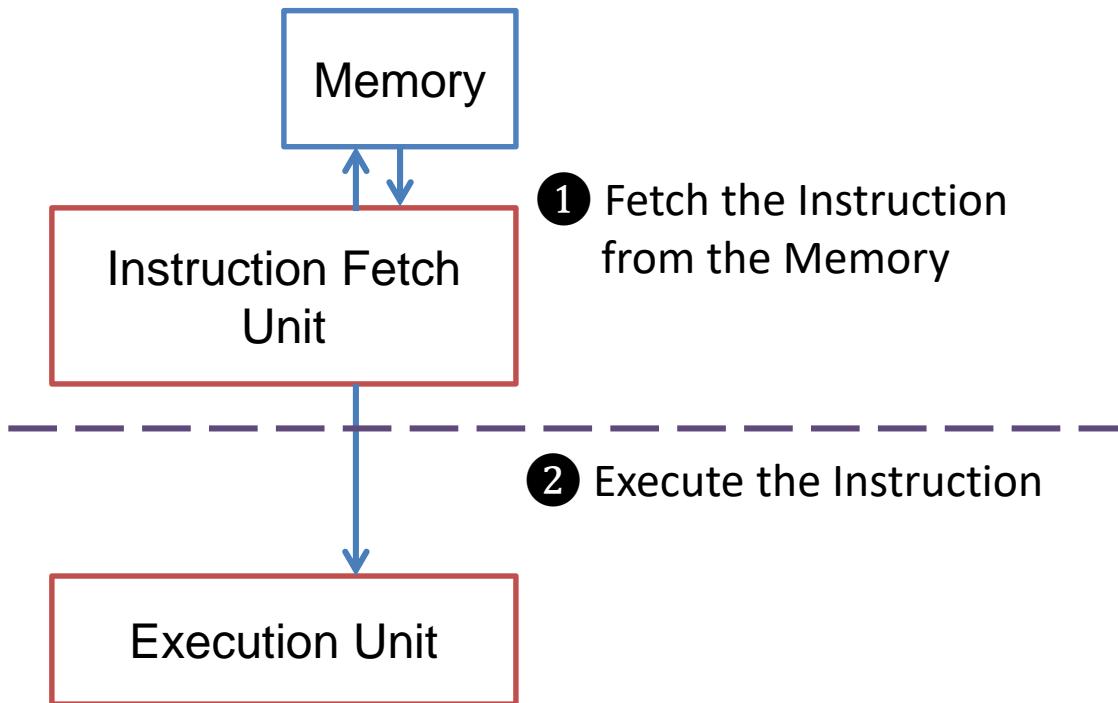
- Computer architects have offered designs that get more performance by **computing operations in parallel**



Performance via Parallelism

- Computer architects have offered designs that get more performance by **computing operations in parallel**

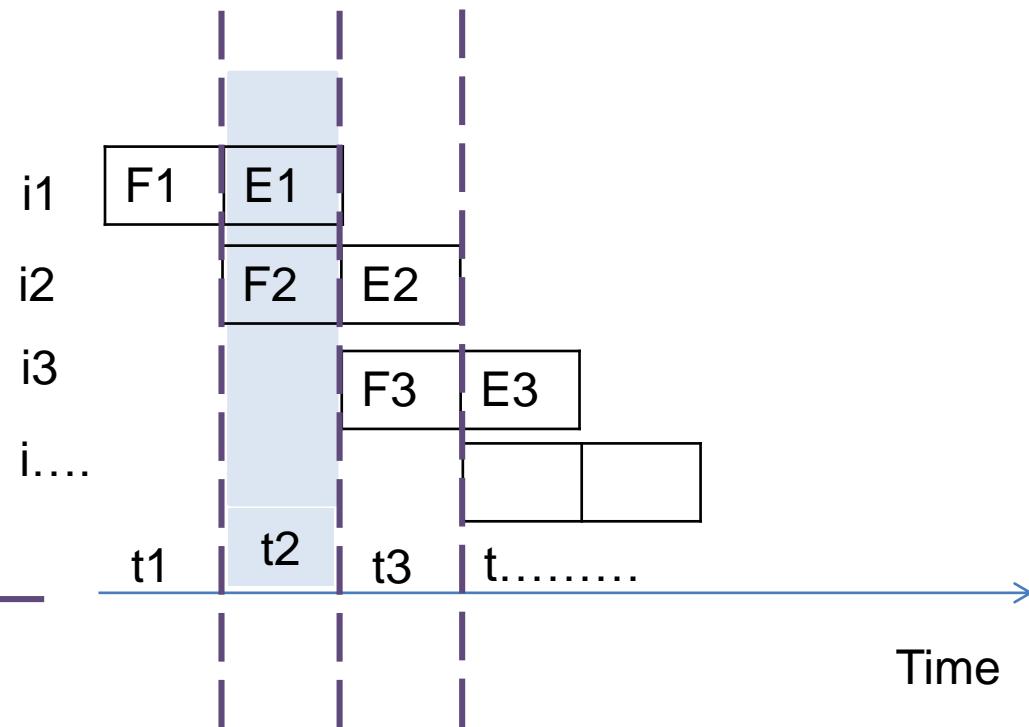
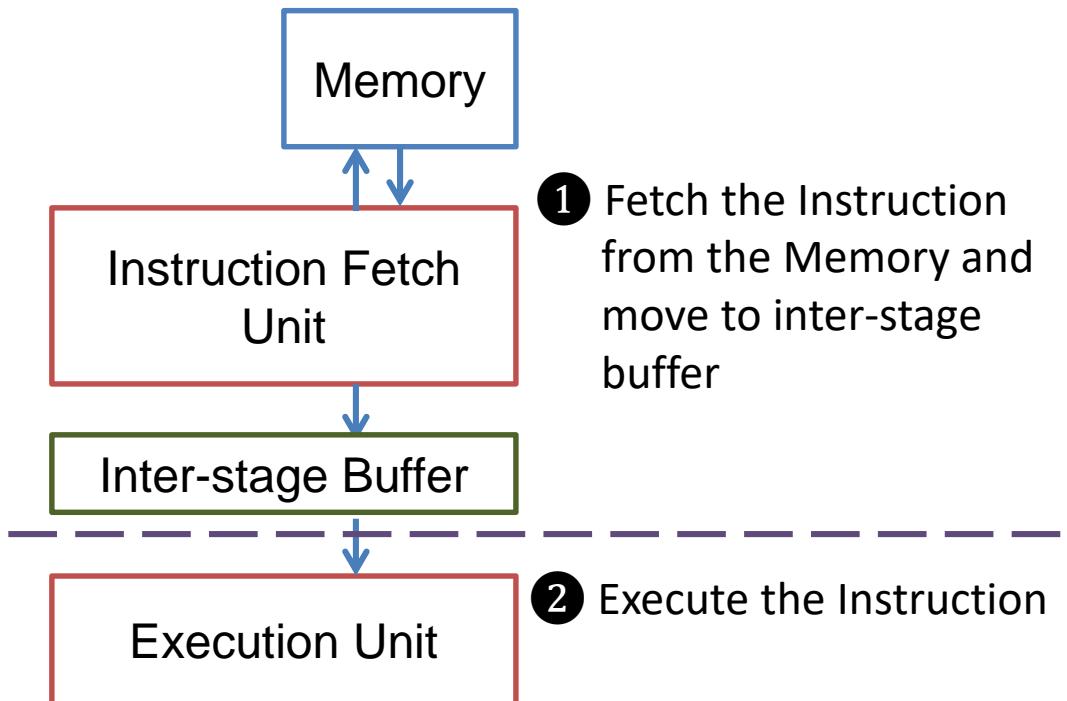
Sequential Operation



Performance via Parallelism

- Computer architects have offered designs that get more performance by computing operations in parallel

Parallelism

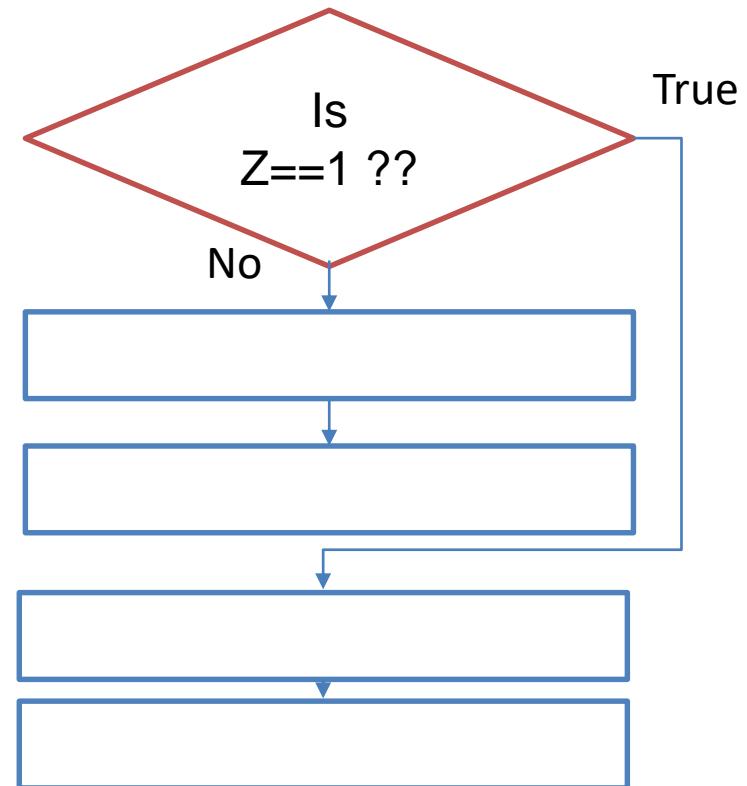


Performance via Prediction

In some cases, it can be faster on average to guess and start working rather than wait until you know for sure, assuming that the mechanism to recover from a mis-prediction is not too expensive and your prediction is relatively accurate.

PC →

Address	Label	Instructions
0x1000		I1
0x1004		I2
0x1008		I3
0x100C		I4
0x1010		BZ Next
0x1014		I5
0x1018		I6....I20
0x101C		I20
0x1020	Next:	I21.....



Execution of branch instruction or branching by the direct modification of PC.

By the time a branch instruction gets into execution stage of pipeline, regardless of branch instruction outcome, It has fetched next successive instruction and are in pipeline. Therefore the RISC-V core will **flush its pipeline**.

Cycle	IF	ID	IE	MA	WB
1	I1	X	X	X	X
2	I2	I1	X	X	X
3	I3	I2	I1	X	X
4	I4	I3	I2	I1	X
5	I5	I4	I3	I2	I1
6					

If branch condition is TRUE
 Branch location gets updated
 in PC in IE stage in RISC-V
 $PC = ALUout$

PC →

Address	Label	Instructions
0x1000		I1
0x1004		I2
0x1008		BZ Next
0x100C		I4
0x1010		I5
0x1014		I6
0x1018		I7
0x101C		I8
0x1020	Next:	I9.....

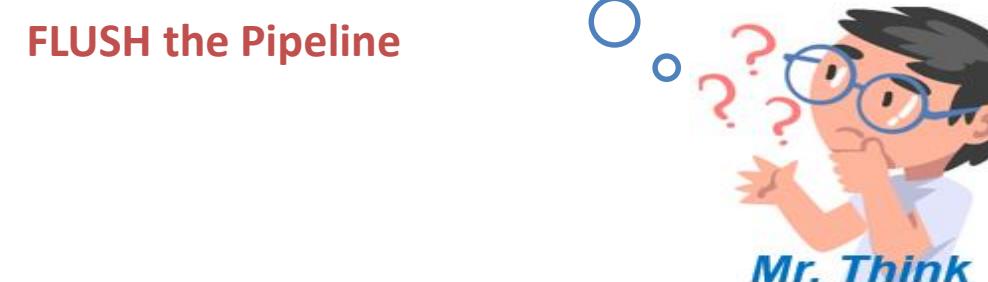
Execution of branch instruction or branching by the direct modification of PC.

By the time a branch instruction gets into execution stage of pipeline, regardless of branch instruction outcome, It has fetched next successive instruction and are in pipeline. Therefore the RISC-V core will **flush its pipeline**.

Cycle	IF	ID	IE	MA	WB
1	I1	X	X	X	X
2	I2	I1	X	X	X
3	I3	I2	I1	X	X
4	I4	I3	I2	I1	X
5	I5	I4	I3	I2	I1
6					

What will happen to I4 and I5 which are in Pipeline ????

FLUSH the Pipeline



Mr. Think

Execution of branch instruction or branching by the direct modification of PC.

By the time a branch instruction gets into execution stage of pipeline, regardless of branch instruction outcome, It has fetched next successive instruction and are in pipeline. Therefore the RISC-V core will **flush its pipeline**.

Cycle	IF	ID	IE	MA	WB
1	I1	X	X	X	X
2	I2	I1	X	X	X
3	I3	I2	I1	X	X
4	I4	I3	I2	I1	X
5			I3	I2	I1
6	I9				I2
7	I10	I9			
8	I11	I10	I9		

I4 and I5 which were in pipeline will be FLUSHED & Refilling of pipeline starts from branched location

PC →

Address	Label	Instructions
0x1000		I1
0x1004		I2
0x1008		BZ Next
0x100C		I4
0x1010		I5
0x1014		I6
0x1018		I7
0x101C		I8
0x1020	Next:	I9.....



Instruction Memory

Address	Label	Instructions
0x1000		I1
0x1004		I2
0x1008		BZ Next
0x100C		I4
0x1010		I5
0x1014		I6
0x1018		I7
0x101C		I8
0x1020	Next:	I9.....

Without Branch Prediction

Cycle	IF	ID	IE	MA	WB
1	I1	X	X	X	X
2	I2	I1	X	X	X
3	I3	I2	I1	X	X
4	I4	I3	I2	I1	X
5			I3	I2	I1
6	I9				I2
7	I10	I9			
8	I11	I10	I9		
9	I12	I11	I10	I9	
10	I13	I12	I11	I10	I9

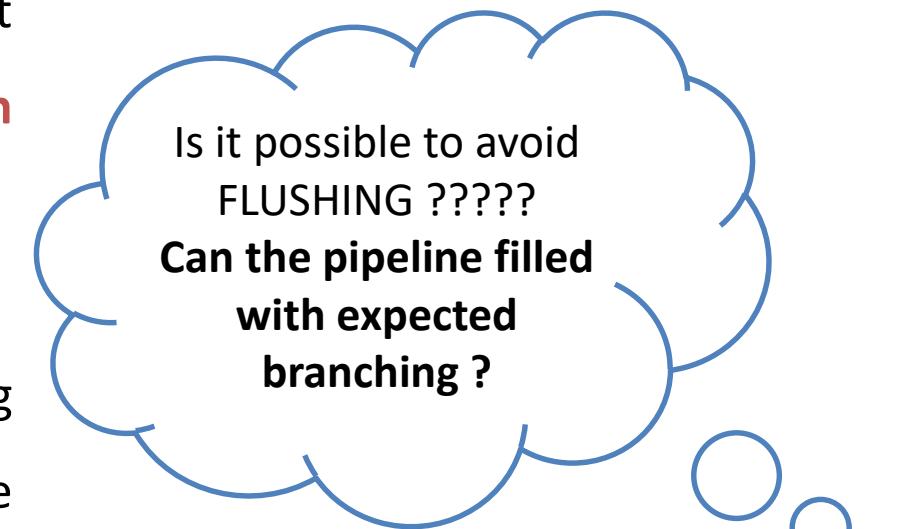
With Branch Prediction

Cycle	IF	ID	IE	MA	WB
1	I1	X	X	X	X
2	I2	I1	X	X	X
3	I3	I2	I1	X	X
4	I9	I3	I2	I1	X
5	I10	I9	I3	I2	I1
6	I11	I10	I9	I3	I2
7	I12	I11	I10	I9	I3
8	I13	I12	I11	I10	I9

Branch Prediction

In computer architecture, a **branch predictor** is a digital circuit that tries to **guess which way a branch will go before this is known definitively**. This reduces effect of pipeline flush.

Branch prediction reduces the effect of a pipeline flush by predicting possible branches and loads the new branch address prior to the execution of the instruction

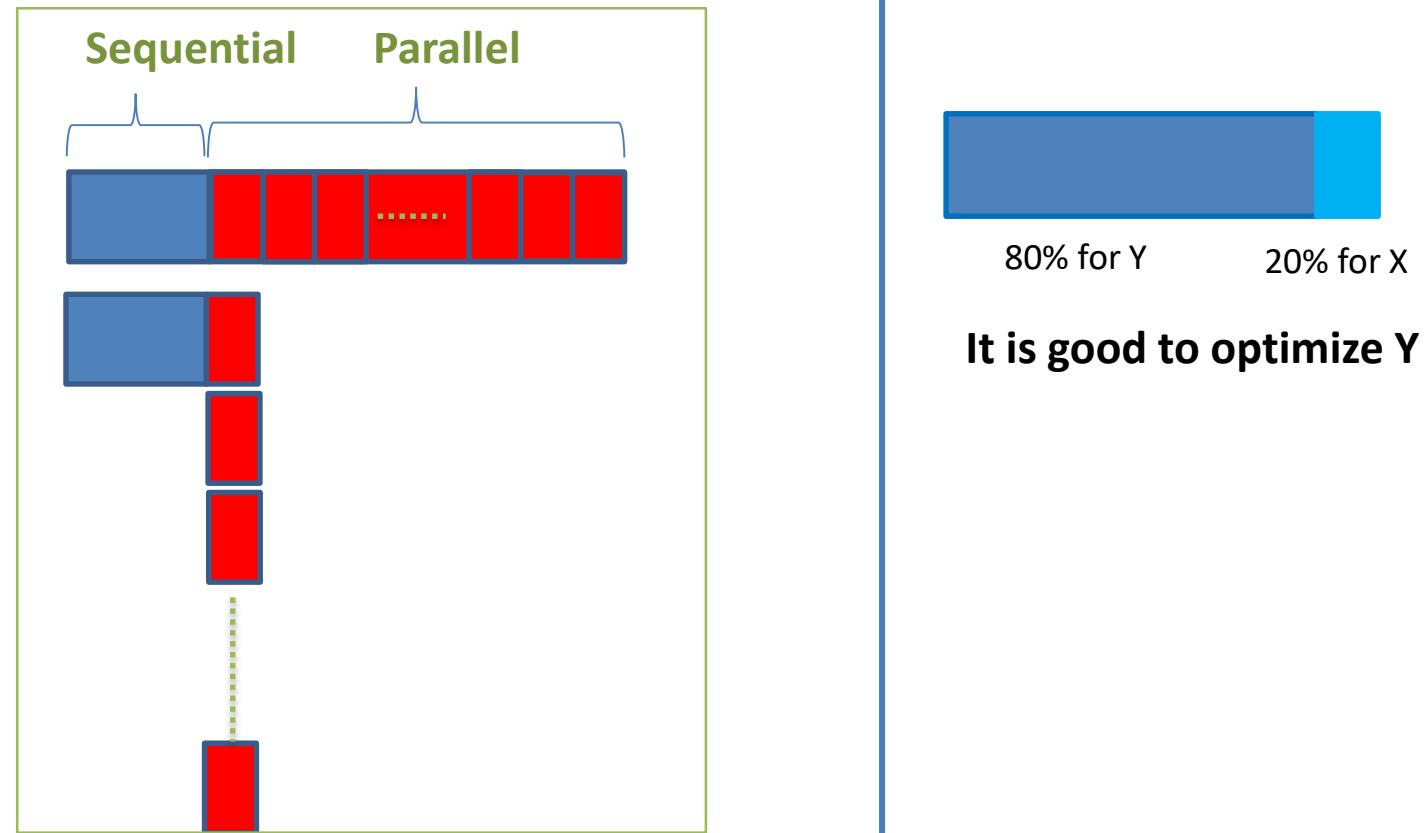


Make the Common Case Fast

Making the **common case fast** will tend to **enhance performance better than optimizing the rare case**. Ironically, the common case is often simpler than the rare case and hence is usually easier to enhance..

Amdahl's Law:

- ✓ A program needs 20 hours to complete
- ✓ No matter what, 1 hour needs to run sequentially
- ✓ Only, rest of the 19 hours can run in parallel

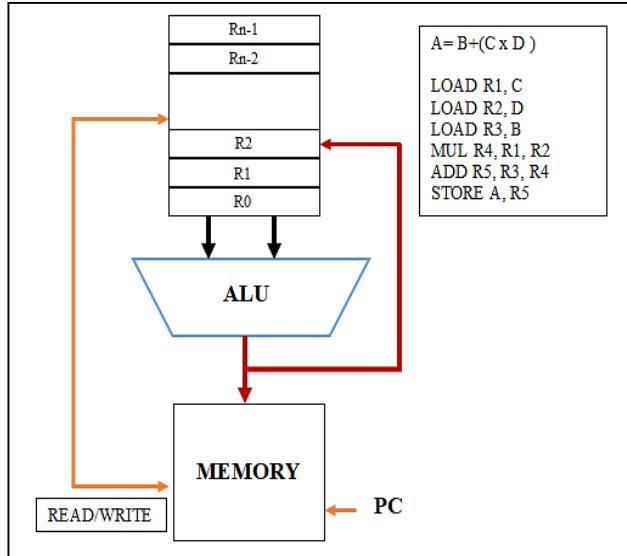


Make the Common Case Fast

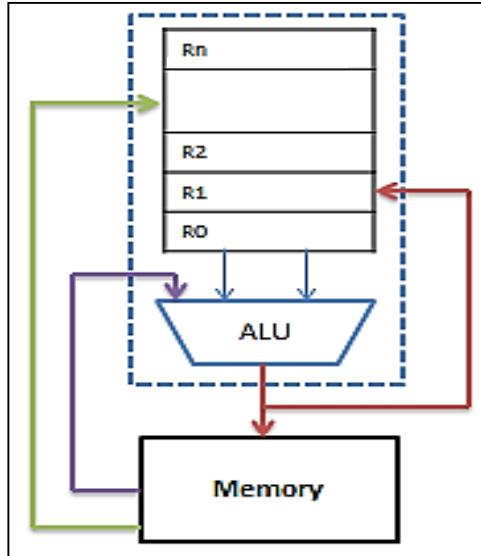
The load-store computer architecture separates instructions into

1. Memory access operations (load and stores)
2. Operations that operate on the data in the register file (register to register or register to immediate).

Load-Store



Direct Memory Access



By not combining memory accesses with data manipulation operations, the processor's complexity is reduced which enables making the common case fast

Dependability via Redundancy

Replicate, If possible

- Computers not only need to be fast; they **need to be dependable**.
- Since any physical device can fail, we make systems **dependable** by including redundant components that can take over when a failure occurs *and* to help detect failures.
- Failing piece does not make whole system fail
- Increasing transistor density reduce the cost of redundancy

Summary

Use Abstraction to Simplify Design

Hierarchy of Memories

Performance via Pipelining

Performance via Parallelism

Performance via Prediction

Make the Common Case Fast

Dependability via Redundancy



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 4 – Technologies for Building Processors and Memory

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Technologies used over time

Year	Technology used in computers	Relative performance/unit cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit	900
1995	Very large-scale integrated circuit	2,400,000
2020	Ultra large-scale integrated circuit	500,000,000,000

FIGURE 1.10 Relative performance per unit cost of technologies used in computers over time. Source: Computer Museum, Boston, with 2013 extrapolated by the authors. See  [Section 1.13](#).

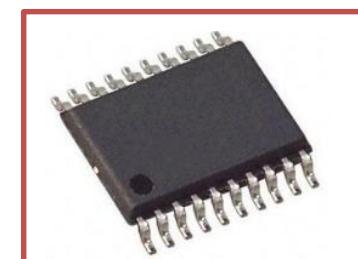
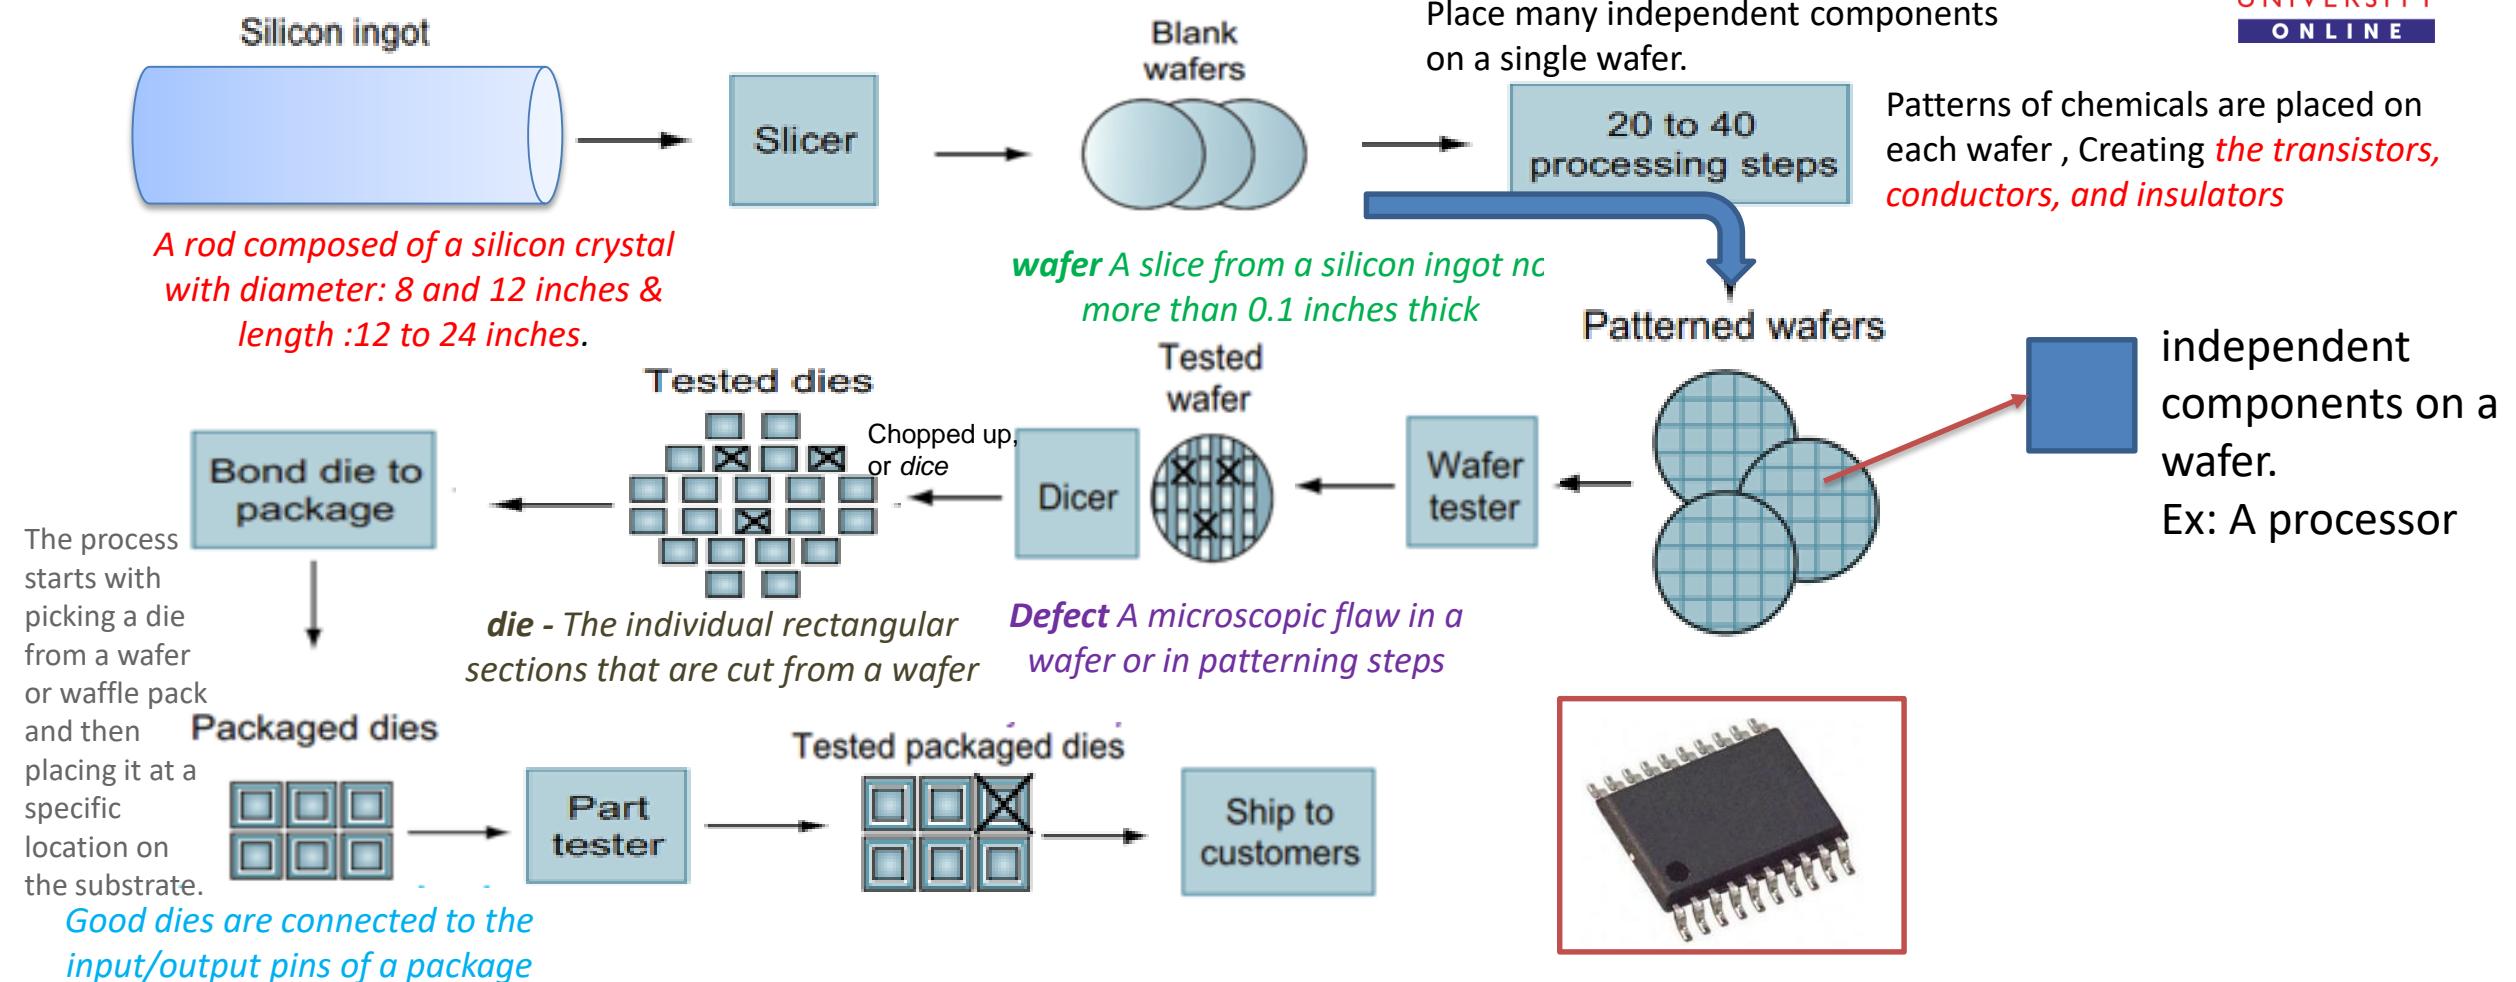
- A **transistor** is simply an on/off switch controlled by electricity.
- The *integrated circuit* (IC) combined dozens to hundreds of transistors into a single chip.
- When Gordon Moore predicted the continuous doubling of resources, he was forecasting the growth rate of the number of transistors per chip.

The Chip manufacturing process

- The manufacture of a chip begins with **silicon**, a substance found in sand. Because silicon does not conduct electricity well, it is called a **semiconductor**.
- With a special chemical process, it is possible to add materials to silicon that allow tiny areas to transform into one of three devices:
 1. **Excellent conductors of electricity** (using either microscopic copper or aluminum wire)
 2. **Excellent insulators from electricity** (like plastic sheathing or glass)
 3. **Areas that can conduct or insulate under specific conditions** (as a switch)
- Transistors fall into the last category.
- **A VLSI circuit**, then, is just **billions of combinations of conductors, insulators, and switches** manufactured in a single small package.

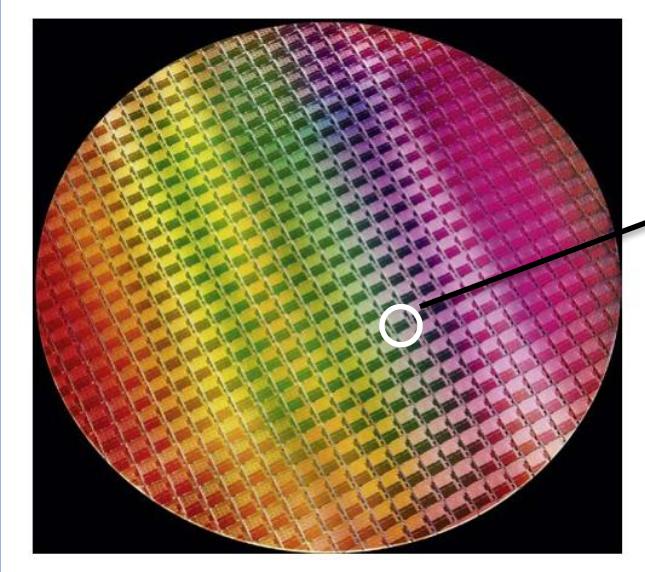


The chip manufacturing process



The chip manufacturing process

A photograph of a wafer containing microprocessors before they have been diced



A 12-inch (300mm) wafer this 10nm wafer contains **10th Gen Intel® Core™ processors**, code-named “Ice Lake” (Courtesy Intel).

An individual microprocessor die.



The processor integrated circuit inside the A12 package. The size of chip is 8.4 by 9.91 mm, and it was manufactured originally in a 7-nm process.

It has

- two identical ARM processors
- four small cores,
- a Graphics Processing Unit (GPU)
- a domain-specific accelerator for neural networks called the NPU
- Second-level cache memory (L2) banks for the big and small cores
- Interfaces to the main memory (DDR DRAM).

(CourtesyTechInsights, www.techinsights.com)

Cost and Prize of the Chip

Yield : The percentage of good dies from the total number of dies on the wafer.

The cost of an integrated circuit can be **expressed in three simple equations:**

$$\text{Dies per wafer} = \frac{\text{Wafer area}}{\text{Die area}} \quad \dots \dots \dots \quad (1)$$

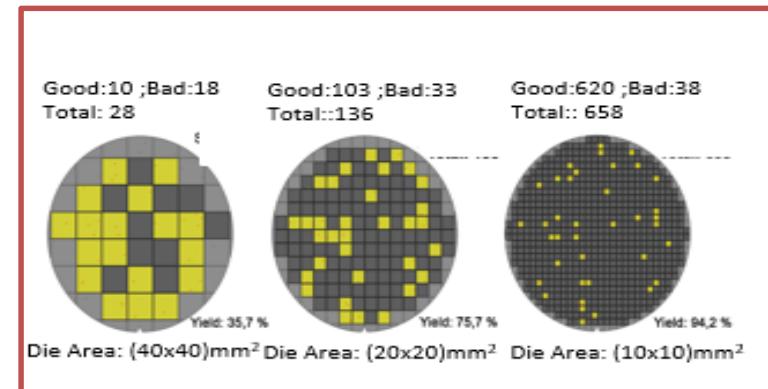
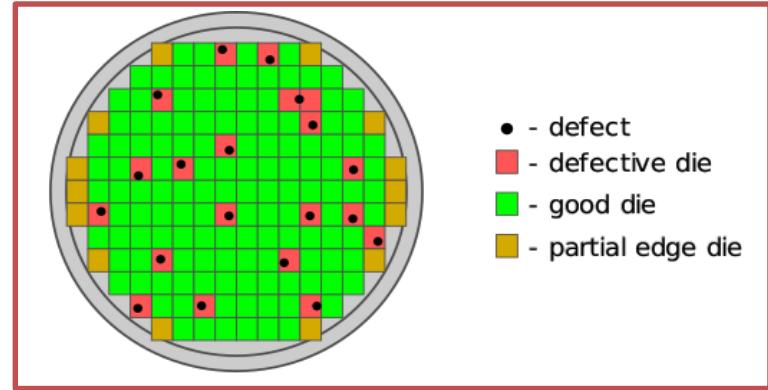
Equation (1) is an approximation, **since it does not subtract the area near the border of the round wafer** that cannot accommodate the rectangular dies

$$\text{Cost per die} = \frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{yield}} \quad \dots \dots \dots \quad (2)$$

Equation (3) **Empirical observations** of Die-yield (Assuming wafer yield is 100%) at IC factories, with the exponent related to the number of critical processing steps.

$$\text{Yield} = \frac{1}{(1 + (\text{Defects per area} \times \text{Die area}))^N} \quad \dots \dots \dots \quad (3)$$

N is the process-complexity factor.



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Cost and Prize of the Chip

Yield : The percentage of good dies from the total number of dies on the wafer.

N in eqn(3) is the process-complexity factor & depends on technology used.

For 130nm processes , $N=4$

For 28 nm processes in 2017, N is 7.5–9.5.

For 16 nm process, N ranges from 10 to 14.

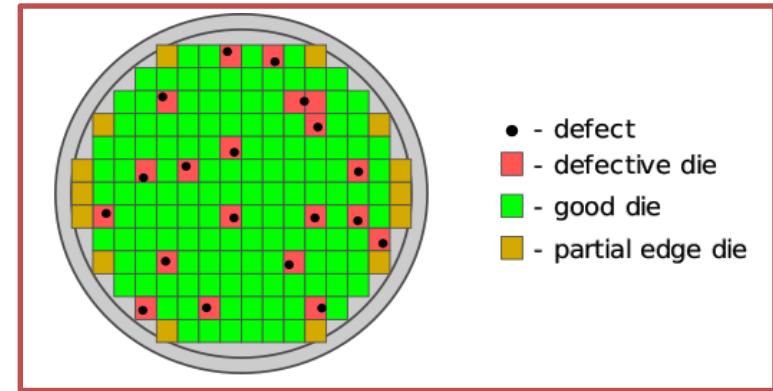
For 40nm process, N ranges 11.5 to 15.5

Note: 1. Factor N should be known to calculate the yield. In case factor N is not given use the following empirical formula i.e., eqn(4).

$$\text{Yield} = \frac{1}{\left(1 + \text{defects per area} \times \frac{\text{die area}}{2}\right)^N} \quad (4)$$

If wafer yield is not 100%, then use the following formula

$$\text{Die yield} = \text{Wafer yield} \times \frac{1}{\left(1 + \text{Defects per unit area} \times \text{Die area}\right)^N} \quad (5)$$



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Cost and Prize of the Chip

Defects per Area can be derived from eqn (4) or eqn(5). From eqn (4)

$$\text{Yield} \left(1 + \text{defects per area} \times \frac{\text{diearea}}{2} \right)^2 = 1$$

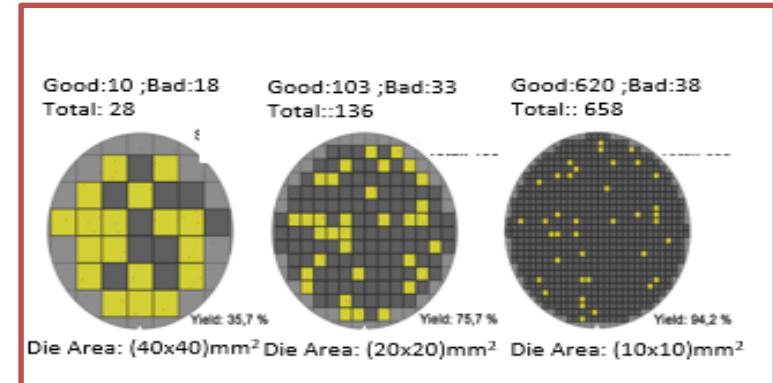
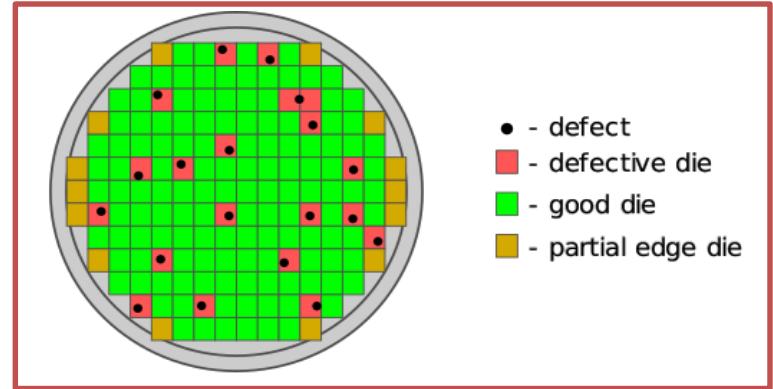
By taking square root

$$\sqrt{\text{Yield}} \left(1 + \text{defects per area} \times \frac{\text{diearea}}{2} \right) = 1$$

$$\left(\sqrt{\text{Yield}} + \sqrt{\text{Yield}} \times \text{defects per area} \times \frac{\text{diearea}}{2} \right) = 1$$

$$\sqrt{\text{Yield}} \times \text{defects per area} \times \frac{\text{diearea}}{2} = 1 - \sqrt{\text{Yield}}$$

$$\text{defects per area} = \left(\frac{1 - \sqrt{\text{Yield}}}{\sqrt{\text{Yield}} \times \frac{\text{diearea}}{2}} \right) \quad \text{--- (6)}$$



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 1

1. Find the number of dies 300 mm (30 cm) wafer for a die that is 1.5 cm on a side and for a die that is 1.0 cm on a side.



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 1 Solution

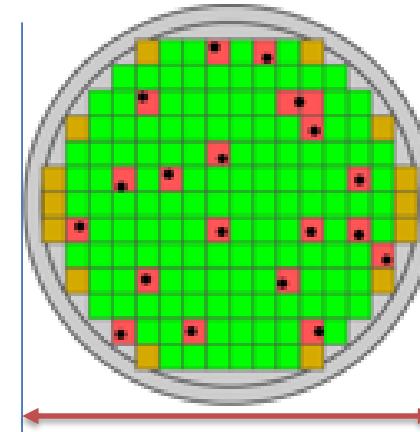
Example Find the number of dies per 300 mm (30 cm) wafer for a die that is 1.5 cm on a side and for a die that is 1.0 cm on a side.

Answer When die area is 2.25 cm^2 :

$$\text{Dies per wafer} = \frac{\pi \times (30/2)^2}{2.25} - \frac{\pi \times 30}{\sqrt{2 \times 2.25}} = \frac{706.9}{2.25} - \frac{94.2}{2.12} = 270$$

Because the area of the larger die is 2.25 times bigger, there are roughly 2.25 as many smaller dies per wafer:

$$\text{Dies per wafer} = \frac{\pi \times (30/2)^2}{1.00} - \frac{\pi \times 30}{\sqrt{2 \times 1.00}} = \frac{706.9}{1.00} - \frac{94.2}{1.41} = 640$$



Diameter of Wafer	300mm
Square Die width	1.5cm
Square Die width	1cm

Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 2

2. Find the die yield for dies that are 1.5 cm on a side and 1.0 cm on a side, assuming a defect of 0.047 per cm^2 and N is 12.



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 2

Example Find the die yield for dies that are 1.5 cm on a side and 1.0 cm on a side, assuming a defect density of 0.047 per cm^2 and N is 12.

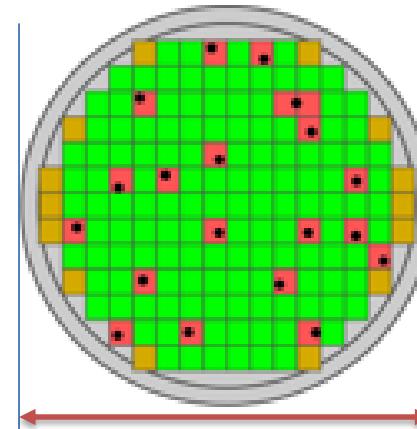
Answer The total die areas are 2.25 and 1.00 cm^2 . For the larger die, the yield is

$$\text{Die yield} = 1 / (1 + 0.047 \times 2.25)^{12} \times 270 = 120$$

For the smaller die, the yield is

$$\text{Die yield} = 1 / (1 + 0.047 \times 1.00)^{12} \times 640 = 444$$

The bottom line is the number of good dies per wafer. Less than half of all the large dies are good, but nearly 70% of the small dies are good.



Square Die width	1.5cm
Square Die width	1cm
Defects/Unit Area	0.047
N	12

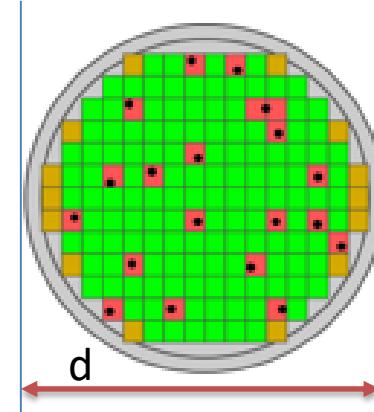
Computer Abstractions and Technology

Technologies for Building Processors and Memory

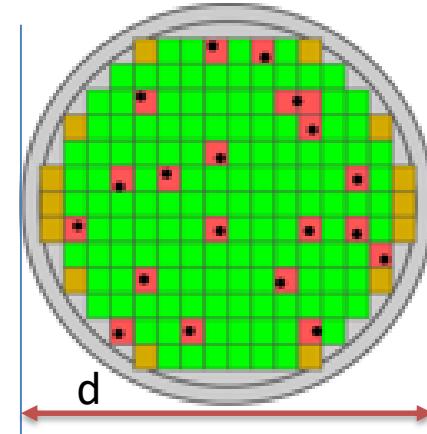
Numerical 3

Assume a 15 cm diameter wafer has a cost of 12, contains 84 dies, and has $0.020 \text{ defects/cm}^2$. Assume a 20 cm diameter wafer has a cost of 15, contains 100 dies, and has $0.031 \text{ defects/cm}^2$.

- Find the yield for both wafers.
- Find the cost per die for both wafers.
- If the number of dies per wafer is increased by 10% and the defects per area unit increases by 15%, find the die area and yield.
- Assume a fabrication process improves the yield from 0.92 to 0.95. Find the defects per area unit for each version of the technology given a die area of 200 mm^2 .



Diameter of Wafer	15cm
Cost of wafer	12
Number of Die	84
Defects/ cm^2	0.020
N	12



Diameter of Wafer	20cm
Cost of wafer	125
Number of Die	100
Defects/ cm^2	0.031
N	12

Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 3

a) Die Are(15cm) = $\left(\frac{\text{Wafer Area}}{\text{Dies per wafer}}\right) = \left(\frac{\pi r^2}{84}\right) = \left(\frac{3.14 \times 7.5^2}{84}\right) = 2.10 \text{ cm}^2$

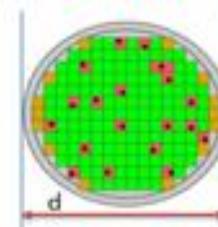
Die Are(20cm) = 3.14 cm^2

Yield(15cm) = $\left(\frac{1}{\left(1 + \text{Defects per area} \times \frac{\text{Die Area}}{2}\right)^2}\right) = \left(\frac{1}{\left(1 + 0.020 \times \frac{2.10}{2}\right)^2}\right) = 0.9593$

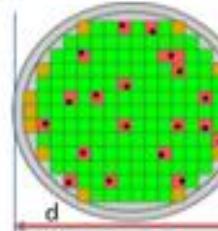
Yield(20cm) = 0.9093

b) Cost per Die (15cm) = $\left(\frac{\text{Cost per Wafer}}{\text{Dies per wafer} \times \text{Yield}}\right) = \left(\frac{12}{84 \times 0.9593}\right) = 0.1489$

Cost per Die (20cm) = $\left(\frac{15}{100 \times 0.9093}\right) = 0.1650$



Diameter of Wafer	15cm
Cost of wafer	12
Number of Die	84
Defects/cm ²	0.020
N	12



Diameter of Wafer	20cm
Cost of wafer	125
Number of Die	100
Defects/cm ²	0.031
N	12

Computer Abstractions and Technology

Technologies for Building Processors and Memory

Cost and Prize of the Chip

Defects per Area can be derived from eqn (4) or eqn(5). From eqn (4)

$$\text{Yield} \left(1 + \text{defects per area} \times \frac{\text{diearea}}{2} \right)^2 = 1$$

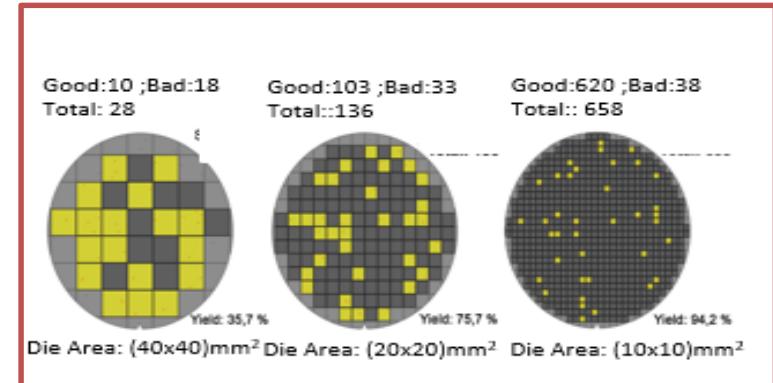
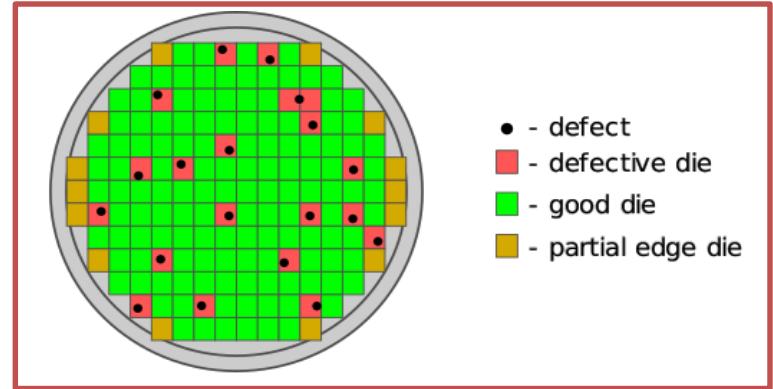
By taking square root

$$\sqrt{\text{Yield}} \left(1 + \text{defects per area} \times \frac{\text{diearea}}{2} \right) = 1$$

$$\left(\sqrt{\text{Yield}} + \sqrt{\text{Yield}} \times \text{defects per area} \times \frac{\text{diearea}}{2} \right) = 1$$

$$\sqrt{\text{Yield}} \times \text{defects per area} \times \frac{\text{diearea}}{2} = 1 - \sqrt{\text{Yield}}$$

$$\text{defects per area} = \left(\frac{1 - \sqrt{\text{Yield}}}{\sqrt{\text{Yield}} \times \frac{\text{diearea}}{2}} \right) \quad \text{--- (6)}$$



Computer Abstractions and Technology

Technologies for Building Processors and Memory

Numerical 3

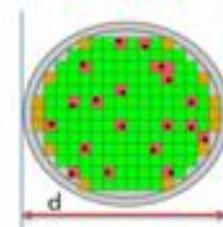
$$c) \text{ Die Area}(15\text{cm}) = \left(\frac{\text{Wafer Area}}{\text{Dies per wafer}} \right) = \left(\frac{\pi r^2}{84} \right) = \left(\frac{3.14 \times 7.52}{84 \times 1.1} \right) = 1.91 \text{ cm}^2$$

$$\text{Die Are}(20\text{cm}) = \left(\frac{3.14 \times 7.52}{100 \times 1.1} \right) = 2.86 \text{ cm}^2$$

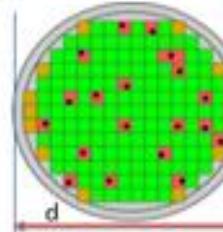
$$\text{Yield}(15\text{cm}) = \left(\frac{1}{\left(1 + \text{Defects per area} \times \frac{\text{Die Area}}{2} \right)^2} \right) = \left(\frac{1}{\left(1 + 0.020 \times 1.15 \times \frac{1.91}{2} \right)^2} \right) = 0.9575$$

$$\text{Yield}(20\text{cm}) = \left(\frac{1}{\left(1 + 0.031 \times 1.15 \times \frac{2.86}{2} \right)^2} \right) = 0.9082$$

$$d) \text{ defects per area} = \left(\frac{1 - \sqrt{\text{Yield}}}{\sqrt{\text{Yield}} \times \frac{\text{diearea}}{2}} \right) = \left(\frac{1 - \sqrt{0.92}}{\sqrt{0.92} \times \frac{2\text{cm}^2}{2}} \right) = 0.043 \text{ defects/cm}^2$$



Diameter of Wafer	15cm
Cost of wafer	12
Number of Die	84
Defects/cm ²	0.020
N	12



Diameter of Wafer	20cm
Cost of wafer	125
Number of Die	100
Defects/cm ²	0.031
N	12

Computer Abstractions and Technology

Technologies for Building Processors and Memory

Fun Activity

<http://www.silicon-edge.co.uk/j/index.php/resources/die-per-wafer>





THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 5 – Performance – Part 1

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Defining Performance

When we say one computer has better performance than another, what do we mean?



Defining Performance

An analogy with passenger airplanes to define Performance

Airplane	Passenger capacity	Cruising range (miles)	Cruising speed (m.p.h.)	Passenger throughput (passengers × m.p.h.)
Boeing 737	240	3000	564	135,360
BAC/Sud Concorde	132	4000	1350	178,200
Boeing 777-200LR	301	9395	554	166,761
Airbus A380-800	853	8477	587	500,711

Which of the planes in this table has the best performance?

Answer : We would first need to define performance.

Defining Performance

Airplane	Passenger capacity	Cruising range (miles)	Cruising speed (m.p.h.)	Passenger throughput (passengers × m.p.h.)
Boeing 737	240	3000	564	135,360
BAC/Sud Concorde	132	4000	1350	178,200
Boeing 777-200LR	301	9395	554	166,761
Airbus A380-800	853	8477	587	500,711

- The plane with the **highest cruising speed** was the Concorde
- The plane with the **longest range** is the Boeing 777- 200LR
- The plane with the **largest capacity** is the Airbus A380-800

Defining Performance

Airplane	Passenger capacity	Cruising range (miles)	Cruising speed (m.p.h.)	Passenger throughput (passengers × m.p.h.)
Boeing 737	240	3000	564	135,360
BAC/Sud Concorde	132	4000	1350	178,200
Boeing 777-200LR	301	9395	554	166,761
Airbus A380-800	853	8477	587	500,711

Let's suppose we define performance in terms of speed.

You could define the fastest plane as the one with the highest cruising speed, taking a single passenger from one point to another in the least time.

If you were interested in transporting 500 passengers from one point to another, however, the Airbus A380-800 would clearly be the fastest, as the last column of the figure shows.

Similarly, we can define computer performance in several distinct ways.

Basic Definitions

Response time also called Execution time:

The total time required for the computer to complete a task, including disk accesses, memory accesses, I/O activities, operating system overhead, CPU execution time, and so on.

Throughput also called Bandwidth:

Another measure of performance, it is the number of tasks completed per unit time.

Basic Definitions

Task : Throughput and Response Time

Do the following changes to a computer system **increase throughput, decrease response time, or both?**

1. Replacing the processor in a computer with a faster version
2. Adding additional processors to a system that uses multiple processors for separate tasks.



Basic Definitions

Task : Throughput and Response Time

Do the following changes to a computer system **increase throughput, decrease response time, or both?**

1. Replacing the processor in a computer with a faster version

Decreasing response time almost always improves throughput. Hence, in case 1, both response time and throughput are improved.

Basic Definitions

Task : Throughput and Response Time

Do the following changes to a computer system **increase throughput, decrease response time, or both?**

2. Adding additional processors to a system that uses multiple processors for separate tasks.

In case 2, no one task gets work done faster, so only throughput increases.

Performance of a Computer

- To maximize performance, we want to **minimize response time or execution time for some task.**
i.e., Lesser Response time Maximum is the performance.
- Thus, we can relate performance and execution time for a computer X:

$$\text{Performance}_X = \frac{1}{\text{Execution time}_X}$$

- For two computers X and Y, if the performance of X is greater than the performance of Y, we have

$$\begin{aligned}\text{Performance}_X &> \text{Performance}_Y \\ \frac{1}{\text{Execution time}_X} &> \frac{1}{\text{Execution time}_Y} \\ \text{Execution time}_Y &> \text{Execution time}_X\end{aligned}$$

Performance of a Computer

- we often want to relate the performance of two different computers quantitatively

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = n$$

- If X is n times as fast as Y, then the execution time on Y is n times as long as it is on X:

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution time}_Y}{\text{Execution time}_X} = n$$

Relative Performance



Computer X runs a program in 10 seconds

Computer Y runs the same program in 15 seconds



How much faster is A than B?

Performance of a Computer

Relative Performance



Computer X runs a program in 10 seconds



Computer Y runs the same program in 15 seconds

How much faster is A than B?

We know that A is n times as fast as B if

$$\frac{\text{Performance}_A}{\text{Performance}_B} = \frac{\text{Execution time}_B}{\text{Execution time}_A} = n$$

Thus the performance ratio is

$$\frac{15}{10} = 1.5$$

and A is therefore 1.5 times as fast as B.

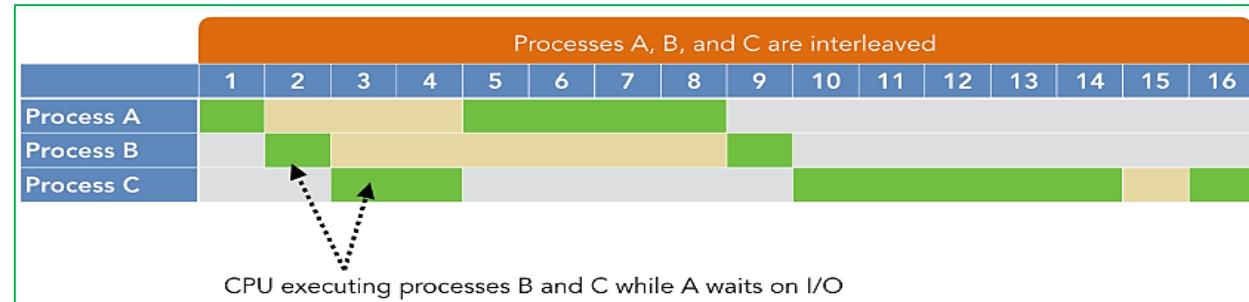
Measuring Performance

Wall clock time, Response time, or Elapsed time: The total time to complete a task, including

- ✓ disk accesses,
- ✓ memory accesses,
- ✓ input/output (I/O) activities,
- ✓ operating system overhead—everything.

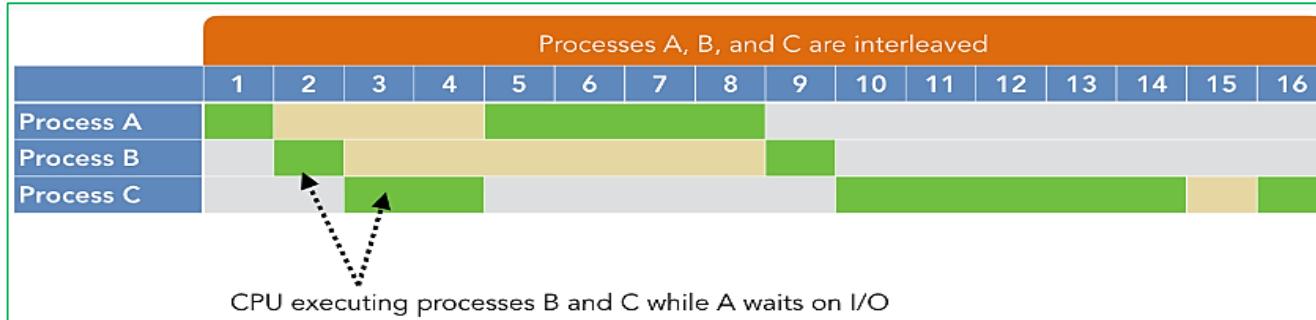
Computers are often shared, however, and a processor may work on several programs simultaneously.

In such cases, the system may try to optimize throughput rather than attempt to minimize the elapsed time for one program.



Measuring Performance

CPU execution time (CPU time) : It is the time the CPU spends computing a particular task and does not include time spent waiting for I/O or running other programs.



CPU time can be further divided into -

User CPU time: the amount of time the processor spends in running your application code

System CPU time: the time spent running code in the operating system kernel on behalf of your program

We will use **CPU performance** to refer to user CPU time and concentrate on the same.

CPU Performance and Its Factors

Different applications are sensitive to Different aspects of the performance of a computer system

Example: Applications running on servers, depend as much on I/O performance, which, in turn, relies on both hardware and software.

CPU Performance and Its Factors

CPU performance - the bottom-line performance measure is CPU execution time.

A simple formula relates the most basic metrics (clock cycles and clock cycle time) to CPU time:

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clock cycle time}}$$

This formula makes it clear that the hardware designer can improve performance by reducing the number of clock cycles required for a program or the length of the clock cycle.

Alternatively, because clock rate and clock cycle time are inverses,

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate}}$$

CPU Performance and Its Factors

Computer A



Our favorite program runs in 10 seconds on computer A, which has a 2 GHz clock.

The designer has determined that a substantial increase in the clock rate is possible, but this increase will affect the rest of the CPU design, causing computer B to require 1.2 times as many clock cycles as computer A for this program. What clock rate should we tell the designer to target?

Computer B



Computer designer build a computer, B, which will run this program in 6 seconds

First find the CPU time for A

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate}}$$

$$\text{CPU time}_A = \frac{\text{CPU clock cycles}_A}{\text{Clock rate}_A}$$

$$10 \text{ seconds} = \frac{\text{CPU clock cycles}_A}{2 \times 10^9 \frac{\text{cycles}}{\text{second}}}$$

$$\text{CPU clock cycles}_A = 10 \text{ seconds} \times 2 \times 10^9 \frac{\text{cycles}}{\text{second}} = 20 \times 10^9 \text{ cycles}$$

CPU Performance and Its Factors

Computer A



Our favorite program runs in 10 seconds on computer A, which has a 2 GHz clock.

The designer has determined that a substantial increase in the clock rate is possible, but this increase will affect the rest of the CPU design, causing computer B to require 1.2 times as many clock cycles as computer A for this program. What clock rate should we tell the designer to target?

Computer B



Computer designer build a computer, B, which will run this program in 6 seconds

Answer: 4 GHz

$$\text{Clock rate}_B = \frac{1.2 \times 20 \times 10^9 \text{ cycles}}{6 \text{ seconds}} = \frac{0.2 \times 20 \times 10^9 \text{ cycles}}{\text{second}} = \frac{4 \times 10^9 \text{ cycles}}{\text{second}} = 4 \text{ GHz}$$



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 5 – Performance – Part 2

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Instruction Performance

- Program is basically set of Instructions.
- The computer had to execute the instructions to run the program, therefore execution time must depend on the number of instructions in a program.
- Therefore, the number of clock cycles required for a program can be written as

$$\text{CPU execution time for a program} = \frac{\text{CPU clock cycles for a program}}{\text{Clockrate}}$$

$$\text{CPU clock cycles} = \text{Instructions for a program} \times \frac{\text{Average clock cycles per instruction}}$$

Instruction Performance

Clock cycles per instruction (CPI):

- The average number of clock cycles each instruction takes to execute.
- Since different instructions may take different amounts of time depending on what they do, CPI is an average of all the instructions executed in the program.
- **CPI provides one way of comparing two different implementations** of the identical instruction set architecture, since the number of instructions executed for a program will, of course, be the same.

Computer Abstractions and Technology

Performance

Instruction Performance

Suppose we have two implementations of the same instruction set architecture.

Computer A



Computer A has a clock cycle time of 250ps and a CPI of 2.0 for some program

Computer B



computer B has a clock cycle time of 500 ps and a CPI of 1.2 for the same program

Which computer is faster for this program and by how much?

First, find the number of processor clock cycles for each computer:

$$\text{CPU clock cycles}_A = I \times 2.0$$
$$\text{CPU clock cycles}_B = I \times 1.2$$

Now we can compute the CPU time for each computer:

$$\begin{aligned}\text{CPU time}_A &= \text{CPU clock cycles}_A \times \text{Clock cycle time} \\ &= I \times 2.0 \times 250 \text{ ps} = 500 \times I \text{ ps}\end{aligned}$$

Computer Abstractions and Technology

Performance

Instruction Performance

Suppose we have two implementations of the same instruction set architecture.

Computer A



Computer A has a clock cycle time of 250ps and a CPI of 2.0 for some program

Computer B



Computer B has a clock cycle time of 500 ps and a CPI of 1.2 for the same program

Which computer is faster for this program and by how much?

Now we can compute the CPU time for each computer:

$$\text{CPU time}_B = I \times 1.2 \times 500 \text{ ps} = 600 \times I \text{ ps}$$

Computer Abstractions and Technology

Performance

Instruction Performance

Suppose we have two implementations of the same instruction set architecture.

Computer A



Computer A has a clock cycle time of 250ps and a CPI of 2.0 for some program

Computer B



computer B has a clock cycle time of 500 ps and a CPI of 1.2 for the same program

Which computer is faster for this program and by how much?

$$\frac{\text{CPU performance}_A}{\text{CPU performance}_B} = \frac{\text{Execution time}_B}{\text{Execution time}_A} = \frac{600 \times I \text{ ps}}{500 \times I \text{ ps}} = 1.2$$

Conclusion:

Computer A is 1.2 times as fast as computer B for this program.

The Classic CPU Performance Equation

Basic performance equation in terms of **instruction count** (the number of instructions executed by the program), **CPI**, and **clock cycle time**:

$$\text{CPU time} = \text{Instruction count} \times \text{CPI} \times \text{Clock cycle time}$$

$$\text{CPU time} = \frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}$$

- The three key factors that affect performance Instruction Count, CPI and Clock rate.
- We can use these formulas to
 - ✓ Compare two different implementations or
 - ✓ To evaluate a design alternative if we know its impact on these three parameters.

Comparing Code Segments

A compiler designer is trying to decide between two code sequences for a computer. The hardware designers have supplied the following facts:

	CPI for each instruction class		
	A	B	C
CPI	1	2	3

For a particular high-level language statement, the compiler writer is considering two code sequences that require the following instruction counts:

Code sequence	Instruction counts for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

Which code sequence executes the most instructions? Which will be faster?
What is the CPI for each sequence?

- Number of Instructions in Code Sequence 1 : $2+1+2=5$
- Number of Instructions in Code Sequence 2 : $4+1+2=6$

Code Sequence 1 executes fewer Instructions

Comparing Code Segments

A compiler designer is trying to decide between two code sequences for a computer. The hardware designers have supplied the following facts:

	CPI for each instruction class		
	A	B	C
CPI	1	2	3

For a particular high-level language statement, the compiler writer is considering two code sequences that require the following instruction counts:

Code sequence	Instruction counts for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

Which code sequence executes the most instructions? Which will be faster?
What is the CPI for each sequence?

- Total number of clock cycles for sequence 1:
 $= 2 \times 1 + 1 \times 2 + 2 \times 3 = 10 \text{ clock cycles}$
- Total number of clock cycles for sequence 1:
 $= 4 \times 1 + 1 \times 2 + 1 \times 3 = 9 \text{ clock cycles}$

So code sequence 2 is faster, even though it executes one extra instruction.

Comparing Code Segments

A compiler designer is trying to decide between two code sequences for a computer. The hardware designers have supplied the following facts:

	CPI for each instruction class		
	A	B	C
CPI	1	2	3

For a particular high-level language statement, the compiler writer is considering two code sequences that require the following instruction counts:

Code sequence	Instruction counts for each instruction class		
	A	B	C
1	2	1	2
2	4	1	1

Which code sequence executes the most instructions? Which will be faster?
What is the CPI for each sequence?

- CPI of Sequence 1 and 2 are

$$CPI_1 = \frac{\text{CPU clock cycles}_1}{\text{Instruction count}_1} = \frac{10}{5} = 2.0$$

$$CPI_2 = \frac{\text{CPU clock cycles}_2}{\text{Instruction count}_2} = \frac{9}{6} = 1.5$$

The basic components of performance and how each is measured.

Components of performance	Units of measure
CPU execution time for a program	Seconds for the program
Instruction count	Instructions executed for the program
Clock cycles per instruction (CPI)	Average number of clock cycles per instruction
Clock cycle time	Seconds per clock cycle

FIGURE 1.15 The basic components of performance and how each is measured.

The performance of a program depends on the algorithm, the language, the compiler, the architecture, and the actual hardware.

Hardware or software component	Affects what?	How?
Algorithm	Instruction count, CPI	<ul style="list-style-type: none">▪ The algorithm determines the number of source program instructions executed and hence the number of processor instructions executed.▪ The algorithm may also affect the CPI, by favoring slower or faster instructions.▪ For example, if the algorithm uses more divides, it will tend to have a higher CPI.

The performance of a program depends on the algorithm, the language, the compiler, the architecture, and the actual hardware.

Hardware or software component	Affects what?	How?
Programming language	Instruction count, CPI	<ul style="list-style-type: none">The programming language certainly affects the instruction count, since statements in the language are translated to processor instructions, which determine instruction count.The language may also affect the CPI because of its features; for example, a language with heavy support for data abstraction (e.g., Java) will require indirect calls, which will use higher CPI instructions.

The performance of a program depends on the algorithm, the language, the compiler, the architecture, and the actual hardware.

Hardware or software component	Affects what?	How?
Compiler	Instruction count, CPI	<ul style="list-style-type: none">▪ The efficiency of the compiler affects both the instruction count and average cycles per instruction, since the compiler determines the translation of the source language instructions into computer instructions.▪ The compiler's role can be very complex and affect the CPI in varied ways.

The performance of a program depends on the algorithm, the language, the compiler, the architecture, and the actual hardware.

Hardware or software component	Affects what?	How?
Instruction set architecture	Instruction count, clock rate, CPI	<ul style="list-style-type: none">The instruction set architecture affects all three aspects of CPU performance, since it affects the instructions needed for a function, the cost in cycles of each instruction, and the overall clock rate of the processor.

IPC (instruction per Clock Cycle):

- Although you might expect that the minimum CPI is 1.0, some processors fetch and execute multiple instructions per clock cycle.
- To reflect that approach, some designers invert CPI to talk about **IPC**, or *instructions per clock cycle*.
- If a processor executes on average two instructions per clock cycle, then it has an IPC of 2 and hence a CPI of 0.5.

Turbo mode:

- Although clock cycle time has traditionally been fixed, to save energy or temporarily boost performance, today's processors can vary their clock rates, so we would need to use the *average* clock rate for a program.
- For example, the Intel Core i7 will temporarily increase clock rate by about 10% until the chip gets too warm.
- **Intel calls this Turbo mode.**



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 6 – The Power Wall

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

The Clock rate and Power for Intel x86 microprocessors over nine generations and 36 years

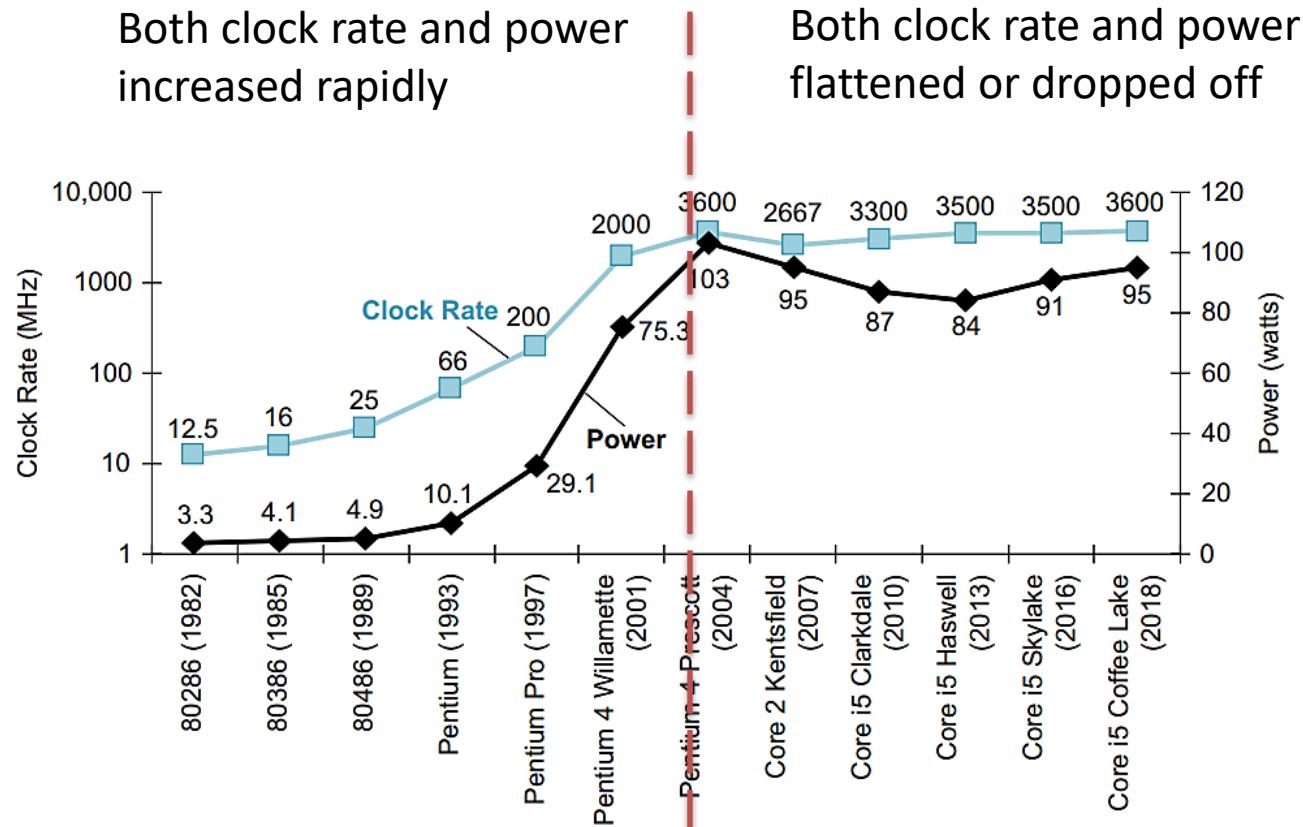


FIGURE 1.16 Clock rate and power for Intel x86 microprocessors over nine generations

- The dominant technology for **integrated circuits** is called **CMOS**
- For CMOS, the primary source of energy consumption is so-called dynamic energy i.e., the energy that is consumed when transistors switch states
- The dynamic energy depends on the **capacitive loading** of each transistor and the **voltage applied**

$$\text{Energy} \propto \text{Capacitive load} \times \text{Voltage}^2$$

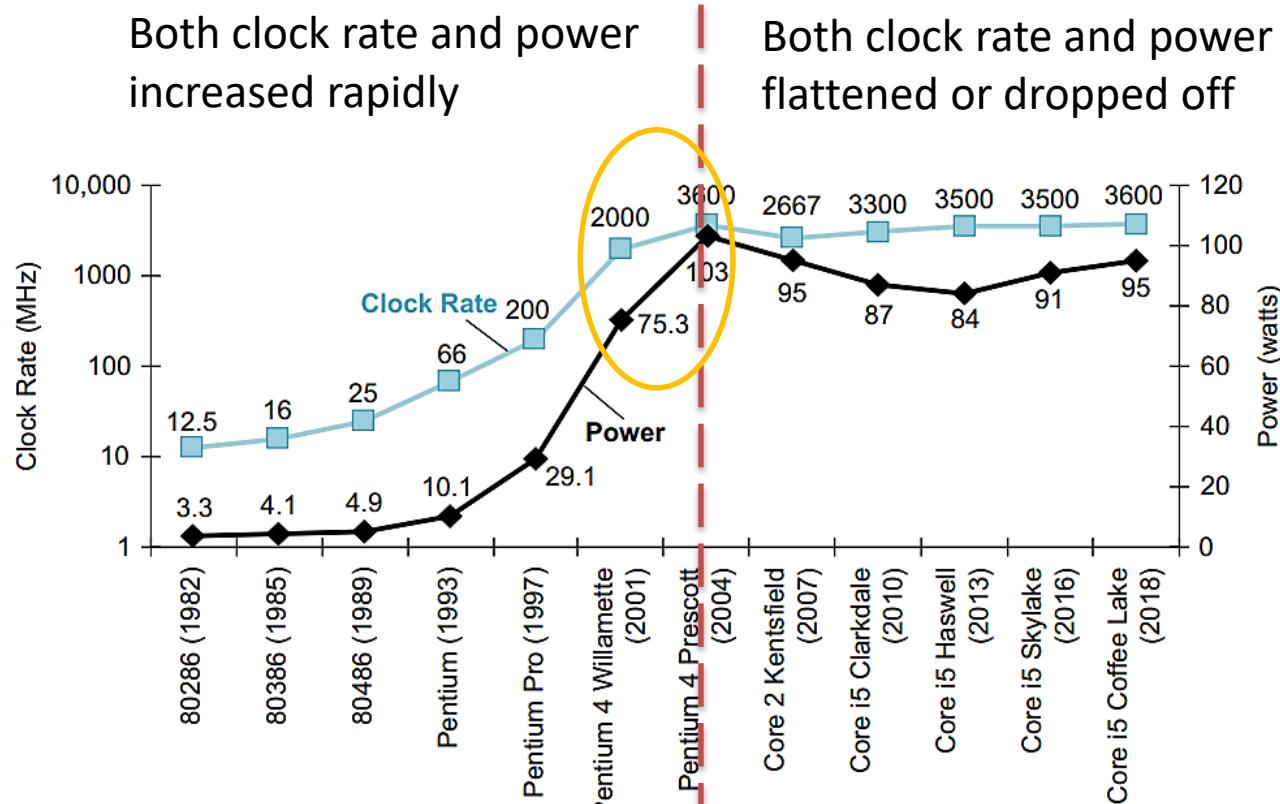
This equation is the energy of a pulse during the logic transition of $0 \rightarrow 1 \rightarrow 0$ or $1 \rightarrow 0 \rightarrow 1$. The energy of a single transition is then

$$\text{Energy} \propto 1/2 \times \text{Capacitive load} \times \text{Voltage}^2$$

The power required per transistor is just the product of energy of a transition and the frequency of transitions:

$$\text{Power} \propto 1/2 \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}$$

The Clock rate and Power for Intel x86 microprocessors over nine generations and 36 years



How could clock rates grow by a factor of 1000 while power increased by only a factor of 30?



Energy and thus power can be reduced by lowering the voltage, which occurred with each new generation of technology, and power is a function of the voltage².

FIGURE 1.16 Clock rate and power for Intel x86 microprocessors over nine generations

Relative Power

Suppose we developed a new, simpler processor that has 85% of the capacitive load of the more complex older processor.

Further, assume that it can adjust voltage so that it can reduce voltage 15% compared to processor B, which results in a 15% shrink in frequency. What is the impact on dynamic power?

$$\frac{\text{Power}_{\text{new}}}{\text{Power}_{\text{old}}} = \frac{\text{Capacitive load} \times 0.85 \times \text{Voltage} \times 0.85^2 \times \text{Frequency switched} \times 0.85}{\text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}}$$

Thus the power ratio is

$$0.85^4 = 0.52$$

Hence, the new processor uses about half the power of the old processor.



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 7 – The Switch from Uniprocessors to Multiprocessors

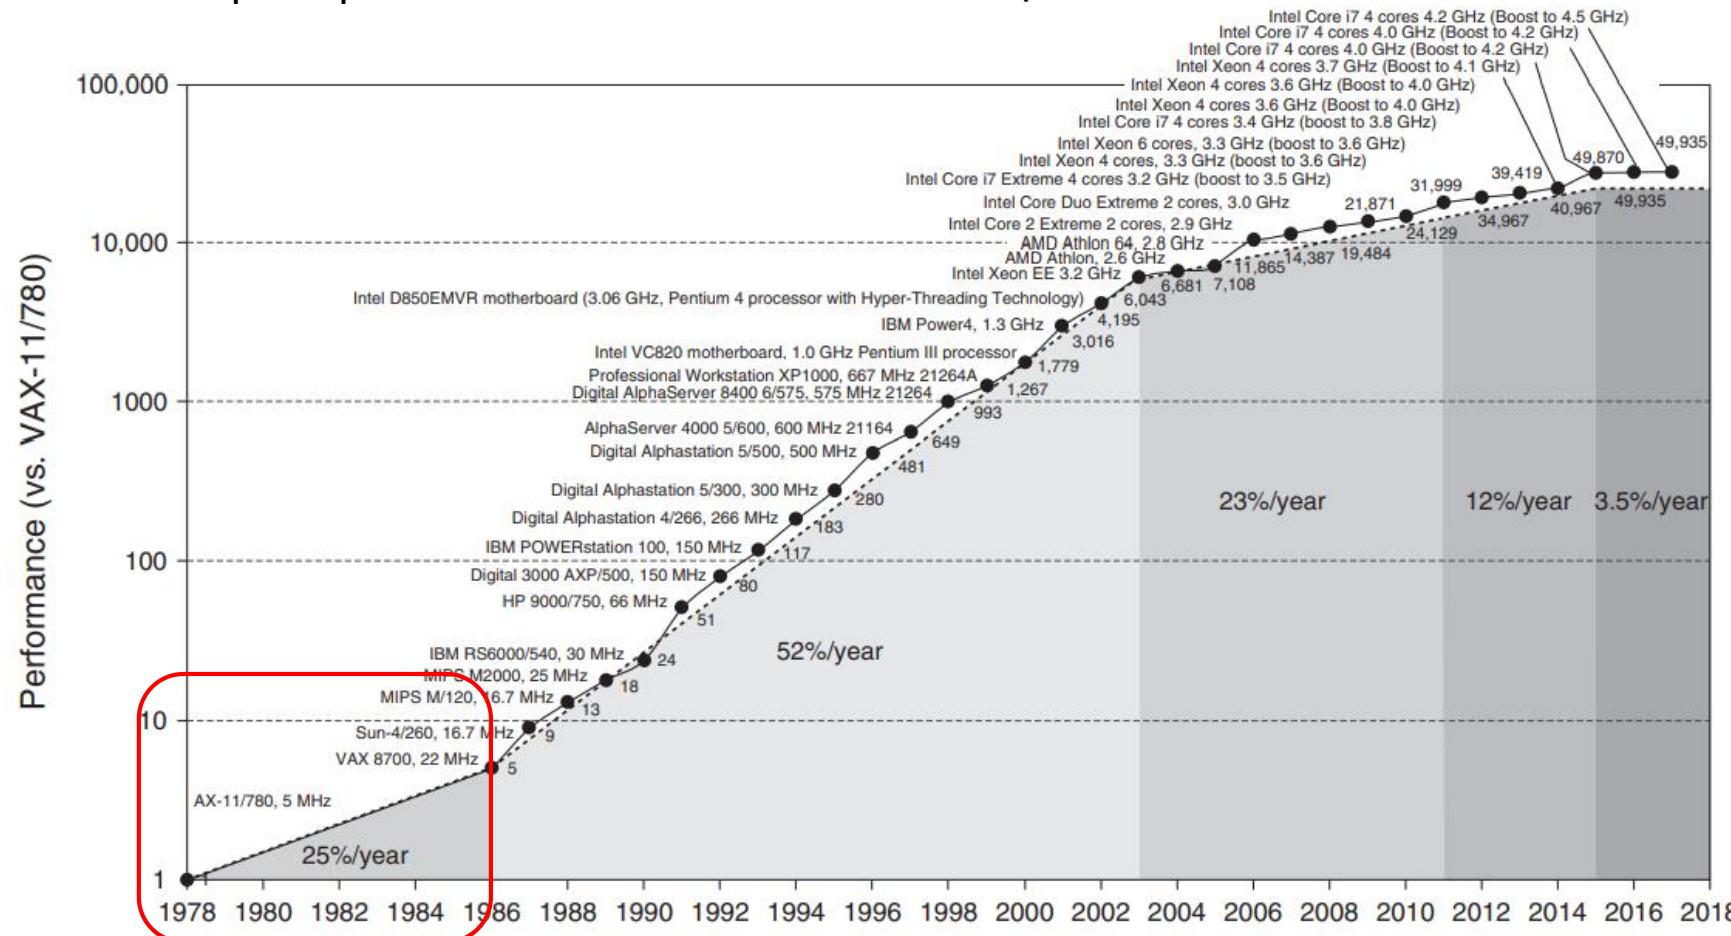
Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

The Sea Change: The Switch from Uniprocessors to Multiprocessors

Growth in Processor Performance since the mid-1980s

This chart plots performance relative to the VAX 11/780

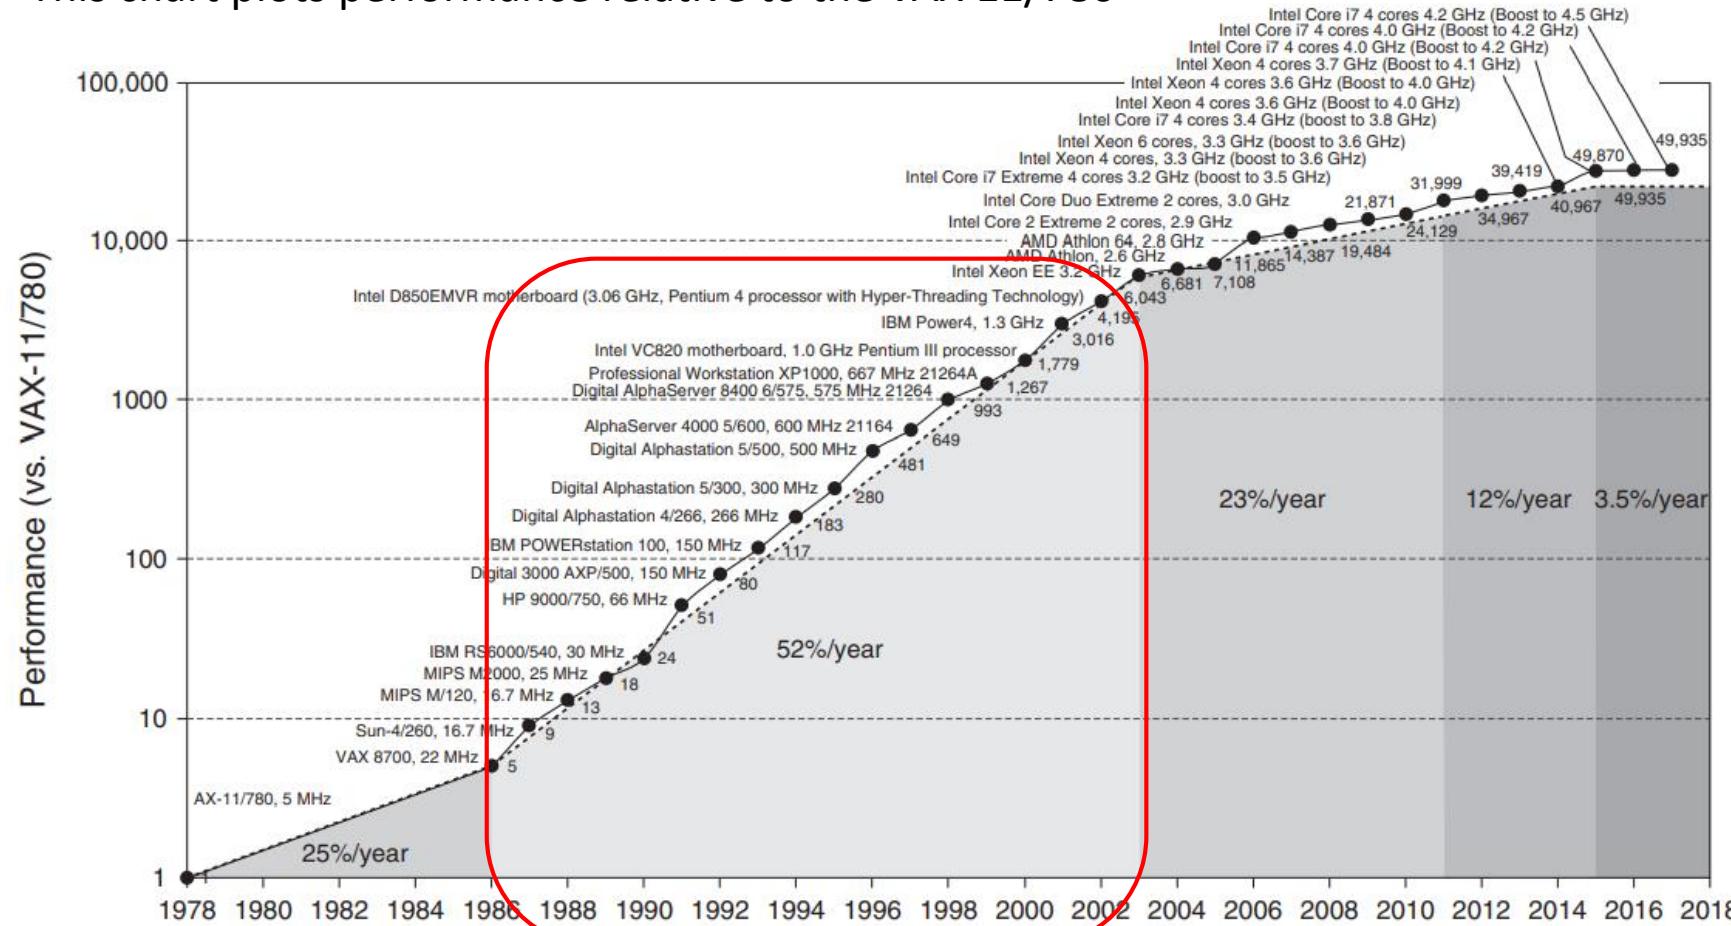


Prior to the mid-1980s, processor performance growth was largely technology driven and averaged about 25% per year.

The Sea Change: The Switch from Uniprocessors to Multiprocessors

Growth in Processor Performance since the mid-1980s

This chart plots performance relative to the VAX 11/780

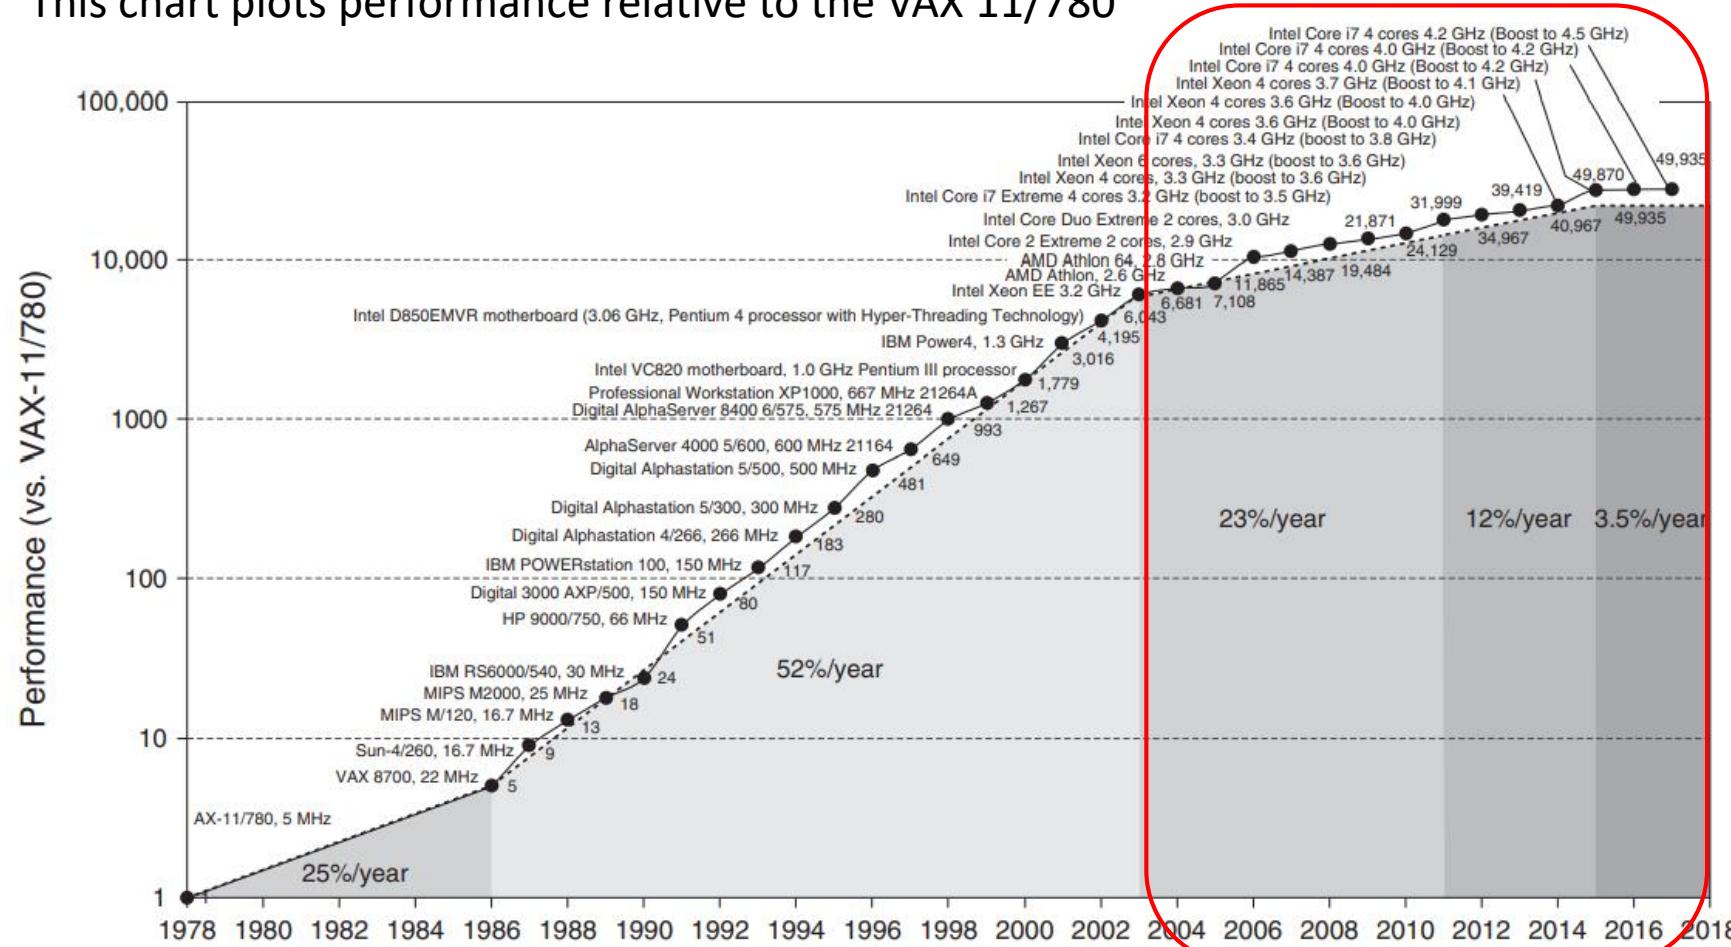


The increase in growth to about 52% since then is attributable to more advanced architectural and organizational ideas

The Sea Change: The Switch from Uniprocessors to Multiprocessors

Growth in Processor Performance since the mid-1980s

This chart plots performance relative to the VAX 11/780



Since 2002, the limits of power, available instruction-level parallelism, and long memory latency have slowed uniprocessor performance recently, to about 3.5% per year.

The Sea Change: The Switch from Uniprocessors to Multiprocessors

What do we do??

From mid 1980s

Continuing to decrease the response time of one program running on the single processor

As of 2006

All desktop and server companies are shipping microprocessors with multiple processors per chip

Benefit is often more on **throughput** than on response time

Companies refer to processors as “cores,” and such microprocessors are generically called multicore microprocessors.

Hence, a “quadcore” microprocessor is a chip that contains four processors or four cores.

The Sea Change: The Switch from Uniprocessors to Multiprocessors

Challenges for Programmers

- Rewrite their programs to take advantage of multiple processors.
- Programmers will have to continue to improve the performance of their code as the number of cores increases

The Sea Change: The Switch from Uniprocessors to Multiprocessors

Challenges for Programmers

Why has it been so hard for programmers to write explicitly parallel programs?

- Parallel programming is by **definition performance programming**, which increases the difficulty of programming.
- Fast for parallel hardware means that the programmer must **divide an application** so that each processor has roughly the same amount to do at the same time, and that the overhead of scheduling and coordination doesn't fritter away the potential performance benefits of parallelism.

The Sea Change: The Switch from Uniprocessors to Multiprocessors

To reflect this sea change in the industry -

- Parallelism and Instructions: *Synchronization*.
- Parallelism and Computer Arithmetic: *Subword Parallelism*.
- Parallelism via Instructions.
- Parallelism and Memory Hierarchies: *Cache Coherence*.
- Parallelism and Memory Hierarchy: *Redundant Arrays of Inexpensive Disks*.



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 8 – Benchmarking the Intel Core i7

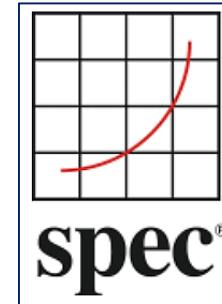
Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

SPEC CPU Benchmark

Workload: A set of programs run on a computer

To evaluate two computer systems, a user would simply compare the execution time of the workload on the two computers.



Benchmarks: Programs specifically chosen to measure performance

The benchmarks form a workload that the user hopes will predict the performance of the actual workload.

SPEC: (System Performance Evaluation Cooperative) is an effort funded and supported by a number of computer vendors to create standard sets of benchmarks for modern computer systems.

Benchmarking the Intel Core i7

SPEC CPU Benchmark

- In 1989, SPEC originally created a benchmark set focusing on processor performance (now called SPEC89).
- The latest is SPEC CPU2017, which consists of a set of **10 integer benchmarks** (SPECspeed 2017 Integer) and **13 floating-point benchmarks** (CFP2006).
- The **integer benchmarks** vary from part of a C compiler to a chess program to a quantum computer simulation.
- The **floating-point benchmarks** include structured grid codes for finite element modeling, particle method codes for molecular dynamics, and sparse linear algebra codes for fluid dynamics.

Benchmarking the Intel Core i7

SPECspeed 2017 Integer benchmarks running on a 1.8 GHz Intel Xeon E5-2650L

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Perl interpreter	perlbench	2684	0.42	0.556	627	1774	2.83
GNU C compiler	gcc	2322	0.67	0.556	863	3976	4.61
Route planning	mcf	1786	1.22	0.556	1215	4721	3.89
Discrete Event simulation - computer network	omnetpp	1107	0.82	0.556	507	1630	3.21
XML to HTML conversion via XSLT	xalancbmk	1314	0.75	0.556	549	1417	2.58
Video compression	x264	4488	0.32	0.556	813	1763	2.17
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	2216	0.57	0.556	698	1432	2.05
Artificial Intelligence: Monte Carlo tree search (Go)	leela	2236	0.79	0.556	987	1703	1.73
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	6683	0.46	0.556	1718	2939	1.71
General data compression	xz	8533	1.32	0.556	6290	6182	0.98
Geometric mean	—	—	—	—	—	—	2.36

Table describes the SPEC integer benchmarks and their execution time on the Intel Core i7 and shows the factors that explain execution time:

- Instruction count
- CPI
- Clock cycle time

Benchmarking the Intel Core i7

SPECspeed 2017 Integer benchmarks running on a 1.8 GHz Intel Xeon E5-2650L

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Perl interpreter	perlbench	2684	0.42	0.556	627	1774	2.83
GNU C compiler	gcc	2322	0.67	0.556	863	3976	4.61
Route planning	mcf	1786	1.22	0.556	1215	4721	3.89
Discrete Event simulation - computer network	omnetpp	1107	0.82	0.556	507	1630	3.21
XML to HTML conversion via XSLT	xalancbmk	1314	0.75	0.556	549	1417	2.58
Video compression	x264	4488	0.32	0.556	813	1763	2.17
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	2216	0.57	0.556	698	1432	2.05
Artificial Intelligence: Monte Carlo tree search (Go)	leela	2236	0.79	0.556	987	1703	1.73
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	6683	0.46	0.556	1718	2939	1.71
General data compression	xz	8533	1.32	0.556	6290	6182	0.98
Geometric mean	-	-	-	-	-	-	2.36

Execution Time:

Execution Time = Instruction Count X CPI X Clock cycle time

Ex:

Perlbench Execution Time : $2684 \times 10^9 \times 0.42 \times 0.556 \times 10^{-9}$
: 627

Benchmarking the Intel Core i7

SPECspeed 2017 Integer benchmarks running on a 1.8 GHz Intel Xeon E5-2650L

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Perl interpreter	perlbench	2684	0.42	0.556	627	1774	2.83
GNU C compiler	gcc	2322	0.67	0.556	863	3976	4.61
Route planning	mcf	1786	1.22	0.556	1215	4721	3.89
Discrete Event simulation - computer network	omnetpp	1107	0.82	0.556	507	1630	3.21
XML to HTML conversion via XSLT	xalancbmk	1314	0.75	0.556	549	1417	2.58
Video compression	x264	4488	0.32	0.556	813	1763	2.17
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	2216	0.57	0.556	698	1432	2.05
Artificial Intelligence: Monte Carlo tree search (Go)	leela	2236	0.79	0.556	987	1703	1.73
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	6683	0.46	0.556	1718	2939	1.71
General data compression	xz	8533	1.32	0.556	6290	6182	0.98
Geometric mean	-	-	-	-	-	-	2.36

Reference Time:

Reference Time = Supplied by SPEC

Ex:

Perlbench Reference Time : 1774

Benchmarking the Intel Core i7

SPECspeed 2017 Integer benchmarks running on a 1.8 GHz Intel Xeon E5-2650L

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Perl interpreter	perlbench	2684	0.42	0.556	627	1774	2.83
GNU C compiler	gcc	2322	0.67	0.556	863	3976	4.61
Route planning	mcf	1786	1.22	0.556	1215	4721	3.89
Discrete Event simulation - computer network	omnetpp	1107	0.82	0.556	507	1630	3.21
XML to HTML conversion via XSLT	xalancbmk	1314	0.75	0.556	549	1417	2.58
Video compression	x264	4488	0.32	0.556	813	1763	2.17
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	2216	0.57	0.556	698	1432	2.05
Artificial Intelligence: Monte Carlo tree search (Go)	leela	2236	0.79	0.556	987	1703	1.73
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	6683	0.46	0.556	1718	2939	1.71
General data compression	xz	8533	1.32	0.556	6290	6182	0.98
Geometric mean	-	-	-	-	-	-	2.36

SPECratio:

$$\text{SPECratio} = \frac{\text{Reference Time}}{\text{Execution Time}}$$

Ex:

Perlbench SPECratio : 1774/627
: 2.83

Benchmarking the Intel Core i7

SPECspeed 2017 Integer benchmarks running on a 1.8 GHz Intel Xeon E5-2650L

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Perl interpreter	perlbench	2684	0.42	0.556	627	1774	2.83
GNU C compiler	gcc	2322	0.67	0.556	863	3976	4.61
Route planning	mcf	1786	1.22	0.556	1215	4721	3.89
Discrete Event simulation - computer network	omnetpp	1107	0.82	0.556	507	1630	3.21
XML to HTML conversion via XSLT	xalancbmk	1314	0.75	0.556	549	1417	2.58
Video compression	x264	4488	0.32	0.556	813	1763	2.17
Artificial Intelligence: alpha-beta tree search (Chess)	deepsjeng	2216	0.57	0.556	698	1432	2.05
Artificial Intelligence: Monte Carlo tree search (Go)	leela	2236	0.79	0.556	987	1703	1.73
Artificial Intelligence: recursive solution generator (Sudoku)	exchange2	6683	0.46	0.556	1718	2939	1.71
General data compression	xz	8533	1.32	0.556	6290	6182	0.98
Geometric mean	-	-	-	-	-	-	2.36

Geometric Mean:

The single number quoted as **SPECspeed 2017 Integer** is the geometric mean of the SPECratios.

SPEC Power Benchmark

- Given the increasing importance of energy and power, SPEC added a benchmark to measure power.
- SPECpower_ssj2008 is the first industry-standard benchmark for measuring both the performance and the power consumption of servers.
- SSJ - Server Side Java

http://www.spec.org/power/docs/SPECpower_ssj2008-Design_ssj.pdf

Benchmarking the Intel Core i7

SPEC Power Benchmark

Target Load %	Performance (ssj_ops)	Average Power (watts)
100%	4,864,136	347
90%	4,389,196	312
80%	3,905,724	278
70%	3,418,737	241
60%	2,925,811	212
50%	2,439,017	183
40%	1,951,394	160
30%	1,461,411	141
20%	974,045	128
10%	485,973	115
0%	0	48
Overall Sum	26,815,444	2,165
$\sum \text{ssj_ops} / \sum \text{power} =$		12,385

FIGURE 1.19 SPECpower_ssj2008 running on a dual socket 2.2 GHz Intel Xeon Platinum 8276L with 192 GiB of DRAM and one 80 GB SSD disk.

Benchmarking the Intel Core i7

SPEC Power Benchmark

Target Load %	Performance (ssj_ops)	Average Power (watts)
100%	4,864,136	347
90%	4,389,196	312
80%	3,905,724	278
70%	3,418,737	241
60%	2,925,811	212
50%	2,439,017	183
40%	1,951,394	160
30%	1,461,411	141
20%	974,045	128
10%	485,973	115
0%	0	48
Overall Sum	26,815,444	2,165
$\sum \text{ssj_ops} / \sum \text{power} =$		12,385

SPEC boils these numbers down to one number, called “**overall ssj_ops per watt.**”

The formula for this single summarizing metric is

$$\text{overall ssj_ops per watt} = \left(\sum_{i=0}^{10} \text{ssj_ops}_i \right) / \left(\sum_{i=0}^{10} \text{power}_i \right)$$

ssj_ops_i is performance at each 10% increment

power_i is power consumed at each performance level

FIGURE 1.19 SPECpower_ssj2008 running on a dual socket 2.2 GHz Intel Xeon Platinum 8276L with 192 GiB of DRAM and one 80 GB SSD disk.

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 9 – Matrix Multiply in Python

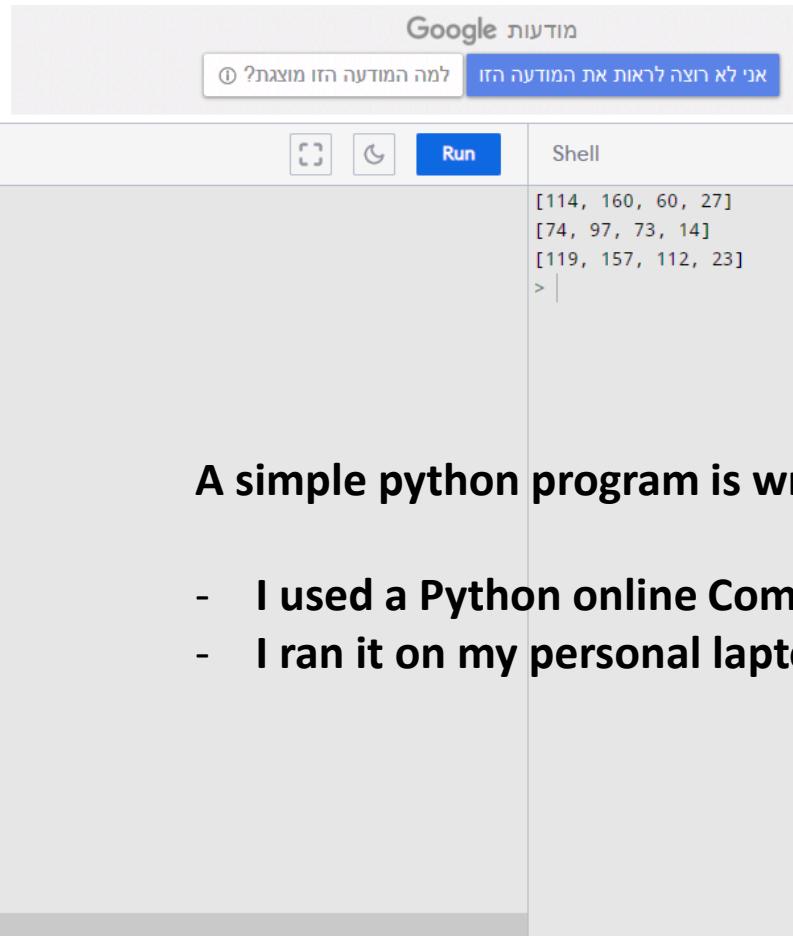
Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering

Python Program to Multiply Two Matrices

Programiz

Python Online Compiler



The screenshot shows a Python online compiler interface. On the left, there's a file selector with 'main.py' selected. Below it is a code editor containing the following Python code:

```
1 # Program to multiply two matrices using nested loops
2
3 # 3x3 matrix
4 X = [[12,7,3],
5     [4 ,5,6],
6     [7 ,8,9]]
7 # 3x4 matrix
8 Y = [[5,8,1,2],
9     [6,7,3,0],
10    [4,5,9,1]]
11 # result is 3x4
12 result = [[0,0,0,0],
13             [0,0,0,0],
14             [0,0,0,0]]
15
16 # iterate through rows of X
17 for i in range(len(X)):
18     # iterate through columns of Y
19     for j in range(len(Y[0])):
20         # iterate through rows of Y
21         for k in range(len(Y)):
22             result[i][j] += X[i][k] * Y[k][j]
23
24 for r in result:
25     print(r)
26
```

On the right, there's a 'Run' button and a 'Shell' output window. The output shows the resulting matrix:

```
[114, 160, 60, 27]
[74, 97, 73, 14]
[119, 157, 112, 23]
```

A simple python program is written to multiply two matrices

- I used a Python online Compiler
- I ran it on my personal laptop

```
for i in xrange(n):
    for j in xrange(n):
        for k in xrange(n):
            C[i][j] += A[i][k] * B[k][j]
```

Machine used	n1-standard-96 server in Google Cloud Engine, which has two Intel Skylake Xeon chips, and each chip has 24 processors or cores and running Python version 3.1.	
Size of Matrix	960 x 960	4096 x 4096
Time taken	5 minutes	6 hours

- ✓ If we used the Numpy library(python) instead, a 960 x 960 matrix multiply would take much less than 1 second instead of 5 minutes.  NumPy

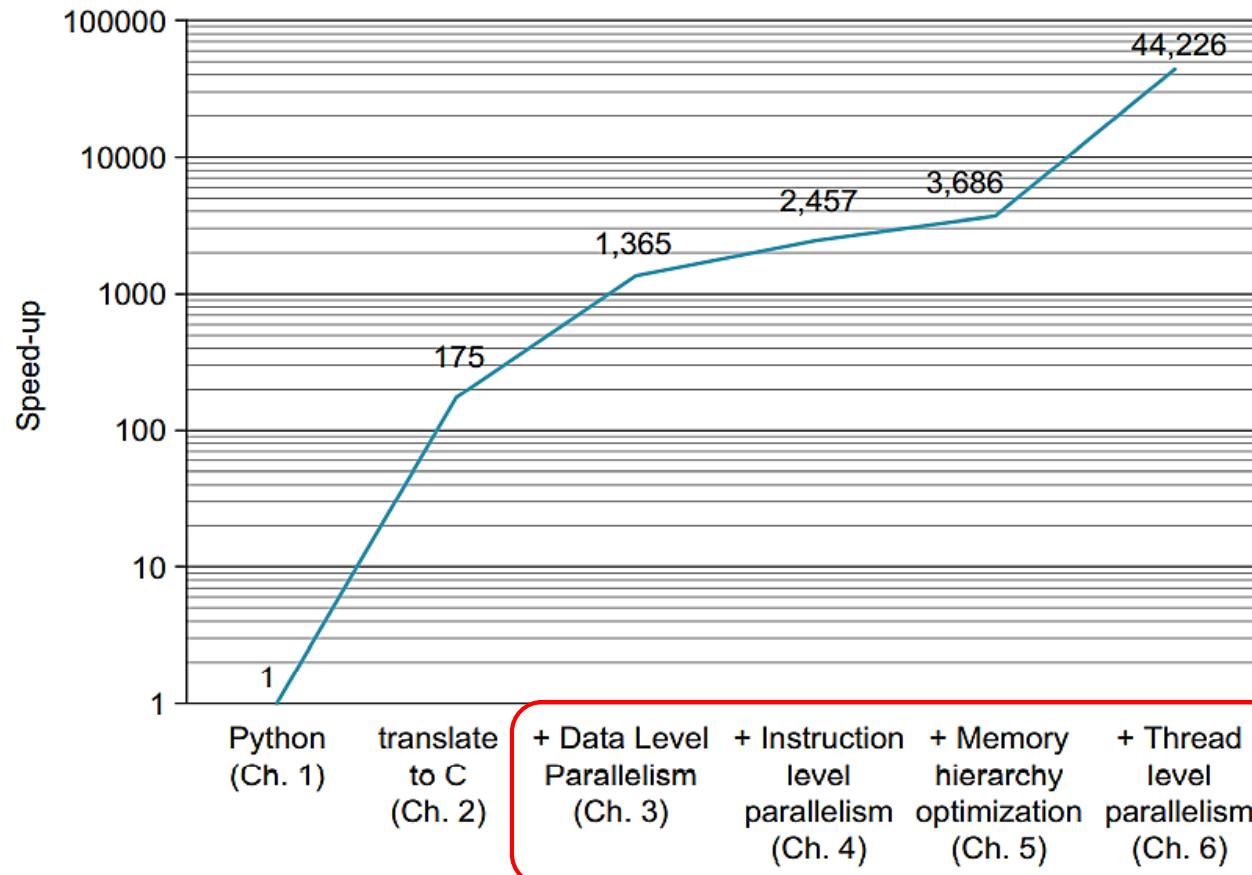
***While it is quick to write the matrix multiply in Python,
who wants to wait that long to get the answer?***



Python Program to Multiply Two Matrices - Optimizations

Text Book Chapter	Techniques used	Increase in performance by a factor of
Chapter 2	Convert the Python version of matrix multiply to a C version	200
Chapter 3	Use subword parallelism via C intrinsics	8
Chapter 4	Use loop unrolling to exploit multiple instruction issue and out-of-order execution hardware	2
Chapter 5	Use cache blocking to increase performance on large matrices	1.5
Chapter 6	Use parallel for loops in OpenMP to exploit multicore hardware	12 to 17

Python Program to Multiply Two Matrices - Optimizations



The last four steps leverage our understanding how the underlying hardware really works in a modern microprocessor and collectively only requires 21 lines of C code.

FIGURE 1.20 Optimizations of matrix multiply program in Python in the next five chapters of this book.



THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822

RISC V ARCHITECTURE

UNIT 1 – Computer Abstractions and Technology

Topic 10 – Fallacies and Pitfalls

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication Engineering



Pitfall

Expecting the improvement of one aspect of a computer to increase overall performance by an amount proportional to the size of the improvement.



Improvement in x by a factor f

=



Improvement in overall performance by a factor f

Suppose a program runs in 100 seconds on a computer, with multiply operations responsible for 80 seconds of this time.

How much do I have to improve the speed of multiplication if I want my program to run five times faster?



Fallacies and Pitfalls

Suppose a program runs in 100 seconds on a computer, with multiply operations responsible for 80 seconds of this time.

How much do I have to improve the speed of multiplication if I want my program to run five times faster? *Multiplication – 80 seconds*

$$\frac{\text{Execution time after improvement}}{\text{Execution time affected by improvement}} = \frac{\text{Amount of improvement}}{\text{Execution time unaffected}}$$

For this problem:

$$\text{Execution time after improvement} = \frac{80 \text{ seconds}}{n} + (100 - 80 \text{ seconds})$$

Since we want the performance to be five times faster, the new execution time should be 20 seconds, giving

$$20 \text{ seconds} = \frac{80 \text{ seconds}}{n} + 20 \text{ seconds}$$
$$0 = \frac{80 \text{ seconds}}{n}$$

That is, there is no amount by which we can enhance-multiply to achieve a fivefold increase in performance, if multiply accounts for only 80% of the workload.



Computers at low utilization use little power.

- Power efficiency matters at low utilizations because server workloads vary.
- Utilization of servers in Google's warehouse scale computer, for example, is between 10% and 50% most of the time and at 100% less than 1% of the time.
- Specially configured computer with the best results in 2020 still uses 33% of the peak power at 10% of the load.
- Systems in the field that are not configured for the SPECpower benchmark are surely worse.



Designing for performance and designing for energy efficiency are unrelated goals.

Since energy is power over time, it is often the case that hardware or software optimizations that take **less time save energy** overall even if the optimization takes a bit more energy when it is used.



Pitfall

Using a subset of the performance equation as a performance metric

- We have already warned about the danger of predicting performance based on simply one of the clock rate, instruction count, or CPI.
- Another common mistake is to use only two of the three factors to compare performance.

One alternative to time is MIPS (million instructions per second).

For a given program, MIPS is simply

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$$

Since MIPS is an instruction execution rate, MIPS specifies performance inversely to execution time;

faster computers have a higher MIPS rating.

There are three problems with using MIPS as a measure for comparing computers.



MIPS specifies the instruction execution rate but does not take into account the capabilities of the instructions.

We cannot compare computers with different instruction sets using MIPS, since the instruction counts will certainly differ.

There are three problems with using MIPS as a measure for comparing computers.



MIPS varies between programs on the same computer; thus, a computer cannot have a single MIPS rating.

$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$$

$$\text{MIPS} = \frac{\frac{\text{Instruction count}}{\text{Instruction count} \times \text{CPI}} \times 10^6}{\text{Clock rate}} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

Change in CPI changes the MIPS rating

There are three problems with using MIPS as a measure for comparing computers.



If a new program executes more instructions but each instruction is faster, MIPS can vary independently from performance!

Consider the following performance measurements for a program:

Measurement	Computer A	Computer B
Instruction count	10 billion	8 billion
Clock rate	4 GHz	4 GHz
CPI	1.0	1.1

- a. Which computer has the higher MIPS rating?
- b. Which computer is faster?





THANK YOU

Mahesh Awati & Vinay Reddy

Department of Electronics and Communication

mahesha@pes.edu

+91 9741172822